

02-control-plane

Control Plane — Go modular monolith, data model, events, coordinator, API

Revision: 1 **Last modified:** 2026-06-25T00:00:00Z

Master technical specification — document 02 of the HelixVPN set. Scope: the **Go control plane** (`helix-go`) — the brain that holds identity, topology and policy truth, compiles ACLs, and pushes desired-state network maps to every edge. This is a SPEC (describe the implementation; do not build the product). Source evidence cited inline by id, e.g. [04_P1 §2], [04_ARCH §3], [02_QWN], [SYNTHESIS §4].

0. Position in the system & what this document owns

HelixVPN is a self-hostable overlay network with a privacy-VPN front end: three roles — **Connector** (network-side outbound-only agent advertising CIDRs) $\rightleftharpoons$ **Gateway** (public VPS: control + data plane) $\rightleftharpoons$ **Client** (Mullvad-style end-user app) [04_ARCH §1.1, 05_YBO, SYNTHESIS §1]. This document specifies **only the control plane** — the Go services that decide *who may reach what* and *which transport to use*, and stream that decision to edges.

It does **not** specify the Rust data plane (doc 01 — moves already-encrypted bytes), the client UI (doc 05), or the platform tunnel shims (doc 06). The agent wire contract (Coordinator protobuf) is *introduced*

here because the control plane is its server, but its byte-level evolution semantics belong to doc 03 (WatchNetworkMap contract); this document gives the canonical .proto and the server-side handlers.

0.1 The governing constraints (every line below obeys these)

#	Invariant	Source
C1	Go is never in the packet path. The control plane computes and streams desired state; the Rust edge forwards bytes. If control is down, existing tunnels keep forwarding (fail-static).	[04_ARCH §2.1, SYNTHESIS §2]
C2	Postgres is the single source of truth for identity/topology/policy; Redis is ephemeral (presence + event bus). Losing Redis loses no durable state.	[04_P1, SYNTHESIS §2]
C3	No-logging by construction. The schema holds <i>no</i> connection/traffic/packet/flow table; a CI schema-lint fails the build if one appears. Only control actions are audited.	[04_ARCH §2.7/§7, 04_P1 §2, SYNTHESIS §7]
C4	Default-deny, need-to-know. A peer is delivered to a node only if a compiled policy rule grants it;	[04_ARCH §3.4/§7, 04_P1 §3]

#	Invariant	Source
	nodes never learn of peers they cannot reach.	
	Push, don't poll. Desired state is event-driven (Tailscale MapResponse-style): snapshot then deltas over a long-lived stream; convergence $p_{99} < 1$ s.	
C5		[04_ARCH §2, 04_P1 §4/§10]
	Device private keys never leave the device. The control plane registers only the 32-byte WG public key.	
C6		[04_P1 §6, SYNTHESIS §7]
	Package boundaries == future service boundaries. A modular monolith whose Phase-2 split is mechanical, not a rewrite.	
C7		[04_P1 §1]
	Multi-tenant isolation enforced at the database via Row-Level Security under a non-superuser role — not merely by <code>WHERE</code> clauses.	
C8		[04_P1 §2.2]

0.2 Settled stack floor for this layer

Go + Gin + PostgreSQL + Redis + Podman (rootless) [05_YBO mandated, confirmed by all 10 analyses, SYNTHESIS §2]. Postgres = truth; Redis = presence + bus; Connect (buf) for agent RPC; sqlc for

compile-checked SQL; goose/atlas for migrations; rootless Podman per constitution §11.4.161 + vasic-digital/containers submodule per §11.4.76 [SYNTHESIS §8].

1. Module architecture (the modular monolith)

One Go binary (helixd), many internal packages, deployed as one rootless container. Packages talk through **interfaces** + **events**, never by importing each other's stores [04_P1 §1, C7].

1.1 Repository / package layout

```
helix-go/
├── cmd/
│   ├── helixd/                # control-plane binary – wires all
modules
│   ├── helixvpctl/            # bootstrap + ops CLI (Cobra)
├── internal/
│   ├── identity/              # tenants, users, OIDC, anonymous
device-enroll tokens
│   ├── registry/              # devices, connectors, advertised
prefixes, presence
│   ├── ipam/                  # overlay IP allocation (ULA /48 per
tenant – D4)
│   ├── pki/                   # WG key registry, short-lived device
certs, rotation/revoke
│   ├── policy/                # ACL model + pure, versioned
compiler
│   ├── coordinator/           # in-mem topology graph, deltas,
WatchNetworkMap streams
│   ├── events/                # bus abstraction (Redis Streams MVP
→ NATS JetStream P2)
│   ├── telemetry/             # Prometheus metrics + control-action
audit sink
│   ├── api/                   # Gin REST (apps) + Connect (agents)
+ WS/SSE
│   └── store/                  # Postgres pool, RLS tenant-tx
helper, sqlc Queries, migrations
├── proto/                     # helix/agent/v1/*.proto → buf
generate (Go/Dart/Rust)
└── openapi/                   # REST contract → generated Dart/TS
```

```

app clients
└─ migrations/                                # goose SQL – the schema authority
(CI-linted, §2.4)

```

1.2 Wiring rules (enforced, not aspirational)

- **R1 — no cross-store imports.** coordinator MUST NOT import internal/store directly for another module's tables; it reads registry/policy/pki through their exported interfaces (§1.3). A pre-build gate (`go list -deps graph check`) fails on a forbidden edge.
- **R2 — coordinator owns no durable tables.** Its only state is the in-memory per-tenant graph + open streams (a cache hydrated from Postgres + events). [04_P1 §4]
- **R3 — every state mutation emits an event.** A write to Postgres that changes topology/policy is followed (same logical unit of work) by an `events.Publish` — the bus is the only path by which coordinator learns of change (§5).
- **R4 — every DB access is tenant-scoped.** No query runs outside `store.WithTenant` (§2.3); a `golangci-lint` custom analyzer flags `raw db.Query` outside store.

1.3 Inter-module interface surface (Go signatures)

```

// internal/registry/iface.go
type Registry interface {
    EnrollDevice(ctx context.Context, in EnrollInput) (Device,
        error)
    SetPrefixes(ctx context.Context, connectorID uuid.UUID, cidrs
        []netip.Prefix) ([]Conflict, error)
    DevicesForTenant(ctx context.Context, t uuid.UUID) ([]Device,
        error)
    MarkPresence(ctx context.Context, deviceID uuid.UUID, online
        bool, rttMS uint32) error
}

// internal/policy/iface.go – pure, deterministic, versioned (§7)
type Compiler interface {
    Compile(ctx context.Context, t uuid.UUID, spec Spec)
        (CompiledPolicy, error) // dry-run safe
    Activate(ctx context.Context, t uuid.UUID, version int64)
        error // flips active
}

type CompiledPolicy struct {
    Version int64
    VisibleTo map[DeviceID]map[DeviceID]struct{} // need-
        to-know peer set (C4)
}

```

```

AllowedIPs map[DeviceID]map[DeviceID]
    []netip.Prefix // coarse WG AllowedIPs
Verdicts    map[DeviceID]map[DeviceID][]PortRule // fine
    port-level edge verdicts
ExitNodes   map[DeviceID]struct{}
}

// internal/pki/iface.go
type PKI interface {
    IssueDeviceCert(ctx context.Context, deviceID uuid.UUID, ttl
        time.Duration) (Cert, error)
    Revoke(ctx context.Context, deviceID uuid.UUID) error // < 1s
        propagation (C5)
}

// internal/events/iface.go – bus-agnostic (D3: Redis Streams now,
    NATS later)
type Bus interface {
    Publish(ctx context.Context, stream string, env Envelope)
        (string, error)
    Subscribe(ctx context.Context, group, consumer string,
        streams ...string) (<-chan Envelope, error)
    Ack(ctx context.Context, stream, group, id string) error
}

```

Decision D3 — event bus [SYNTHESIS §3, 04_P1 §5].

Recommendation: Redis Streams for the MVP, NATS

JetStream for scale (Phase 2). The Bus interface above is the seam: the envelope (§5.2) is bus-agnostic, so the swap is a transport change, not a rewrite. Camp B (NATS from day one [DSK, KMI]) buys multi-region fan-out earlier but adds an operational dependency the self-host-first MVP does not need yet — Redis is already mandated for presence [05_YBO], so MVP reuses it. We adopt Camp A (CLD) and keep `events.Bus` as the abstraction.

2. Data model — Postgres DDL with Row-Level Security

The durable store holds **identity, topology, and policy** — never connection/traffic logs (C3). Eleven tenant-scoped tables + enums.

2.1 ER diagram

erDiagram

```
TENANTS ||--o{ USERS : "owns"
TENANTS ||--o{ DEVICES : "owns"
TENANTS ||--o{ GROUPS : "owns"
TENANTS ||--o{ POLICIES : "owns"
TENANTS ||--|| OVERLAY_POOLS : "has one"
TENANTS ||--o{ AUDIT_EVENTS : "records"
USERS ||--o{ DEVICES : "may own"
DEVICES ||--o| CONNECTORS : "is-a (kind=connector)"
CONNECTORS ||--o{ ADVERTISED_PREFIXES : "advertises"
GROUPS ||--o{ GROUP_MEMBERS : "contains"
DEVICES ||--o{ GROUP_MEMBERS : "member of"
DEVICES ||--o{ DEVICE_CERTS : "holds"
```

```
TENANTS {
    uuid id PK
    text name
    timestampz created_at
}
USERS {
    uuid id PK
    uuid tenant_id FK
    text email "nullable (anon)"
    text oidc_sub "nullable (SSO)"
    user_role role
}
DEVICES {
    uuid id PK
    uuid tenant_id FK
    uuid user_id FK "nullable"
    device_kind kind
    bytea wg_pubkey "32B Curve25519"
    inet overlay_ip "tenant ULA /48"
    timestampz last_seen_at "coarse"
    timestampz revoked_at
}
CONNECTORS {
    uuid device_id PK,FK
    uuid tenant_id FK
    text site_name
}
ADVERTISED_PREFIXES {
```

```

        uuid id PK
        uuid tenant_id FK
        uuid connector_id FK
        cidr cidr
        boolean enabled
    }
    GROUPS {
        uuid id PK
        uuid tenant_id FK
        text name
    }
    GROUP_MEMBERS {
        uuid group_id PK,FK
        uuid device_id PK,FK
    }
    POLICIES {
        uuid id PK
        uuid tenant_id FK
        jsonb spec
        bigint version "monotonic per tenant"
        timestampz compiled_at
    }
    OVERLAY_POOLS {
        uuid tenant_id PK,FK
        cidr ula_prefix "fd7a:helix:<rand>::/48"
        bigint next_host
    }
    DEVICE_CERTS {
        uuid id PK
        uuid tenant_id FK
        uuid device_id FK
        text serial
        timestampz not_after
        boolean revoked
    }
    AUDIT_EVENTS {
        bigint id PK
        uuid tenant_id FK
        text actor
        text action
        text target
        jsonb meta
    }
}

```


2.2 Core DDL

```
-- ===== tenancy & identity =====
CREATE TABLE tenants (
  id          uuid PRIMARY KEY DEFAULT gen_random_uuid(),
  name        text NOT NULL,
  created_at  timestampz NOT NULL DEFAULT now()
);

CREATE TYPE user_role AS ENUM ('admin','operator','member');

CREATE TABLE users (
  id          uuid PRIMARY KEY DEFAULT gen_random_uuid(),
  tenant_id   uuid NOT NULL REFERENCES tenants(id) ON DELETE
              CASCADE,
  email       text,                                -- nullable: anonymous
              device-token users (C6 privacy)
  oidc_sub    text,                                -- nullable: when SSO is used
  role        user_role NOT NULL DEFAULT 'member',
  created_at  timestampz NOT NULL DEFAULT now(),
  UNIQUE (tenant_id, email),
  UNIQUE (tenant_id, oidc_sub)
);

-- ===== devices (clients AND connectors) =====
CREATE TYPE device_kind AS ENUM ('client','connector');

CREATE TABLE devices (
  id          uuid PRIMARY KEY DEFAULT gen_random_uuid(),
  tenant_id   uuid NOT NULL REFERENCES tenants(id) ON DELETE
              CASCADE,
  user_id     uuid REFERENCES users(id) ON DELETE SET NULL,
  kind        device_kind NOT NULL,
  name        text NOT NULL,
  wg_pubkey   bytea NOT NULL,                      -- 32-byte Curve25519
              public key (C6: never the private)
  overlay_ip  inet NOT NULL,                       -- allocated from tenant ULA /48 (D4)
  os          text,                                -- ios|android|linux|
              windows|macos|harmonyos|aurora
  enrolled_at timestampz NOT NULL DEFAULT now(),
  last_seen_at timestampz,                          -- COARSE presence;
              refreshed from heartbeat, not per-packet (C3)
  revoked_at  timestampz,
  UNIQUE (tenant_id, wg_pubkey),
  UNIQUE (tenant_id, overlay_ip)
);
```

```
);
```

```
CREATE TABLE connectors (  
  device_id uuid PRIMARY KEY REFERENCES devices(id) ON DELETE  
    CASCADE,  
  tenant_id uuid NOT NULL REFERENCES tenants(id) ON DELETE  
    CASCADE,  
  site_name text NOT NULL  
);
```

```
CREATE TABLE advertised_prefixes (  
  id uuid PRIMARY KEY DEFAULT gen_random_uuid(),  
  tenant_id uuid NOT NULL REFERENCES tenants(id) ON DELETE  
    CASCADE,  
  connector_id uuid NOT NULL REFERENCES connectors(device_id) ON  
    DELETE CASCADE,  
  cidr cidr NOT NULL,  
  enabled boolean NOT NULL DEFAULT true,  
  created_at timestampz NOT NULL DEFAULT now()  
);  
CREATE INDEX ON advertised_prefixes (tenant_id, connector_id);
```

```
-- ===== groups (for policy) =====
```

```
CREATE TABLE groups (  
  id uuid PRIMARY KEY DEFAULT gen_random_uuid(),  
  tenant_id uuid NOT NULL REFERENCES tenants(id) ON DELETE  
    CASCADE,  
  name text NOT NULL,  
  UNIQUE (tenant_id, name)  
);  
CREATE TABLE group_members (  
  tenant_id uuid NOT NULL REFERENCES tenants(id) ON DELETE  
    CASCADE,  
  group_id uuid NOT NULL REFERENCES groups(id) ON DELETE  
    CASCADE,  
  device_id uuid NOT NULL REFERENCES devices(id) ON DELETE  
    CASCADE,  
  PRIMARY KEY (group_id, device_id)  
);
```

```
-- ===== policy (source + compiled marker) =====
```

```
CREATE TABLE policies (  
  id uuid PRIMARY KEY DEFAULT gen_random_uuid(),  
  tenant_id uuid NOT NULL REFERENCES tenants(id) ON DELETE  
    CASCADE,  
  spec jsonb NOT NULL,          -- the declarative ACL  
    document ($7)  
  version bigint NOT NULL,      -- monotonic per tenant
```

```

    active          boolean NOT NULL DEFAULT false,
    compiled_at     timestampz,
    created_at      timestampz NOT NULL DEFAULT now(),
    UNIQUE (tenant_id, version)
);
-- at most one active policy per tenant (instant rollback = re-
  activate an older version)
CREATE UNIQUE INDEX one_active_policy_per_tenant
  ON policies (tenant_id) WHERE active;

-- ===== overlay IPAM =====
CREATE TABLE overlay_pools (
  tenant_id  uuid PRIMARY KEY REFERENCES tenants(id) ON DELETE
    CASCADE,
  ula_prefix cidr NOT NULL,           -- e.g.
    fd7a:helix:<rand>::/48 (D4 - §3.1)
  next_host  bigint NOT NULL DEFAULT 2 -- ::1 reserved for
    gateway
);

-- ===== PKI =====
CREATE TABLE device_certs (
  id          uuid PRIMARY KEY DEFAULT gen_random_uuid(),
  tenant_id   uuid NOT NULL REFERENCES tenants(id) ON DELETE
    CASCADE,
  device_id   uuid NOT NULL REFERENCES devices(id) ON DELETE
    CASCADE,
  serial      text NOT NULL,
  not_after   timestampz NOT NULL,
  revoked     boolean NOT NULL DEFAULT false,
  created_at  timestampz NOT NULL DEFAULT now()
);
CREATE INDEX ON device_certs (tenant_id, device_id) WHERE NOT
  revoked;

-- ===== audit (CONTROL actions only – never traffic, C3)
  =====
CREATE TABLE audit_events (
  id          bigint GENERATED ALWAYS AS IDENTITY PRIMARY KEY,
  tenant_id   uuid NOT NULL REFERENCES tenants(id) ON DELETE
    CASCADE,
  actor       text NOT NULL,           -- user id / "system"
  action      text NOT NULL,           -- e.g.
    "device.revoke", "policy.activate"
  target      text,
  ts          timestampz NOT NULL DEFAULT now(),
  meta        jsonb

```

```
);
```

```
CREATE INDEX ON audit_events (tenant_id, ts DESC);
```

Design notes. - devices holds **both** clients and connectors (kind); connectors is a 1:1 detail table so a connector inherits all device machinery (enrollment, cert, revoke) [04_P1 §2]. The asymmetry (a connector advertises prefixes; a client does not) lives in advertised_prefixes. - group_members carries tenant_id redundantly (vs. the source's two-column PK) **so RLS can be enforced on it directly** — a refinement over [04_P1 §2.1], which left the join table un-RLS'd. The composite PK (group_id, device_id) is preserved. - policies.active + the partial unique index replace the source's implicit “newest is active”; this makes **instant rollback** a single UPDATE ... SET active=true WHERE version=\$n inside a tx that clears the prior active — see §7.4. - **Absent by design (C3):** no connections, sessions, flows, traffic, packets, netflow, bandwidth_samples. Their absence is a *guarantee*, lint-enforced (§2.4).

2.3 Row-Level Security (tenant isolation at the database, C8)

Every tenant-scoped table gets the same policy. The app sets app.tenant_id per transaction; Postgres enforces isolation even if a query forgets a WHERE [04_P1 §2.2].

```
-- applied to EVERY tenant-scoped table:
-- users, devices, connectors, advertised_prefixes, groups,
-- group_members,
-- policies, overlay_pools, device_certs, audit_events
ALTER TABLE devices ENABLE ROW LEVEL SECURITY;
ALTER TABLE devices FORCE ROW LEVEL SECURITY;
-- applies even to the table owner
CREATE POLICY tenant_isolation ON devices
  USING      (tenant_id = current_setting('app.tenant_id')::uuid)
  WITH CHECK (tenant_id =
    current_setting('app.tenant_id')::uuid);

-- the application connects as a NON-SUPERUSER, NON-OWNER role
-- (superusers/owners bypass RLS):
CREATE ROLE helix_app NOLOGIN;
GRANT SELECT, INSERT, UPDATE, DELETE ON ALL TABLES IN SCHEMA
  public TO helix_app;
GRANT USAGE, SELECT ON ALL SEQUENCES IN SCHEMA public TO
  helix_app;
-- helixd authenticates as a LOGIN role that is a member of
-- helix_app, never as the owner.
```

FORCE ROW LEVEL SECURITY is the hardening refinement over the source: it closes the gap where the table owner (which migrations run as) silently bypasses RLS. helixd therefore MUST run as helix_app (or a member), never as the migration owner. This is asserted by a startup check that runs `SELECT rolsuper, rolbypassrls FROM pg_roles WHERE rolname = current_user` and aborts if either is true.

The tenant-scoped transaction helper — the **only** way helixd touches the DB (R4):

```
// internal/store/tenant.go
//
// WithTenant runs fn inside a transaction whose RLS scope is
// pinned to tenantID.
// SET LOCAL is transaction-scoped, so it CANNOT leak across
// pooled connections.
func (s *Store) WithTenant(ctx context.Context, tenantID
    uuid.UUID,
    fn func(q *db.Queries) error) error {
    tx, err := s.pool.BeginTx(ctx, pgx.TxOptions{})
    if err != nil {
        return err
    }
    defer func() { _ = tx.Rollback(ctx) }() // no-op after a
        successful Commit

    // set_config(name, value, is_local=true) == SET LOCAL;
    // transaction-scoped.
    if _, err := tx.Exec(ctx,
        "SELECT set_config('app.tenant_id', $1, true)",
        tenantID.String()); err != nil {
        return fmt.Errorf("pin tenant: %w", err)
    }
    if err := fn(db.New(tx)); err != nil {
        return err
    }
    return tx.Commit(ctx)
}
```

A handful of operations (tenant creation, cross-tenant ops) need an explicit `WithSystem` helper that sets `app.tenant_id` to a sentinel and uses a separate `helix_sys` role with narrowly-granted privileges — never a superuser. Migrations run as the schema owner (`helix_owner`) out-of-band of request handling.

2.4 No-logging CI schema-lint (C3 enforced mechanically)

A pre-build gate parses migrations/ and the live schema and **fails the build** if any table, column, or view matches a forbidden traffic-logging shape [04_ARCH §7, 04_P1 §11.4, SYNTHESIS §7]:

```
// tools/schemalint/main.go (run in pre-build; constitution
    §11.4.27/§1.1)
var forbiddenTables = regexp.MustCompile(
    `(?!i)\b(connections?|sessions?|flows?|traffic|packets?|
        netflow|bandwidth_samples?|dns_queries?)\b`)
var forbiddenCols = regexp.MustCompile(
    `(?!i)\b(src_ip|dst_ip|dest_ip|src_port|dst_port|bytes_(in|
        out)|packet_count|payload|snr_host)\b`)
// Exception allowlist: overlay_ip + wg_pubkey on `devices`
    (identity, not traffic); cidr on
// advertised_prefixes (topology). The lint asserts these are the
    ONLY ip/port-shaped columns.
```

Paired §1.1 mutation: a test migration that adds a connections(src_ip, dst_ip, bytes) table **MUST** make the lint FAIL; removing it **MUST** make it pass. The lint is the runtime signature (§11.4.108) that proves C3 is live, not merely promised.

3. IPAM — overlay addressing & the subnet-collision decision (D4)

3.1 The problem and the decision

A single user joins N private RFC1918 networks (the HelixVPN differentiator [SYNTHESIS §1]). Two joined home/lab LANs both numbering 192.168.1.0/24 collide — the client cannot tell which 192.168.1.50 it means. Only three analyses actually solve this [SYNTHESIS §3 D4].

Decision D4 — IP-subnet collision. Recommendation: IPv6 ULA /48 per tenant + Tailscale 4via6 mapping (Camp A / CLD [04_ARCH §3.4]).

- **Camp A — ULA /48 + 4via6:** allocate each tenant a random fd7a:helix:<rand>::/48 (overlay_pools.ula_prefix). Every device gets a stable overlay address from it (devices.overlay_ip). Colliding IPv4 LANs are reached through a deterministic IPv6 encoding (4via6):

fd7a:helix:<rand>:<site-id>::/96 + the 32-bit IPv4 host → a unique IPv6 the client can always disambiguate. The control plane assigns a per-connector site-id and ships the 4via6 mapping inside the network map.

- **Camp B — CGNAT 100.64.0.0/10, 1:1 NAT per network**
[GMI, KMI]: each joined network is remapped 1:1 into a slice of the carrier-grade range at the connector; the client always sees 100.x. Simpler for IPv4-only clients but exhausts /10 space at scale and needs per-network NAT bookkeeping in the data path.

Why A: it is collision-proof by construction (a /48 is astronomically unlikely to clash), needs no stateful NAT in the packet path (keeps the edge dumb, C1), and is the documented Mullvad/Tailscale precedent. **Trade-off honestly stated:** legacy IPv4-only apps inside a client cannot dial a raw 192.168.1.50 without a local resolver shim that rewrites to the 4via6 address — the client core ships exactly that shim (doc 01 routing layer). For Phase 1 the lean path is: overlay addressing is IPv6-ULA; 4via6 for colliding LANs lands with multi-network policy in Phase 1, CGNAT remains the Phase-2 fallback for pure-IPv4 connectors.

3.2 Allocation algorithm (deterministic, gap-free, concurrency-safe)

```
// internal/ipam/alloc.go
//
// AllocOverlayIP hands out the next host address from the
//   tenant's ULA /48.
// Runs inside WithTenant; the UPDATE ... RETURNING is atomic so
//   two concurrent
// enrollments never get the same address (the row lock serialises
//   them).
func (a *IPAM) AllocOverlayIP(ctx context.Context, q *db.Queries,
    t uuid.UUID) (netip.Addr, error) {
    // SELECT ula_prefix, next_host FROM overlay_pools WHERE
    //   tenant_id=$1 FOR UPDATE;
    // UPDATE overlay_pools SET next_host = next_host + 1 WHERE
    //   tenant_id=$1 RETURNING next_host;
    pool, err := q.BumpOverlayNextHost(ctx, t)
    if err != nil {
        return netip.Addr{}, err
    }
    base := pool.UlaPrefix.Addr().As16()
    fd7a:helix:rrrr::
```

```

    host := uint64(pool.NextHost) //
        monotonic per tenant; ::1 = gateway
    // place the host id in the low 64 bits of the /48 (sites use
        bits [48..64))
    return embedHost(base, host), nil
}

```

overlay_pools.next_host is monotonic and never reused (a revoked device's address is *not* recycled in Phase 1 — simpler and avoids stale-route hazards; a /48 gives $\sim 2^{80}$ hosts, so exhaustion is not a real concern). Site-IDs for 4via6 are allocated from a parallel per-connector counter when a connector attaches.

4. Agent contract — protobuf (Coordinator over Connect)

The static map.json of Phase 0 becomes a streamed, versioned, delta-capable protobuf served over **Connect** (works as gRPC, gRPC-Web, and the Connect protocol — one service serves native agents over HTTP/2 and a future WASM Console over HTTP/1.1) [04_P1 §3]. The byte-level evolution rules are doc 03; the canonical .proto is here.

```

// proto/helix/agent/v1/agent.proto
syntax = "proto3";
package helix.coordinator.v1;
option go_package = "github.com/vasic-digital/helix-go/gen/helix/
    agent/v1;agentv1";

service Coordinator {
    // One-time device enrollment (exchanges an enroll token for
        identity + cert).
    rpc Enroll(EnrollRequest) returns (EnrollResponse);

    // The spine of the system: open once, get a snapshot, then a
        delta stream (C5).
    rpc WatchNetworkMap(WatchRequest) returns (stream MapUpdate);

    // Connectors push their advertised prefixes (also settable via
        Console/API).
    rpc AdvertisePrefixes(AdvertiseRequest) returns
        (AdvertiseResponse);

    // Lightweight heartbeat / status (presence, current transport,
        rtt). NO traffic data (C3).
    rpc ReportStatus(StatusReport) returns (StatusAck);
}

```



```

}

message EnrollRequest {
    string    enroll_token = 1;    // issued by Console/identity,
    short-lived, single-use
    bytes     wg_pubkey     = 2;    // device-generated; private key
    never leaves device (C6)
    string    os            = 3;
    string    name          = 4;
    DeviceKind kind         = 5;
}

enum DeviceKind { DEVICE_KIND_UNSPECIFIED = 0; CLIENT = 1;
                  CONNECTOR = 2; }

message EnrollResponse {
    string    device_id     = 1;
    string    overlay_ip    = 2;    // allocated by IPAM, e.g.
    "fd7a:helix:1::2" (D4)
    bytes     device_cert   = 3;    // short-lived mTLS cert for the
    control channel
    GatewayInfo gateway     = 4;
}

message WatchRequest {
    string    device_id     = 1;
    uint64    known_version = 2;    // 0 => full snapshot; else =>
    deltas since (cheap reconnect)
}

// Either a full snapshot or an incremental delta. Agents
// reconcile to it.
message MapUpdate {
    uint64    version = 1;
    oneof body {
        NetworkMap snapshot = 2;
        MapDelta    delta   = 3;
        KeepAlive    keepalive = 4;    // periodic; proves liveness
        without state churn
    }
}

message NetworkMap {
    Node      self          = 1;
    GatewayInfo gateway     = 2;
    repeated Peer peers     = 3;    // peers this node MAY reach
    (already policy-filtered, C4)
    DnsConfig dns           = 4;
}

```

```

    TransportPolicy transport = 5;
}

message Node { string device_id = 1; string overlay_ip = 2; }
message GatewayInfo {
    string endpoint = 1; // "gw.example:443"
    bytes wg_pubkey = 2;
    string masque_sni = 3; // host to present for MASQUE/
    HTTP-3 masquerade (doc 01, D1)
}
message Peer {
    string device_id = 1;
    bytes wg_pubkey = 2;
    repeated string allowed_ips = 3; // compiled from policy:
    only what this node may reach (C4)
    string endpoint = 4; // via gateway relay in
    MVP; direct candidates in P2
    bool is_connector = 5;
    repeated Via6Route via6 = 6; // 4via6 mappings for
    colliding IPv4 LANs (D4)
}
message Via6Route { string ipv4_cidr = 1; string via6_prefix =
    2; }
message DnsConfig { repeated string resolvers = 1; repeated
    string search = 2; }

message TransportPolicy {
    // ordered escalation ladder the client walks on handshake
    failure (doc 01 §Transport)
    repeated string order = 1; // ["plain-
    udp", "lwo", "masque-h3"]
    bool allow_user_override = 2;
}

message MapDelta {
    repeated Peer upsert_peers = 1;
    repeated string remove_peer_ids = 2;
    TransportPolicy transport = 3; // present only if changed
    DnsConfig dns = 4; // present only if changed
}

message AdvertiseRequest { string device_id = 1; repeated string
    cidrs = 2; }
message AdvertiseResponse { bool accepted = 1; repeated string
    conflicts = 2; }

message StatusReport {
    string device_id = 1;

```

```

string transport = 2;           // current transport in use
                                (presence/health only)
uint32 rtt_ms      = 3;
// deliberately NO bytes/flows/destinations – presence + health
                                only (C3)
}
message StatusAck {}
message KeepAlive {}

```

Load-bearing properties [04_P1 §3]: `known_version` makes reconnects cheap (resume from a version, no full resync); the server sends **snapshot then deltas**; peers arrive **already policy-filtered** so a client never learns of nodes it cannot reach (need-to-know, C4). `buf` generate emits Go (server), Dart (client core), and Rust (edge/connector) stubs from this one file — no hand-written clients, so the codebases cannot drift [04_ARCH §4.2].

5. Event backbone — Redis Streams contracts (D3)

Redis Streams is the MVP bus; the §1.3 `events.Bus` interface keeps the envelope bus-agnostic so NATS JetStream is a Phase-2 transport swap [04_P1 §5, SYNTHESIS §3].

5.1 Streams & consumer groups

Stream	Producers	Consumer groups
<code>events:devices</code>	identity, registry	coordinator, telemetry, audit
<code>events:routes</code>	registry	coordinator, telemetry
<code>events:policy</code>	policy	coordinator, audit
<code>events:presence</code>	api (heartbeats)	coordinator (TTL/online state)
<code>events:gateway</code>	edge health probes	coordinator, telemetry

5.2 Envelope (every event — bus-agnostic)

```

{
  "id":          "<redis-stream-id>",

```

```

"type":      "device.revoked",
"tenant_id": "uuid",
"ts":       "RFC3339",
"actor":    "user-uuid|system",
"payload":  { "...": "type-specific" },
"trace_id": "for correlation"
}

```

5.3 Event taxonomy (concrete)

Type	Payload	Coordinator reaction
device.enrolled	{device_id, kind, overlay_ip}	add node; if policy grants, push to peers' maps
device.online / .offline	{device_id}	update presence; peers see relay availability
device.revoked	{device_id}	remove node; push peer-removal delta to everyone who saw it; edge drops sessions
connector.attached	{device_id, site}	register connector + allocate site-id (D4)
connector.prefixes.changed	{connector_id, cidrs[]}	recompute routes; push to nodes whose policy includes it
route.conflict.detected	{cidr, connector_ids[]}	flag overlapping-CIDR; surface in Console (§7.3)
policy.updated	{version}	trigger compile (dry-run already done)
policy.compiled	{version}	

Type	Payload	Coordinator reaction
gateway.failover	{from, to}	recompute tenant visibility; push deltas re-point affected nodes' gateway endpoint

5.4 Producer / consumer mechanics + DLQ

```
// produce (registry/policy/etc. via events.Bus)
xid, _ := rdb.XAdd(ctx, &redis.XAddArgs{Stream: "events:policy",
    Values: env}).Result()

// consume – durable consumer group, at-least-once
res, _ := rdb.XReadGroup(ctx, &redis.XReadGroupArgs{
    Group:  "coordinator", Consumer: hostID,
    Streams: []string{"events:policy", ">"}, Count: 64, Block: 5
    * time.Second,
}).Result()
// handle each, then XACK. Stuck/un-acked entries (a crashed
// consumer) are reclaimed by a
// background sweeper using XAUTOCLAIM – the dead-letter recovery
// path:
claimed, _, _ := rdb.XAutoClaim(ctx, &redis.XAutoClaimArgs{
    Stream: "events:policy", Group: "coordinator", Consumer:
    hostID,
    MinIdle: 30 * time.Second, Start: "0-0", Count: 64,
}).Result()
// After N delivery attempts (PEL entry's delivery count), route
// to events:policy:dlq and alert
// (telemetry counter helix_events_dlq_total) rather than spin
// forever.
```

At-least-once delivery + **idempotent reactions** (every delta is recomputed from current graph state, so a replayed event is harmless) [04_P1 §5.4]. XAUTOCLAIM + a per-entry delivery-count ceiling gives a real dead-letter path — replacing Phase 0's file-watch and any cron-restart loop. This composes with constitution §11.4.147 (no-work-loss): a crashed coordinator's PEL entries are reclaimed, never silently dropped.

6. The coordinator — building & streaming maps

The coordinator is the brain. It owns no durable tables (R2); it **reacts to events**, recomputes affected maps, and pushes minimal deltas down open streams [04_P1 §4].

6.1 Responsibilities

1. Maintain an in-memory, per-tenant **topology graph** (devices, connectors, prefixes, groups, compiled policy), hydrated from Postgres on boot and kept fresh by events.
2. Hold the set of **open WatchNetworkMap streams** (one per connected agent).
3. On each relevant event: determine the **affected agents**, compute each one's **delta**, push it. Convergence target: **p99 < 1 s** from event to delta-on-wire (a measured SLO, §10).

6.2 Map computation (per node)

```
buildMap(node):  
    reachable := policy.PeersVisibleTo(node)           # need-to-know  
    filter (C4)  
    peers := []  
    for p in reachable:  
        allowed := policy.AllowedIPs(node -> p)       # compiled  
        CIDRs+ports  
        via6 := policy.Via6Routes(node -> p)          # 4via6 for  
        colliding IPv4 LANs (D4)  
        peers.append(Peer{p.pubkey, allowed, relayEndpoint(p),  
p.isConnector, via6})  
    return NetworkMap{self, gateway, peers, dns(node),  
transportPolicy(node)}
```

relayEndpoint in MVP always routes peer traffic **through the gateway** (hub-and-spoke); direct P2P is a Phase-2 optimization [04_P1 §4.2]. This keeps the MVP data path identical to the proven Phase 0 slice.

6.3 Streaming + delta loop (Go)

```
// internal/coordinator/watch.go  
func (c *Coordinator) WatchNetworkMap(ctx context.Context,  
    req *connect.Request[agentv1.WatchRequest],
```

```

stream *connect.ServerStream[agentv1.MapUpdate]) error {

dev := authDevice(ctx) // from mTLS
    device cert (§8.2)
if dev.Revoked { // belt-and-
    suspenders: revoked never streams
    return connect.NewError(connect.CodePermissionDenied,
        errRevoked)
}
sub := c.subscribe(dev.TenantID, dev.ID) // registers
    stream + flips presence online
defer
    sub.Close() // flips presence
    offline on disconnect

// 1. snapshot (or resume from known_version – diff against
    the client's last seen)
cur := c.version(dev.TenantID)
if req.Msg.KnownVersion == 0 || req.Msg.KnownVersion > cur {
    if err := stream.Send(&agentv1.MapUpdate{Version: cur,
        Body: &agentv1.MapUpdate_Snapshot{Snapshot:
            c.buildMap(dev)}}); err != nil {
        return err
    }
} else {
    for _, d := range c.deltasSince(dev,
        req.Msg.KnownVersion) { // catch-up deltas
        if err := stream.Send(deltaUpdate(d)); err != nil {
            return err
        }
    }
}

// 2. delta loop: events arrive on sub.C, keepalive on a
    ticker
ka := time.NewTicker(20 * time.Second); defer ka.Stop()
for {
    select {
    case <-ctx.Done(): return ctx.Err()
    case d := <-sub.C: if err :=
        stream.Send(deltaUpdate(d)); err != nil { return err }
    case <-ka.C: if err := stream.Send(keepAlive());
        err != nil { return err }
    }
}
}

```

6.4 Fan-out: events → minimal affected set

The coordinator consumes the event streams (§5), updates its graph, then computes which open streams are affected and enqueues a delta to each. A `policy.compiled` event recomputes visibility for the whole tenant; a `connector.prefixes.changed` event touches only nodes whose policy grants that connector. **Compute the minimal affected set — never broadcast a full resync on a small change** [04_P1 §4.4]. Each device’s per-stream send queue is bounded; a slow consumer that fills its queue is dropped and forced to reconnect-with-snapshot (back-pressure, not memory growth — the §10 24h-soak SLO).

7. Policy model & compiler

Tailscale-ACL-flavored, declarative, **default-deny**, **fail-closed** [04_P1 §7]. This is the access-control brain that turns “1 user → N networks” into compiled per-node visibility.

7.1 Source document (stored in `policies.spec.jsonb`)

```
{
  "groups": {
    "group:admins": ["alice@corp", "bob@corp"],
    "group:contractors": ["carol@ext"]
  },
  "hosts": {
    "warehouse-cams": "10.10.0.0/24",          // served by connector
A    "office-lan":      "192.168.50.0/24"      // served by connector
B
  },
  "acls": [
    { "action": "accept", "src": ["group:admins"], "dst":
["*:*"] },
    { "action": "accept", "src": ["group:contractors"], "dst":
["warehouse-cams:554,80"] }
  ],
  "exitNodes": ["group:admins"]                // who may use the
gateway as a full-tunnel exit
}
```


7.2 Compilation algorithm (pure, deterministic, versioned)

```
compile(tenant, spec) -> CompiledPolicy:
  resolve groups -> sets of device_ids (via users + group_members)
  resolve hosts -> CIDRs (cross-check vs advertised_prefixes;
emit route.conflict if ambiguous)
  for each device d:
    visible[d] = {}                                # peers d may reach
    (need-to-know, C4)
    for rule in acls where d in resolve(rule.src):
      for dst in rule.dst:
        targets = expand(dst)                      # connector(s)
serving the CIDR, peer devices, or exit
      for t in targets:
        visible[d] += t
        allowedIPs[d][t] += (dst.cidr)            # coarse
WG AllowedIPs (CIDR only)
        verdicts[d][t] += (dst.cidr, dst.ports) # fine
port-level edge verdict map
  return { visible, allowedIPs, verdicts, exitNodes }
```

- Two compiled artifacts, because **WG AllowedIPs is CIDR-only, not port-aware** [04_P1 §7.2]:
 1. coarse per-peer AllowedIPs on the data path; (2) a fine **port-level verdict map** enforced at the edge via nftables/eBPF (doc 01 §routing). Output is exactly what `coordinator.buildMap` consumes.
- **Pure + versioned**: same spec $\Rightarrow$ same output (a property test asserts determinism). Bump `policies.version`, emit `policy.compiled`, `coordinator` diffs and pushes.

7.3 Validation (fail-closed)

`policy.update` runs the compiler **dry-run first** and rejects on: unknown group/host; a dst CIDR not covered by any `advertised_prefixes` row; a rule that would grant a *revoked* device; or an `exitNodes` entry resolving to a connector (connectors are not exits). Only a clean dry-run inserts the new `policies` row [04_P1 §7.3]. A `route.conflict.detected` (two connectors advertising overlapping CIDRs) does not block compile but surfaces in the Console and is resolved by `4via6` site disambiguation (D4) or operator choice.

7.4 Activation & instant rollback

```
// internal/policy/activate.go – atomic flip; prior version stays
// compiled for instant rollback
func (p *PolicyService) Activate(ctx context.Context, t
    uuid.UUID, version int64) error {
    return p.store.WithTenant(ctx, t, func(q *db.Queries) error {
        if err := q.DeactivateAllPolicies(ctx, t); err != nil {
            return err } // active=false
        if err := q.ActivatePolicyVersion(ctx,
            db.ActivateParams{t, version}); err != nil {
            return
            err //
            active=true, sets compiled_at
        }
        // R3: emit on the bus so the coordinator recomputes +
        // pushes
        _, err := p.bus.Publish(ctx, "events:policy",
            events.New("policy.compiled", t, "system",
                map[string]any{"version": version}))
        return err
    })
}
```

The partial unique index `one_active_policy_per_tenant` (§2.2) guarantees the flip is atomic. A bad policy is rolled back by `Activate(t, olderVersion)` — instant, no recompile, because every version's compiled artifact is reproducible from its spec.

8. API surface (Gin REST + Connect + WS/SSE) & authz

One multiplexed HTTP server [04_P1 §8].

Audience	Protocol	Examples
Apps (Access / Connector / Console)	REST via Gin	POST /v1/enroll-tokens, GET /v1/devices, POST /v1/ policies, POST /v1/policies/ {v}/activate, GET /v1/ networks
Live UI	WebSocket / SSE	GET /v1/stream → device.online, route.changed, handshake.failing

Audience	Protocol	Examples
Agents	Connect/gRPC	Coordinator.WatchNetworkMap,.Enroll,.Advertis

- **One server, multiplexed:** Gin serves REST/WS; Connect handlers mount alongside on the same listener (Connect speaks HTTP/2 for native agents, downgrades to HTTP/1.1 for browsers — so a future WASM Console calls the same Coordinator service) [04_P1 §8].
- **Contracts generated:** OpenAPI → Dart/TS app clients; buf generate → Go/Dart/Rust agent stubs. No hand-written clients [04_ARCH §4.2].

8.1 REST authz (RBAC + RLS backstop)

REST uses OIDC session **or** API token, then RBAC (admin > operator > member); RLS is the database-level backstop (C8). Gin middleware chain:

```
// internal/api/middleware.go
r.Use(authn())           // validate OIDC/session/API-token → sets
                        ctx user + tenant
r.Use(rbac())            // role gate per route group
r.POST("/v1/policies", requireRole("admin","operator"),
      h.CreatePolicy)
r.POST("/v1/devices/:id/revoke", requireRole("admin"),
      h.RevokeDevice)
r.GET ("/v1/devices", requireRole("admin","operator","member"),
      h.ListDevices)

// every handler runs its DB work through
      WithTenant(ctx.TenantID, ...): RLS is the floor even
// if rbac() is misconfigured – defense in depth.
```

Route group	admin	operator	member
policies (create/activate)	✓	✓	×
devices (revoke)	✓	×	×
devices/networks (read)	✓	✓	✓
enroll-tokens (mint)	✓	✓	×

8.2 Agent authz (device mTLS)

Agent RPCs (WatchNetworkMap/AdvertisePrefixes/ReportStatus) authenticate with the short-lived **device mTLS cert** issued at enrollment (§9). Enroll is the only unauthenticated RPC — it instead validates a single-use, hashed **enroll token**. authDevice(ctx) resolves

the cert serial → device_certs row → devices row, and rejects if revoked or expired. A revoked device's open stream is force-closed within the convergence SLO (§9.3).

9. Identity, enrollment & PKI

9.1 Two identity modes (both MVP) [04_P1 §6.1]

- **Managed (OIDC):** tenant admins log into the Console via any OIDC IdP (Keycloak/Authentik); users belong to a tenant, devices to users.
- **Anonymous (privacy mode):** a tenant mints **device enroll tokens** with no email/SSO — the Mullvad “account number, no PII” posture (C6). The device gets identity + cert with zero personal data stored (users.email/oidc_sub stay NULL).

9.2 Enrollment flow (the device never surrenders its private key, C6)

sequenceDiagram

autonumber

participant Console

participant Identity as identity (helixd)

participant IPAM as ipam

participant PKI as pki

participant Bus as events bus

participant Device as Access/Connector

Console->>Identity: create enroll token (admin/operator)

Identity-->>Identity: store HASHED single-use token

Identity-->>Console: token / QR

Console-->>Device: operator hands token (paste / scan)

Device-->>Device: generate WG keypair locally (private NEVER leaves, C6)

Device->>Identity: Enroll(token, wg_pubkey, os, kind)

[Connect RPC]

Identity-->>Identity: verify+consume token (atomic)

Identity->>IPAM: AllocOverlayIP(tenant)

IPAM-->>Identity: overlay_ip (ULA /48, D4)

Identity->>PKI: IssueDeviceCert(device, 24h)

PKI-->>Identity: short-lived mTLS cert

Identity-->>Identity: persist device row (WithTenant)

Identity->>Bus: publish device.enrolled

Identity-->>Device: EnrollResponse{device_id, overlay_ip,

```
cert, gateway}
```

```
Device->>Identity: open WatchNetworkMap (mTLS w/ cert)
```

9.3 PKI specifics [04_P1 §6.3]

- **WG keys:** device-generated; only the **public** key is registered. The control plane never sees a private key (C6).
 - **Control-channel mTLS:** pki issues a short-lived (≈ 24 h) device cert; the agent uses it to authenticate the RPC channel and auto-renews over the existing channel before expiry.
 - **Revoke (< 1 s):** `device.revoked` (a) removes the device's WG peer from every relevant map (delta push), (b) is enforced at the edge (kernel WG peer removed), (c) marks the cert revoked. Revocation latency target == convergence SLO (§10).
 - **Root / KMS:** the tenant CA key is the one true secret; back it up (KMS or offline). It and Postgres are the only stateful things to protect (architecture §10). Credentials handled per constitution §11.4.10 (never git-tracked, never logged).
-

10. End-to-end reconciliation (the < 1 s promise) & SLOs

10.1 Policy-apply reconciliation sequence

```
sequenceDiagram
```

```
    autonumber
```

```
    participant Admin as Console (admin)
```

```
    participant API as api (Gin)
```

```
    participant Policy as policy
```

```
    participant Store as Postgres (RLS)
```

```
    participant Bus as Redis Streams
```

```
    participant Coord as coordinator
```

```
    participant Edge as Rust edge / agents
```

```
    Admin->>API: POST /v1/policies (new spec)
```

```
    API->>Policy: Compile(dry-run) -- fail-closed (§7.3)
```

```
    Policy->>Store: INSERT policies(version, spec) [WithTenant]
```

```
    Admin->>API: POST /v1/policies/{v}/activate
```

```
    API->>Policy: Activate(v) -- atomic flip (§7.4)
```

```
    Policy->>Store: deactivate all + activate v (one tx)
```

```
    Policy->>Bus: XADD events:policy {policy.compiled, v}
```

```
    Bus-->>Coord: XReadGroup(coordinator) delivers event
```

```

Coord-->>Coord: recompute tenant visibility (minimal affected
set, §6.4)
  loop each affected open WatchNetworkMap stream
    Coord-->>Edge: stream.Send(MapDelta) -- upsert/remove
peers, AllowedIPs
  end
  Edge-->>Edge: reconcile WG peers + nftables/eBPF verdict map
  Note over Admin,Edge: target p99 < 1s, event→enforced, ZERO
restarts

```

This turns the architecture promise (“policy change reflected on all affected edges in < 1 s, no restarts” [04_ARCH]) into a concrete, measurable behavior — the event-driven, authorized descendant of Phase 0’s G6 file-delta reconcile [04_P1 §9].

10.2 Measured SLOs (the anti-bluff acceptance numbers) [04_P1 §10]

Metric	Target	How measured
event → delta-on-wire	p99 < 1 s	coordinator histogram helix_reconcile_seconds from event-receive to stream.Send
device revoke → edge enforced	< 1 s	revoke-to-WG-peer-removed timer on the edge
coordinator memory @ 10k streams	bounded, no leak over 24 h soak	process_resident_memory_bytes slope ≈ 0
enrollment round-trip	< 500 ms	API histogram
policy compile (1k devices)	< 200 ms	compiler benchmark

10.3 Test strategy (per constitution §11.4.27 / §11.4.5 / §1.1)

- **Unit:** policy compiler determinism (same spec ⇒ byte-identical CompiledPolicy); IPAM no-double-alloc under concurrency.
- **Store:** RLS tests — assert tenant A literally cannot read tenant B’s rows even with a crafted query and FORCE ROW LEVEL SECURITY on, running as helix_app [04_P1 §10].
- **Integration:** spin Postgres + Redis on-demand via the vasic-digital/containers submodule (§11.4.76 — *not* ad-hoc docker run); drive enroll → advertise → policy → WatchNetworkMap, assert the

delta stream content and the < 1 s SLO with captured evidence (§11.4.69/§11.4.107).

- **Soak:** N simulated agents holding streams; flap policies; assert convergence SLO and zero coordinator memory growth over 24 h.
 - **Schema-lint meta-test:** the §2.4 §1.1 mutation (add a connections table → lint FAILs).
-

11. Container & deployment manifests (rootless Podman, §11.4.161/§11.4.76)

The control plane deploys as rootless Podman quadlets [SYNTHESIS §2/§8]. Manifests illustrate the runtime shape; `helixvpnc1 bootstrap` renders them.

11.1 Podman quadlet (the MVP deploy unit)

```
# deploy/quadlet/helixd.container (rootless; read-only rootfs;
NET no extra caps)

[Unit]
Description=HelixVPN control plane
After=helix-postgres.service helix-redis.service
Requires=helix-postgres.service helix-redis.service

[Container]
Image=ghcr.io/vasic-digital/helixd:helixvpn-0.1.0 # §11.4.151
release-prefixed tag
# helixd connects to Postgres as the NON-superuser, NON-owner role
so RLS is enforced (C8):
Environment=HELIX_DB_URL=postgres://helix_app_login@helix-
postgres/helix?sslmode=verify-full
Environment=HELIX_REDIS_URL=redis://helix-redis:6379/0
PublishPort=443:8443
ReadOnly=true
NoNewPrivileges=true
DropCapability=ALL
Volume=helix-ca:/var/lib/helix/ca:Z # tenant CA key (KMS-
backed in prod)

[Service]
Restart=always

[Install]
WantedBy=default.target
```

11.2 Docker Compose (developer-loop equivalent)

```
# deploy/compose/dev.yml – local dev only; prod is rootless Podman
  quadlets above
services:
  postgres:
    image: postgres:16
    environment:
      POSTGRES_DB: helix
      # migrations run as owner; helixd connects as
      helix_app_login (RLS-enforced)
    command: ["postgres", "-c", "row_security=on"]
  redis:
    image: redis:7
  helixd:
    build: ../../helix-go
    depends_on: [postgres, redis]
    environment:
      HELIX_DB_URL: postgres://helix_app_login@postgres/helix?
        sslmode=disable
      HELIX_REDIS_URL: redis://redis:6379/0
    ports: ["8443:8443"]
```

11.3 Kubernetes (Phase-2 HA shape — stateless coordinators)

```
# deploy/k8s/helixd-deployment.yaml – Phase 2: coordinators are
  stateless (graph rebuilt from PG
# on boot + events), so they scale horizontally; Patroni-managed
  PG + NATS JetStream replace
# single PG + Redis. Shown to prove the Phase-1 seams permit it
  without reshaping. [04_P2]
apiVersion: apps/v1
kind: Deployment
metadata: { name: helixd }
spec:
  replicas: 3
  selector: { matchLabels: { app: helixd } }
  template:
    metadata: { labels: { app: helixd } }
    spec:
      containers:
        - name: helixd
          image: ghcr.io/vasic-digital/helixd:helixvpn-0.2.0
          securityContext: { runAsNonRoot: true,
            readOnlyRootFilesystem: true,
```



```

allowPrivilegeEscalation: false,
capabilities: { drop: ["ALL"] } }
env:
  - { name: HELIX_DB_URL, valueFrom: { secretKeyRef: {
name: helix-db, key: url } } }
  - { name: HELIX_BUS_URL, valueFrom: { secretKeyRef: {
name: helix-bus, key: url } } }
ports: [ { containerPort: 8443 } ]

```

12. Helix-ecosystem & constitution wiring (this layer)

Submodule / clause	Use in the control plane
vasic-digital/containers (§11.4.76/§11.4.161)	sole container orchestration: deploy quadlets + boot Postgres/ Redis on-demand for integration tests — no ad-hoc docker/podman [SYNTHESIS §8]
helix_qa + challenges (§11.4.27/ §11.4.5/.69/.107)	anti-bluff QA: the §10.3 integration/soak suites + a Challenge that drives enroll → policy → delta and asserts the < 1 s SLO with captured evidence
docs_chain (§11.4.106)	mechanizes this spec's HTML/ PDF/DOCX export + workable- items (phases → tasks become items) sync
security (§7)	reviews RLS/mTLS/CA-key handling; pairs with the §2.4 schema-lint
§11.4.93/.95	each Phase-1 task/subtask below becomes a git-tracked workable- items DB row
§11.4.151/.155	release-prefixed image/tag (helixvpn-0.1.0) + recording prefixes for QA evidence
§11.4.156	NO active CI — gates run locally (pre-build + meta-test)

13. Phase → task → subtask plan (control plane, Phase 1)

Derived from [04_P1] and the §0.1 invariants. Each leaf is a workable item (§11.4.93). Ordering is risk-descending (§11.4.132): RLS + no-log lint first (they are the irreversible correctness floor).

- **CP-T1 — store + schema + RLS (the truth floor).**
 - CP-T1.1 migrations/ goose DDL (§2.2) + enums; helix_app/ helix_owner/helix_sys roles.
 - CP-T1.2 FORCE ROW LEVEL SECURITY + tenant_isolation on all 10 tenant tables; startup rolsuper/rolbypassrls abort check.
 - CP-T1.3 store.WithTenant / WithSystem; sqlc Queries; RLS cross-tenant denial test.
 - CP-T1.4 tools/schemalint (§2.4) + paired §1.1 mutation (add connections → FAIL).
- **CP-T2 — proto + Connect server skeleton.**
 - CP-T2.1 proto/helix/agent/v1/agent.proto (§4) + buf generate (Go/Dart/Rust).
 - CP-T2.2 Connect handlers mounted on the Gin listener; authDevice mTLS resolver (§8.2).
- **CP-T3 — identity + enrollment + PKI.**
 - CP-T3.1 OIDC + anonymous enroll-token mint/store-hashed/ consume-atomic (§9.1).
 - CP-T3.2 Enroll RPC: verify token → IPAM alloc → device-cert issue → persist → emit device.enrolled (§9.2).
 - CP-T3.3 device-cert rotation + revoke (< 1 s edge enforcement, §9.3).
- **CP-T4 — ipam (D4).**
 - CP-T4.1 overlay_pools ULA /48 provisioning on tenant create; AllocOverlayIP (§3.2).
 - CP-T4.2 per-connector site-id + 4via6 route emission for colliding IPv4 LANs.
- **CP-T5 — events bus (D3).**
 - CP-T5.1 events.Bus Redis Streams impl: XADD/XReadGroup/XACK + envelope (§5.2).
 - CP-T5.2 XAUTOCLAIM DLQ sweeper + helix_events_dlq_total (§5.4).
- **CP-T6 — policy compiler.**
 - CP-T6.1 spec parser + group/host resolution + advertised-prefix cross-check (§7.2).
 - CP-T6.2 two-artifact compile (AllowedIPs + port verdict map) + determinism property test.

- CP-T6.3 dry-run fail-closed validation (§7.3) + atomic activate/rollback (§7.4).
- **CP-T7 — coordinator.**
 - CP-T7.1 in-mem per-tenant graph hydrate-on-boot + event apply (§6.1).
 - CP-T7.2 WatchNetworkMap stream loop: snapshot/resume/delta/keepalive + bounded send queue (§6.3).
 - CP-T7.3 minimal-affected-set fan-out (§6.4) + helix_reconcile_seconds histogram.
- **CP-T8 — api (REST + WS/SSE) + authz.**
 - CP-T8.1 Gin REST routes + OpenAPI gen; RBAC middleware (§8.1).
 - CP-T8.2 GET /v1/stream WS/SSE live-event fan-out to the Console.
- **CP-T9 — helixvpncctl + deploy.**
 - CP-T9.1 Cobra bootstrap/migrate/tenant/enroll-token ops.
 - CP-T9.2 quadlet rendering (§11.1) + on-demand-infra integration harness via containers.
- **CP-T10 — SLO + QA gates.** §10.2 histograms wired; §10.3 integration + 24 h soak + Challenge; the 8-criterion MVP DoD asserted end-to-end [04_P1 §11].

13.1 MVP definition-of-done the control plane must satisfy [04_P1 §11, SYNTHESIS §4]

Self-host from zero; enroll connector + client; reach an authorized LAN host **and deny an unauthorized one**; auto-escalate transport when WG is blocked (coordinator ships the TransportPolicy ladder; the escalation itself is doc 01); **policy edit reflected < 1 s, no restart; revoke < 1 s**; kill-switch + DNS-leak (client side, gated by core state the map drives); **no durable connection log exists** (§2.4 lint green); the three apps drive it all via generated clients.

14. Phase-2 forward seams (why this is additive, not a rewrite) [04_P1 §12]

The MVP interfaces extend without reshaping: TransportPolicy.order gains shadowsocks/udp-over-tcp/daita; Peer.endpoint gains direct-P2P NAT-traversal candidates so traffic stops always relaying through the gateway; coordinator federates across regional instances over **NATS JetStream** (the events.Bus swap, D3); pki adds the PQ pre-shared layer;

the Connect service gains multi-hop path fields; and HarmonyOS/Aurora app shims slot under the unchanged Flutter UI + Rust core. Phase 2 is additive because Phase 1 drew the seams in the right places (R1–R4, the Bus/Compiler/PKI/Registry interfaces).

Sources

[04_P1] HelixVPN-Phase1-MVP.md (control-plane architecture, DDL, RLS, proto, coordinator, events, policy, API) · [04_ARCH] HelixVPN-Architecture-Refined.md (control/data separation, ULA /48 + 4via6, no-logging, default-deny) · [04_P0] HelixVPN-Phase0-Spike.md (G6 file-delta reconcile ancestor) · [04_P2] HelixVPN-Phase2-Parity.md (NATS, HA, multi-region seams) · [05_YBO] (Go+Gin+PG+Redis+Podman mandate) · [02_QWN], [DSK], [GMI], [KMI], [MST] (D3/D4 divergence camps) · [SYNTHESIS] §§1–9 (stack floor, key decisions, ecosystem wiring, constitution bindings).