

helix-core Orchestrator & State Machine

helix-core Orchestrator & State Machine

Revision: 1 Last modified: 2026-06-25T00:00:00Z

Master technical specification — Volume 2 (Data Plane), nano-detail deep-dive. This document **deepens** the *Orchestrator*, *status stream*, and *auto transport selection* section of doc 01 [01-DP §5.1, §5.2] into an implementation-ready specification of the `crates/helix-core` orchestrator: its three async loops, the `tokio::broadcast` status stream, network-map reconciliation, the reconnection backoff state machine, the kill-switch + DNS coupling, and the `--mode=connector` posture that shares the same core. SPEC-ONLY: it describes **what to build**, not the shipping product. Sources cited inline by id — [01-DP §N] = doc 01 final/01-data-plane.md; [04_ARCH §N] = 04_VPN_CLD/HelixVPN-Architecture-Refined.md; [04_P0 §N] = 04_VPN_CLD/HelixVPN-Phase0-Spike.md; [04_P2 §N] = Phase-2 refined doc; [SYNTHESIS §N] = the cross-document synthesis; [01-LAD §N] = v02-data-plane/transport-selection-ladder.md; [research-masque] = the cited research digests. Any claim not grounded in the evidence base is marked UNVERIFIED per constitution §11.4.6.

0. Position, ownership, and invariants

0.1 What this document owns

The orchestrator (`crates/helix-core`, file `helix-core/src/orchestrator.rs` [01-DP §12]) is the **conductor** of the data plane on every role. It owns the wiring that turns the four frozen Phase-0 contracts — the Transport trait [01-DP §3.1], `helix-wg`'s `WgPeer` [01-DP §4], the `RouteMap`

reconciler input [01-DP §6.2], and the TunnelStatus broadcast enum [01-DP §5.2] — into a running tunnel, and it re-converges them without a process restart as state changes [04_ARCH §4.4, 04_P0 §4.5/§10]. Concretely it owns:

1. The **three async loops** — `tun` → `wg` → (`daita`) → `transport`, `transport` → `wg` → `tun`, `timer tick` — and the `tokio::select!` driver that multiplexes them (§2, §3).
2. The **`tokio::sync::broadcast status stream`** that surfaces `Connecting` | `Handshaking` | `Connected{transport,rtt}` | `Reconnecting` | `Down{reason}` to the FFI, the ladder, and aggregate metrics (§4) [01-DP §5.2, 04_P0 §4.5].
3. The **orchestrator state machine** that decides which TunnelStatus is current and gates every side effect (kill-switch open/close, DNS apply/revert, ladder kick) (§5).
4. **Network-map reconciliation** — `diff desired-vs-actual RouteMap`, converge peers / routes / transport policy / DNS, no restart, < 1 s convergence (§6) [04_ARCH §4.4].
5. The **reconnection backoff state machine** — jittered exponential re-dial with roam detection, feeding the ladder (§7) [04_P2 §1.4, 04_P0 §8].
6. The **kill-switch + DNS coupling** — OS firewall + tunnel-DNS side effects driven by the core state machine, never hand-edited (§8) [04_ARCH §5.6 table, §7].
7. The **connector posture** (`--mode=connector`) — the same core with a forwarding/NAT routing posture instead of default-route capture (§9) [04_P0 §4.5].

0.2 What this document does NOT own


- The Transport trait, TransportConfig, `dial()`, per-transport wire framing — frozen in doc 01 [01-DP §3]; the orchestrator consumes opaque `Box<dyn Transport>` handles.
- The **ladder selection algorithm** (which rung, per-network memory, regional prior, manual pin) — owned by `helix-core/src/ladder.rs` and specified in [01-LAD]. This document owns the *driver* that asks the ladder for the next transport and reacts to its outcome; it treats `ladder::select()` / `ladder::escalate()` as upstream calls (§7.4).
- WireGuard crypto (`helix-wg` Tunn state machine) [01-DP §4]; routing/ACL compilation (`helix-route`) [01-DP §6–§7]; DAITA shaping (`helix-daita`) [01-DP §9] — the orchestrator drives their hooks but does not define them.
- The WatchNetworkMap protobuf that *delivers* the RouteMap (doc 03); the orchestrator consumes the map's *shape* and treats a file-watch on `map.json` as the Phase-0 stand-in for the stream [04_P0 §10].
- The FFI bridge that re-emits TunnelStatus to Dart (doc 05); this document owns the broadcast *sender*, not the FFI *receiver*.

0.3 Invariants this document inherits and tightens

#	Invariant	Source	Orchestrator-specific tightening
I1	Transport never sees plaintext.	[01-DP I1]	Loop A hands WG-encrypted bytes to <code>Transport::send</code>; loop B hands <code>Transport::recv</code> bytes to <code>WgPeer</code>. The orchestrator never logs nor inspects either side's payload (§3, §10).
I2	Unreliable datagrams, never an ordered stream.	[01-DP I2]	The loops carry one datagram per iteration; a dropped <code>recv</code> is not retried at the orchestrator layer — WG owns loss semantics.
I3	Control plane never in the packet path; fail-static.	[01-DP I3, 04_ARCH §2.1]	If the reconcile source (stream / file-watch) dies, the existing loops keep forwarding; reconciliation is a side task, never on the hot path (§6.5).
I5	No-logging by construction; aggregate counters only.	[01-DP I5]	The status stream carries ephemeral state, never durable per-connection records; the only persisted artefact is the ladder's per-network memory hash (§4.5, §10) [01-LAD §6.4].
O-I10	Single state authority — exactly one <code>OrchState</code> is current; every <code>TunnelStatus</code> emission and every kill-switch / DNS side effect is a function of an <code>OrchState</code> transition, never an ad-hoc emit.	derived from [01-DP §5.2]	§5; prevents UI flicker and firewall races.
O-I11	Kill-switch fails closed — on any transition <i>into</i> Reconnecting/Down, the firewall stays/goes closed until a Connected	[04_ARCH §5.6 "no leaks if tunnel drops", §7]	§8.2, §8.5.

#	Invariant	Source	Orchestrator-specific tightening
	transition re-opens it; a crash leaves the closed rules installed.		
O-112	Reconcile is declarative + restart-free — a map delta mutates live peer/route/policy state in place; the process is never restarted to apply a change.	[04_ARCH §4.4, 04_P0 §10]	§6.
O-113	Connector and client run the identical core — only the routing posture differs, selected by one flag; no forked loop bodies.	[04_P0 §4.5, 01-DP I4]	§9.

1. The orchestrator at a glance



Three data-moving loops, one control state machine, two side-effect sinks (FFI + metrics), two control inputs (ladder outcomes, reconcile deltas), and one OS-effect sink (kill-switch/DNS). Everything below is the nano-detail of these eight boxes.

2. Top-level types & lifecycle

2.1 Orchestrator struct and construction

```
// helix-core/src/orchestrator.rs
use std::sync::Arc;
use std::net::IpAddr;
use tokio::sync::{broadcast, mpsc, watch, Mutex};
use tokio_util::sync::CancellationToken;
use bytes::Bytes;

use helix_transport::{Transport, TransportConfig, TransportError,
    dial};
use helix_wg::{WgPeer, WgVerdict};
use helix_tun::TunDevice;
use helix_route::{RouteMap, CompiledPolicy};

/// Run-role selected once at startup; the rest of the core is
/// identical (0-I13).
#[derive(Clone, Copy, Debug, PartialEq, Eq)]
pub enum Mode { Client, Connector }

/// All construction inputs. Resolved from CLI flags + the first
/// RouteMap snapshot.
pub struct OrchestratorConfig {
    pub mode: Mode,
    pub tun: TunSpec, // client: create/inject
    // fd; connector: forward iface + NAT
    pub initial_map: RouteMap, // first snapshot (doc
    // 03) or static map.json (Phase 0)
    pub transport_policy: TransportPolicy, // ladder order/pin/
    // budget [01-LAD $3]
    pub kill_switch: KillSwitchConfig, // $8
    pub dns: DnsConfig, // $8.3
    pub tunables: Tunables, // $11
}

/// The live orchestrator. Owns the broadcast sender; loops hold
/// `Arc<OrchShared>`.
pub struct Orchestrator {
    shared: Arc<OrchShared>,
    cancel: CancellationToken, // cooperative shutdown
    // of all loops ($2.4)
    join: Vec<tokio::task::JoinHandle<()>>,
}

/// Shared state every loop reads; mutated only via the state
/// machine (0-I10) and reconciler.
pub struct OrchShared {
    pub mode: Mode,
    pub status_tx: broadcast::Sender<TunnelStatus>, // $4
    pub state: Mutex<StateCell>, //
    // OrchState authority ($5)
}
```

```

pub active_tx: Mutex<Option<Arc<dyn Transport>>>, //
    current rung; swapped by re-dial (§7)
pub wg:
    Mutex<WgRegistry>, // peers
    keyed by pubkey (§6.3)
pub tun: TunDevice,
pub daita: Mutex<Option<helix_daita::Daita>>, // None
    unless opted in [01-DP §9]
pub policy: Mutex<CompiledPolicy>, //
    connector/edge verdict map (§9.3)
pub counters: Counters, //
    aggregate only (I5)
}

impl Orchestrator {
    /// Build, install the kill-switch in its *armed-closed-until-
    /// connected* posture (§8.2),
    /// spawn the three loops + reconciler + ladder driver, and
    /// emit `Connecting`.
    pub async fn start(cfg: OrchestratorConfig) -> Result<Self,
        OrchestratorError> { /* ... */ }

    /// Subscribe to the live status stream (FFI, metrics, tests
    /// all use this) [01-DP §5.2].
    pub fn subscribe(&self) -> broadcast::Receiver<TunnelStatus> {
        self.shared.status_tx.subscribe()
    }

    /// Push a new desired-state map (Phase 1: from the
    /// WatchNetworkMap stream; Phase 0:
    /// from the file-watch). Triggers reconciliation (§6).
    pub async fn apply_map(&self, map: RouteMap) -> Result<(),
        OrchestratorError> { /* ... */ }


    /// User action: pin a transport (manual override) – re-arms
    /// the ladder (§7.4) [01-LAD §8].
    pub async fn set_pin(&self, pin: Option<TransportKind>)
        { /* ... */ }

    /// Graceful stop: cancel loops, FIN the transport, revert
    /// kill-switch+DNS, emit Down (§2.4).
    pub async fn stop(self) -> Result<(), OrchestratorError> { /*
        ... */ }
}

```

active_tx is Arc<dyn Transport> (not Box) so loop A and loop B share **one** handle and a re-dial can atomically swap it under the mutex without tearing either loop (§7.3). The Transport trait's send/recv take &self [01-DP §3.1], so Arc<dyn Transport> is the correct shared ownership.

2.2 Spawn topology



Six tasks. The state-machine driver is the single writer of `OrchState`; every other task *requests* a transition by sending an `OrchEvent` on an mpsc channel the driver owns (§5.3), so O-I10 holds without lock contention on the hot loops.

2.3 Internal control channels

```

/// Loops and the reconciler never mutate OrchState directly –
they post events.
pub enum OrchEvent {
    HandshakeStarted, // loop B saw first
    inbound after dial
    HandshakeComplete { rtt_ms: u32 }, // WG session
    established (§5.4)
    TransportError(TransportError), // loop A/B carrier
    failed
    PeerTimeout { pubkey: [u8; 32] }, // tick: no recv
    within keepalive window
    MapApplied { changed: MapDelta }, // reconciler
    converged (§6)
    RoamDetected { new_local: IpAddr }, // local iface
    changed (§7.5)
    PinChanged(Option<TransportKind>), // user manual
    override
    Stop, // graceful shutdown
}

pub struct StateCell {
    pub state: OrchState, // §5.1

```

```

    pub since: tokio::time::Instant,           // dwell time for
        budgets/metrics
    pub backoff: Backoff,                       // §7.2
    pub event_tx: mpsc::Sender<OrchEvent>,      // loops clone this
}

```

2.4 Cancellation & shutdown ordering

stop() cancels the CancellationToken. Every loop's tokio::select! includes a `_ = self.cancel.cancelled() => break` arm (§3.4), so all three drain on the next iteration. Shutdown order is **strict** to honour O-I11 (no leak window during teardown):

1. Transition OrchState → ShuttingDown (suppresses re-dial; ladder driver exits).
 2. `active_tx.close()` (FIN the QUIC/TCP carrier idempotently) [01-DP §3.1].
 3. **Keep** the kill-switch closed; revert DNS to system resolvers (§8.4).
 4. Cancel loops A/B/C/reconciler; join all handles with a bounded timeout (Tunables.shutdown_grace, default 2 s).
 5. On user-initiated stop only: revert the kill-switch to *off* (open) — the user asked to disconnect, so traffic flows on the underlay again (§8.6). On crash/abnormal exit the closed rules are left installed (O-I11).
 6. Emit Down { reason: "stopped" } last.
-

3. The three loops (nano-detail)

The orchestrator runs three loops [04_P0 §4.5, 01-DP §5.1]. Each is a single tokio::task with an internal loop { tokio::select! { ... } }. The bodies are cancel-safe: Transport::recv and TunDevice::recv are documented cancel-safe [01-DP §3.1], and WgPeer calls are synchronous CPU work between awaits.

3.1 Loop A — outbound: tun → wg → (daita) → transport

```

async fn loop_a(shared: Arc<OrchShared>, cancel:
    CancellationToken) {
    let mut scratch = vec![0u8; 65535];           // reused
        encrypt buffer; no per-pkt alloc
    loop {
        tokio::select! {
            biased;                               // cancel
            checked first
            _ = cancel.cancelled() => break,
            pkt = shared.tun.recv() => {           // plaintext
                IP packet from the OS (client)
            }
        }
    }
}

```



```

        let ip_pkt = match pkt { Ok(p) => p, Err(_) =>
continue };
        // 1. WG encrypt (helix-wg owns the Tunn state
machine) [01-DP §4]
        let mut wg = shared.wg.lock().await;
        let verdict =
wg.peer_for_dst(&ip_pkt)                // route by
AllowedIPs (I6)
                .map(|p|
p.handle_tun_out(&ip_pkt, &mut scratch));
        let enc: &[u8] = match verdict {
            Some(WgVerdict::WriteToTransport(b)) => b,
            Some(WgVerdict::Nothing) | None =>
continue, // no session yet / no route
            Some(WgVerdict::WriteToTun(_)) =>
continue, // impossible here
            Some(WgVerdict::Err(_)) => {
shared.counters.wg_enc_err.inc(); continue }
        };
        let mut dg = Bytes::copy_from_slice(enc);

drop(wg); //
release before await
        // 2. DAITA shaping (optional, above WG / below
transport) [01-DP §9]
        if let Some(d) =
shared.daita.lock().await.as_mut() {
            for act in d.on_packet(&mut dg,
tokio::time::Instant::now()) {
                if let ScheduledAction::SendNow(cover) =
act { let _ = send_dg(&shared, cover).await; }
            }
        }
        // 3. hand to the active transport (I1: opaque
encrypted bytes)
        if let Err(e) = send_dg(&shared, dg).await {
post_err(&shared, e).await; }
    }
}

}

/// Single send path; swaps cleanly if a re-dial replaced
active_tx mid-flight (§7.3).
async fn send_dg(shared: &OrchShared, dg: Bytes) -> Result<(),
TransportError> {
    let tx =
shared.active_tx.lock().await.clone() //
Arc clone; cheap
        .ok_or(TransportError::Closed)?;
    match tx.send(dg).await {
        Err(TransportError::Oversize(n)) => {
shared.counters.oversize.inc();

        Err(TransportError::Oversize(n)) } //
        §6.6 MTU edge → caller lowers MTU
    }
}

```

```

        other => other,
    }
}

```

Key points: (1) the wg mutex is **released before** any `.await` so loop B can run while loop A is in DAITA/transport — no cross-loop deadlock (§12 cross-feature analysis); (2) a WG-encrypt that yields `Nothing` (handshake not yet complete) silently drops the user packet — WG will buffer the handshake-trigger internally and the app's TCP/QUIC will retransmit; (3) `Oversize` is propagated, never silently truncated (I2) [01-DP §10].

3.2 Loop B — inbound: transport → wg → tun

```

async fn loop_b(shared: Arc<OrchShared>, cancel:
    CancellationToken) {
    let mut scratch = vec![0u8; 65535];
    loop {
        let tx = match shared.active_tx.lock().await.clone() {
            Some(t) => t, None => {
                tokio::time::sleep(IDLE_BACKOFF).await; continue }
        };
        tokio::select! {
            biased;
            _ = cancel.cancelled() => break,
            r = tx.recv() => {
                let dg = match r {
                    Ok(b) => b,
                    Err(e) => { post_err(&shared,
                        e).await;           // §5/§7 drive re-dial
                }
                tokio::time::sleep(IDLE_BACKOFF).await; continue }
            };
            let mut wg = shared.wg.lock().await;
            match wg.handle_transport_in(&dg, &mut scratch) {
                WgVerdict::WriteToTun(plain) =>
                {
                    // decrypted IP packet
                    // connector: forward into served CIDR +
                    NAT (§9.2); client: write TUN
                    match shared.mode {
                        Mode::Client => { let _ =
                            shared.tun.send(plain).await; }
                        Mode::Connector => {
                            forward_into_lan(&shared, plain).await; }
                    }
                }
                WgVerdict::WriteToTransport(reply) =>
                {
                    // handshake reply → send back
                    let b = Bytes::copy_from_slice(reply);
                    drop(wg);
                    let _ = tx.send(b).await;
                    post_event(&shared,
                        OrchEvent::HandshakeStarted).await;
                }
            }
        }
    }
}

```

```

                                WgVerdict::Nothing => { /* keepalive / dupe
*/ }
                                WgVerdict::Err(_) =>
shared.counters.wg_dec_err.inc(),
                                }
                                }
                                }
                                }
}

```

Loop B is where **handshake progress** is observed: the first inbound datagram after a dial posts `HandshakeStarted` ($\rightarrow$ `Handshaking`), and `helix-wg` reporting an established session (via `WgPeer::session_established()` polled in loop C, §3.3) posts `HandshakeComplete{rtt}` ($\rightarrow$ `Connected`). The transport carrier connecting is **not** "connected" — the WG handshake is (I2 tightening, mirrors [01-LAD I2]).

3.3 Loop C — timer tick: tick $\rightarrow$ keepalive / rekey / liveness

```

async fn loop_c(shared: Arc<OrchShared>, cancel:
    CancellationToken) {
    let mut scratch = vec![0u8; 4096];
    let mut iv =
        tokio::time::interval(shared_tick_period()); // 100–
        250 ms [01-DP §4]
    iv.set_missed_tick_behavior(tokio::time::MissedTickBehavior::Skip); //
        no burst after stall
    loop {
        tokio::select! {
            biased;
            _ = cancel.cancelled() => break,
            _ = iv.tick() => {
                let mut wg = shared.wg.lock().await;
                for peer in wg.iter_mut() {
                    match peer.tick(&mut scratch)
                    {
                        // emits keepalive/rekey [01-DP §4]
                        WgVerdict::WriteToTransport(b) => {
                            let dg = Bytes::copy_from_slice(b);
                            drop_then_send(&shared, dg);
                        }
                        _ => {}
                    }
                }
                // liveness: WG passive-keepalive is 10 s;
                declare a peer stale if no
                // authenticated recv within keepalive_timeout
                (§7.1, default 25 s – WG's
                // REKEY-after-no-recv heuristic). UNVERIFIED
                exact value  $\rightarrow$  see §11 Tunables.
                if peer.last_recv_age() >
                shared.tunables.keepalive_timeout {
                    post_event(&shared,
                    OrchEvent::PeerTimeout { pubkey: peer.pubkey() }).await;

```

```

    }
    if peer.session_just_established() {
        post_event(&shared,
OrchEvent::HandshakeComplete { rtt_ms:
peer.rtt_ms() }).await;
    }
}
}
}
}
}
}
}
}

```

The 100–250 ms cadence is the [01-DP §4] guidance;

MissedTickBehavior::Skip prevents a tick storm after the runtime was starved (host-safety adjacent to §12.6). PeerTimeout is the trigger that moves a healthy Connected tunnel into Reconnecting (§5.4, §7.1).

3.4 Why three loops and not one select!

Splitting into three tasks (rather than one mega-select!) gives: independent backpressure (a slow TUN write cannot stall inbound decrypt); per-loop cancel-safety reasoning; and the ability for loop B to sleep(IDLE_BACKOFF) while active_tx is None during a re-dial without blocking loop A's drain. The cost is the shared mutexes; they are held only across synchronous helix-wg calls (microseconds), never across .await, so contention is bounded (§13 performance budget). IDLE_BACKOFF is a short fixed 10 ms poll while disconnected — **not** the reconnection backoff of §7 (that governs *re-dial attempts*, this governs the busy-wait while a re-dial is in flight).

4. The status stream (tokio::broadcast)

4.1 The enum (frozen contract, mirrored across the FFI)

```

// helix-core/src/status.rs (mirrored across the FFI boundary,
// doc 05) [01-DP §5.2, 04_P0 §4.5]
#[derive(Clone, Debug, PartialEq, Eq)]
pub enum TunnelStatus {
    Connecting, // dialling the
                first/seeded rung
    Handshaking, // carrier up,
                WG handshake in flight
    Connected { transport: String, rtt_ms: u32 }, // e.g.
                ("masque-h3", 23)
    Reconnecting, // a working
                tunnel dropped; re-dialling
}

```

```

    Down { reason: String },                                     // terminal for
        this attempt; see §4.4 reasons
}

```

Connected.transport is the Transport::kind() string [01-DP §3.1] ("plain-udp", "masque-h3", ...) so the UI shows *which* rung won [01-LAD §5.4]. rtt_ms is the WG-handshake RTT EWMA from TransportHealth.rtt_ewma_ms [01-DP §3.1], refreshed on each Connected re-emit (§4.3).

4.2 Channel construction & capacity

```

// capacity sized for: FFI receiver + metrics receiver + ladder/
// test receivers, with slack.
const STATUS_CHANNEL_CAP: usize = 64;
let (status_tx, _rx0) =
    broadcast::channel::<TunnelStatus>(STATUS_CHANNEL_CAP);

```

tokio::sync::broadcast is a **bounded ring** [01-DP §5.2]: each receiver has an independent read cursor; a slow receiver that falls STATUS_CHANNEL_CAP behind gets a RecvError::Lagged(n) on its next recv() and is fast-forwarded to the oldest retained value. This is acceptable for a *status* stream (the receiver only ever needs the **latest** state), but the FFI/UI receiver **MUST** treat Lagged as "re-read latest", never as an error (§4.6 receiver contract). Status is **edge-coalesced** (§4.3): emissions are rare (state changes), so 64 slots is generous; lag only happens if a receiver is wedged.

4.3 Emit discipline — coalesce, never spam

Rule (O-I10): the **only** code that calls status_tx.send(...) is the state-machine driver, exactly once per OrchestrationState transition, **after** the side effects of that transition have been applied (kill-switch/DNS first, then emit — so a UI that reacts to Connected can trust the firewall is already open). Re-emitting the *same* logical state is suppressed except for Connected, which is re-emitted when rtt_ms changes by more than Tunables.rtt_emit_delta_ms (default 5 ms) so the UI's latency chip updates without flicker. Reconnecting is emitted once per drop, not once per failed re-dial (§7.3).

4.4 Down.reason closed-ish vocabulary

Down.reason is a human/machine string; the orchestrator uses a **stable prefix** so tests and metrics can classify without parsing prose (§11.4.6 no-guessing):

reason value	Meaning	Emitted from
"stopped"	user/graceful stop	§2.4
"pinned-transport-failed"	manual pin rung failed; no auto-escalate	[01-LAD §8], §5.4
"ladder-exhausted"	every rung in order failed its budget	§7.4
"auth-failed"	WG handshake rejected (revoked/bad key)	§5.4, §8.5
"no-route"	RouteMap has no peer for the requested destination	§6.4
"host-fatal"	TUN open / firewall-apply failed at startup	§8.2

Reconnecting (not Down) is the state while *any* rung remains untried within budget; Down is reached only when the attempt is genuinely terminal (O-I10, §5.4).

4.5 What the stream MUST NOT carry (I5)

No per-packet, per-flow, per-peer-endpoint, or timestamped-per-connection record ever crosses the broadcast. `Connected.transport` is a *kind* label, not an endpoint; `rtt_ms` is an aggregate EWMA. Durable telemetry derived from the stream is **aggregate only**: "transport X succeeded after N escalations in region R" [01-DP §5.3 step 6, 01-LAD §9] — never the SSID, never the gateway IP. This is the data-plane half of "no-logging as code" [SYNTHESIS §7].

4.6 Receiver contract (FFI / metrics / tests)

```
// Reference consumer loop (doc 05 FFI + the test harness both
// follow this).
let mut rx = orch.subscribe();
loop {
  match rx.recv().await {
    Ok(s) =>
      ui_apply(s), // render latest
    Err(broadcast::error::RecvError::Lagged(_)) =>
      continue, // skip-to-latest, NOT an error
    Err(broadcast::error::RecvError::Closed) =>
      break, // orchestrator gone
  }
}
```


5. The orchestrator state machine

5.1 States

```
#[derive(Clone, Debug, PartialEq, Eq)]
pub enum OrchState {
    Idle, // constructed, not started (pre-
        start only)
    Connecting, // dialling a rung (carrier
        handshake in flight)
    Handshaking, // carrier up; WG Noise IK
        handshake in flight
    Connected { kind: TransportKind, rtt_ms: u32 },
    Reconnecting { since_drop: tokio::time::Instant }, // re-
        dialling within budget
    ShuttingDown, // graceful teardown in progress
        (§2.4)
    Down { reason: DownReason }, // terminal for this attempt
}
```

OrchState is the internal authority; TunnelStatus (§4.1) is its FFI-facing projection (Idle/ShuttingDown have no public projection — they map to Connecting/Down respectively at the FFI). The split keeps internal bookkeeping (e.g. `since_drop`) off the wire.

5.2 State diagram



5.3 Driver task (single writer of OrchState)

```
async fn state_driver(shared: Arc<OrchShared>, mut ev:
    mpsc::Receiver<OrchEvent>,
    cancel: CancellationToken) {
    while let Some(event) = tokio::select! {
        _ = cancel.cancelled() => None,
        e = ev.recv() => e,
    } {
        let mut cell = shared.state.lock().await;
        let next = transition(&cell.state, &event,
            &shared).await; // pure-ish decision (§5.4)
        if let Some(ns) = next {
            apply_side_effects(&cell.state, &ns,
                &shared).await; // kill-switch/DNS FIRST (§4.3)
            cell.since = tokio::time::Instant::now();
            cell.state = ns.clone();
            emit_status(&shared,
                &ns).await; // THEN emit (O-I10)
        }
    }
}
```

transition() is the table of §5.2 encoded as a match (&state, &event); it returns Option<OrchState> (None = ignore the event in this state, e.g. a stray HandshakeComplete while already Connected). apply_side_effects() is where O-I11 (kill-switch) and §8 (DNS) live, ordered before the emit so the FFI never observes Connected before the firewall opened.

5.4 Decision details (the non-obvious transitions)

- **Handshaking → Down{auth-failed} is terminal-without-retry.** A WG handshake that is *cryptographically rejected* (peer revoked, key rotated, device.revoked propagated [04_ARCH §4.4]) will fail identically on every rung — escalating the ladder is pointless and would mask a real revocation. helix-wg distinguishes "no reply" (network/DPI → retry/escalate) from "handshake rejected" (auth → Down). UNVERIFIED: boringtun's public API for surfacing an *explicit* reject vs a silent timeout — if unavailable, the orchestrator treats a handshake that never completes within handshake_timeout as the *escalate* class and can only infer auth-failure from the control plane (doc 02) signalling a revoke; recorded as a follow-up forensic item per §11.4.6.
- **Connected → Reconnecting on PeerTimeout.** A live tunnel that goes silent (carrier still "up" but WG sees no authenticated rcv past keepalive_timeout) is the roam / NAT-rebind / silent-drop case [04_P0 §8 "flap the client iface; recovery < 3 s"]. The kill-switch closes immediately (O-I11) so no traffic leaks during the gap.
- **Reconnecting budget.** Each re-dial consults the ladder (§7.4); Reconnecting → Down only when the ladder reports exhaustion *and* the

backoff cap is hit (§7.2) — never on a single failure (avoids premature Down).

6. Network-map reconciliation

6.1 The reconciler contract

Agents are **declarative reconcilers** [04_ARCH §4.4, 04_PO §10]: they diff a *desired* RouteMap (pushed by the coordinator via WatchNetworkMap, doc 03; a file-watch on map.json in Phase 0) against *actual* live state and converge — bring peers up/down, switch transport policy, update routes, update DNS — **without restarting** (O-I12). Convergence target: a change reflected on all affected state in < 1 s [04_ARCH §4.4].

6.2 RouteMap (the desired state, frozen shape) [01-DP §6.2]

```
pub struct RouteMap {
    pub self_overlay: IpAddr,           // this node's
        overlay IP (ULA /48, D4)
    pub peers: Vec<PeerRoute>,         // already policy-
        filtered (need-to-know, I6)
    pub dns: Vec<IpAddr>,              // tunnel DNS
        servers (§8.3)
    pub transport_policy: TransportPolicy, // per-leg ladder
        order/pin/budget [01-LAD §7]
    pub daita:
        Option<DaitaMachines>,         // optional shaping config
        (data, not code) [01-DP §9]
}

pub struct PeerRoute {
    pub wg_pubkey: [u8; 32],
    pub endpoint_candidates:
        Vec<SocketAddr>, // §8 NAT traversal (Phase 2)
    pub allowed_ips:
        Vec<IpNet>,      // overlay prefixes this peer is
        next hop for
    pub via_connector: Option<SiteId>, // 4via6 site
        disambiguation [01-DP §6.1]
}
```

6.3 Diff algorithm

```
/// Pure diff: what changed between actual (live WgRegistry/
    policy/dns) and desired (new map).
pub struct MapDelta {
    pub peers_added: Vec<PeerRoute>,
    pub peers_removed: Vec<[u8; 32]>, // pubkeys to tear
        down (revocation path)
```


```

pub peers_changed: Vec<PeerRoute>,           // allowed_ips /
    endpoint changed
pub dns_changed: bool,
pub policy_changed: bool,
pub transport_policy_changed: bool,
}

fn diff(actual: &WgRegistry, dns: &[IpAddr], pol:
    &CompiledPolicy, desired: &RouteMap) -> MapDelta {
    // keyed by wg_pubkey; O(n) set difference on a sorted pubkey
    // index
    /* ... */
}

```

6.4 Apply (converge) — restart-free



A `peers_removed` entry whose pubkey is the **active** peer (the gateway you are tunnelled through) is the revocation-of-self case: tear the session, close the kill-switch, emit `Down{reason:"auth-failed"}` (§5.4). A `peers_changed` that only touches `allowed_ips` updates the routing map **without** resetting the WG session (the key is unchanged) — this is the "policy edit < 1 s, no restart" MVP gate (AC5) [04_P1 §11.2].

6.5 Reconcile is a side task (I3 fail-static)

The reconciler runs as its own `tokio::task`; if its **source** dies (stream disconnect, file-watch error), the three loops keep forwarding on the last-applied map (fail-static [04_ARCH §2.1]). The reconciler reconnects to the stream with its own backoff (independent of §7's data-plane backoff) and re-applies a fresh snapshot on reconnect. A reconcile failure never tears the tunnel.

6.6 MTU coupling on transport change

When `transport_policy_changed` causes the ladder to later swap the active transport (§7.3), the inner WG MTU is recomputed: `inner_mtu = min(active.effective_mtu(), path_mtu)` [01-DP §10]. The orchestrator

pushes the new MTU into helix-tun and into each WgPeer's fragmentation threshold. An in-flight Oversize error (§3.1) during the swap window triggers an immediate MTU lower rather than a silent drop (I2).

7. Reconnection backoff state machine

7.1 Triggers into reconnect

A reconnect begins on any of: TransportError from loop A/B (carrier died), PeerTimeout from loop C (silent WG drop), or RoamDetected (local iface changed — §7.5). All three post events that move Connected → Reconnecting (§5.4). Target recovery time is < 3 s for an interface flap / roam [04_P0 §8].


7.2 The backoff struct (jittered exponential, capped)

```
pub struct Backoff {
    base: Duration,           // first re-dial delay
    cap: Duration,            // ceiling
    factor: f64,               // multiplier per failed attempt
    jitter: f64,               // 0.0..=1.0 fraction of full jitter
                             (decorrelated)
    attempt: u32,              // resets to 0 on Connected
    rng: SmallRng,
}

impl Backoff {
    /// Decorrelated jitter (AWS-style): next ∈ [base, min(cap,
    /// prev*factor)] then jitter.
    fn next_delay(&mut self) -> Duration { /* ... */ }
    fn reset(&mut self) { self.attempt = 0; }
}
```

Recommended Tunables (derived, calibrated in Phase 0 against the < 3 s roam gate — values are the *spec recommendation*, not measured; §11.4.6): base = 250 ms, cap = 30 s, factor = 2.0, jitter = 1.0 (full decorrelated jitter to avoid thundering-herd re-dials when a gateway restarts and many clients reconnect together). The **first** re-dial after a drop is near-immediate (base) so a transient blip recovers well under 3 s; only a persistently dead network climbs toward the 30 s cap.

7.3 Re-dial driver + atomic transport swap



The swap is **atomic** under active_tx's mutex: loop A's send_dg (§3.1) and loop B's per-iteration active_tx.lock().await.clone() (§3.2) always see either the old or the new Arc, never a torn handle. While active_tx == None, loop A's sends return Closed (the user's app retransmits) and loop B idles on IDLE_BACKOFF.

7.4 Ladder coupling (this doc drives, [01-LAD] decides)

The re-dial driver calls ladder::next_rung() / ladder::escalate() [01-LAD §5]; the ladder owns *which* rung (per-network memory, regional prior, pin). The orchestrator owns *when* to re-dial (backoff timing) and *whether* the attempt is terminal. ladder-exhausted (every rung in order failed its budget) + backoff cap reached ⇒ Down{ladder-exhausted} (§5.4). A pinned transport that fails ⇒ Down{pinned-transport-failed} with **no** escalation (L-I7 [01-LAD §8]).

7.5 Roam detection (fast-path)

```
// A platform iface watcher (netlink on Linux; NWPathMonitor/  
    ConnectivityManager on mobile via  
// the shim, doc 06) posts RoamDetected{new_local} when the  
    device's primary address changes.
```

On RoamDetected the orchestrator **interrupts** any in-flight backoff sleep and re-dials *immediately* at the remembered-good rung for the *new* network fingerprint [01-LAD §6], because WireGuard's roaming will re-latch onto the new source address as soon as a handshake completes [01-DP §8.2]. This is the difference between a 30 s backoff-climb and the < 3 s roam recovery gate. RoamDetected resets Backoff.attempt = 0.

8. Kill-switch + DNS coupling

8.1 Principle

The kill-switch and DNS-leak protection are **owned by the core state machine**, not hand-edited firewall rules [04_ARCH §5.6 carry-forward note, §7]. Every open/close and every DNS apply/revert is a side effect of an OrchState transition (O-I10, §5.3), so there is no window where the UI shows "connected" but the firewall is open, or "disconnected" while DNS still points at the tunnel.

8.2 Kill-switch posture & arming

```
pub struct KillSwitchConfig {  
    pub mode: KillSwitchMode,           // Strict |  
        Permissive | Off  
    pub allow_lan: bool,                 // permit RFC1918  
        LAN even when closed (printers etc.)  
    pub allow_gateway_endpoints: Vec<IpNet>, // the gateway IPs  
        the transport must reach (MUST stay open)  
}  
pub enum KillSwitchMode {  
    Strict,           // closed unless Connected; arm-closed at start  
        (no leak before first connect)  
  
        Permissive, // closed only after a first successful  
        Connected (then on every drop)  
    Off,             // never installs firewall rules  
}
```

On Idle → Connecting (start()), Strict mode installs the **closed** ruleset *immediately* — default-drop all egress except (a) the allow_gateway_endpoints (so the transport can reach :443), (b) DHCP/ND

for link bring-up, (c) optionally `allow_lan`. This is "no leaks before the first handshake even completes". Permissive waits until the first Connected before it will ever close (so a never-connecting client on a captive-portal network isn't bricked). A firewall-apply failure at start $\Rightarrow$ `Down{host-fatal}` (§4.4) — never proceed with a believed-but-unverified kill-switch (§11.4.6).

8.3 DNS coupling

```
pub struct DnsConfig {
  pub servers: Vec<IpAddr>,      // tunnel DNS (from
    RouteMap.dns, §6.2)
  pub block_plaintext_53:
    bool, // drop off-tunnel :53 (UDP+TCP) when Connected
    (leak guard)
  pub apply:
    DnsApplyMethod, // resolvconf | systemd-resolved | /
    etc/resolv.conf | shim (doc 06)
}
```

On Handshaking $\rightarrow$ Connected: set tunnel DNS (`RouteMap.dns`) via the platform method, and (if `block_plaintext_53`) install firewall rules dropping any `:53` egress that is **not** to the tunnel DNS over the tunnel — closing the classic DNS-leak [04_ARCH §5.6 "force tunnel DNS; blocks plaintext :53 off-tunnel"]. On Connected $\rightarrow$ Reconnecting/Down: the DNS-leak block **stays** (closed-fails-safe) until either reconnect re-applies tunnel DNS or a user-stop reverts to system resolvers (§8.6).

8.4 Coupling table (the side-effect matrix)

Transition	Kill-switch	DNS	Source
Idle $\rightarrow$ Connecting (Strict)	install closed (allow gw endpoints only)	leave system DNS	§8.2
Handshaking $\rightarrow$ Connected	open (allow all egress via tunnel)	apply tunnel DNS + block off-tunnel :53	[04_ARCH §5.6]
Connected $\rightarrow$ Reconnecting	close (O-I11)	keep tunnel- DNS block (fail-closed)	§8.1
* $\rightarrow$ Down{reason $\neq$ stopped}	stay closed	keep block	O-I11
ShuttingDown $\rightarrow$ Down{stopped} (user)	revert to off (open)	revert to system DNS	§8.6
crash / panic		block remains	O-I11

Transition	Kill-switch	DNS	Source
	closed rules remain installed		


8.5 auth-failed is closed-and-terminal

A revoked device (auth-failed, §5.4) reaches Down with the kill-switch **closed** and DNS **blocked** — a revoked client cannot leak traffic onto the underlay by virtue of being revoked. Recovery requires explicit user action (re-enroll) which goes through start () again.

8.6 The only open path

The kill-switch is reverted to *open* on **exactly one** path: a user-initiated graceful stop (Stop → ShuttingDown → Down{stopped}, §2.4 step 5). Every other terminal — crash, ladder-exhausted, auth-failed, host-fatal — leaves it closed (O-I11). This is the literal "no leaks if the tunnel drops" guarantee [04_ARCH §5.6].

8.7 Kill-switch + DNS state sub-machine



9. Connector mode (`--mode=connector`)

9.1 Same core, different posture (O-I13)

The connector is the **same** Orchestrator with `Mode::Connector` [04_P0 §4.5, 01-DP §5.1]. The three loops, the status stream, reconciliation, backoff, and the transport layer are **byte-for-byte identical** to the client (I4). Only two things differ:

client: [app traffic] → TUN → WG encrypt → [DAITA] → Transport
→ gateway
connector: gateway → Transport → WG decrypt → forward into
served CIDR (and reverse)

- The connector does **not** capture the device default route; instead it (a) advertises its served CIDRs (e.g. `10.10.0.0/24`) in its RouteMap, and (b) NATs/forwards decapsulated packets into the LAN interface and back [04_P0 §4.5, 01-DP §5.1].
- The connector is **outbound-only** — it dials the gateway; it never listens for inbound connections (the founding no-port-forward constraint [SYNTHESIS §1]).

9.2 Forwarding & NAT (loop B branch)

loop_b's `WriteToTun` branch (§3.2) dispatches on `mode`: Client writes the TUN; Connector calls `forward_into_lan`:

```
async fn forward_into_lan(shared: &OrchShared, plain: Bytes) {  
    // 1. ACL gate: the decapsulated packet's (src,dst,proto,port)  
    //    must match the compiled  
    //    verdict map (default-deny, I6) [01-DP §7] BEFORE it  
    //    touches the LAN.  
    if shared.policy.lock().await.verdict(&plain) !=  
        Verdict::Allow {  
        shared.counters.policy_drop.inc(); return;  
    }  
    // 2. SNAT into the served CIDR and emit on the LAN iface  
    //    (Linux: the veth/eth handle).  
    shared.tun.forward_lan(plain).await; // helix-tun owns the  
    //    masquerade + reverse mapping  
}
```

The **reverse** path (LAN host → connector → gateway → client) enters via the LAN iface, is matched to the originating WG session, WG-encrypted (loop A's `peer_for_dst` selects the gateway peer), and shipped back over the transport. The connector therefore runs loop A driven by the **LAN iface** rather than a TUN; `TunSpec` abstracts both (`TunDevice::recv` yields the next inbound packet regardless of source) so loop A's body is unchanged (O-I13).

9.3 Verdict map on the connector

The connector installs the **compiled policy** (CompiledPolicy, [01-DP §7]) as an nftables/eBPF verdict map and re-installs it on a policy_changed reconcile delta (§6.4) in < 1 s, no restart [04_ARCH §4.4]. Default-deny, fail-closed: if the compiled policy is absent/stale, drop (I6). This is the "split-horizon / microsegmentation by default" property [01-DP §7].

9.4 Status semantics on the connector

The connector emits the **same** TunnelStatus stream (its tunnel to the gateway is what Connected{transport, rtt} describes). The Connector Flutter flavor (doc 05) subscribes to it exactly as the Access app does — the FFI/UI code is shared (O-I13). Down on a connector means "the connector lost its gateway tunnel"; the kill-switch posture is typically Off on a connector (it is a network appliance, not a privacy client) — KillSwitchMode is a config knob (§8.2, §11), not hardcoded.

10. Error taxonomy

```
// helix-core/src/error.rs
#[derive(thiserror::Error, Debug)]
pub enum OrchestratorError {
    #[error("tun open/inject failed: {0}")]
    TunOpen(String),           // → Down{host-fatal}
    #[error("kill-switch apply failed: {0}")]
    KillSwitch(String),        // → Down{host-fatal} (§8.2)
    #[error("dns apply failed: {0}")]
    Dns(String),               // → Down{host-fatal}
    #[error("transport: {0}")]
    Transport(#[from] TransportError), // → re-dial/escalate
    #[error("wg: {0}")]
    Wg(String),                // enc/dec error (counter, non-fatal)
    #[error("map reconcile: {0}")]
    Reconcile(String),         // side task; never tears tunnel (§6.5)
    #[error("ladder exhausted")]
    LadderExhausted,           // → Down{ladder-exhausted}
    #[error("pinned transport failed")]
    PinFailed,                 // → Down{pinned-transport-failed}
    #[error("shutdown timeout")]
    ShutdownTimeout,           // join exceeded grace (§2.4)
}
```

Classification (drives §5.4): TunOpen/KillSwitch/Dns are **host-fatal** at startup (the core refuses to run a tunnel it cannot leak-protect). Transport errors map to the re-dial/escalate path (§7). Wg enc/dec errors are per-

packet, counted, never fatal (a single malformed datagram does not tear the tunnel). Reconcile errors are side-task-only (I3). TransportError variants map to ladder behaviour exactly as [01-DP §3.1 / 01-LAD §5]: DialTimeout/HandshakeFailed/EndpointBlocked → escalate; Closed → re-dial same rung; Oversize → lower MTU (§6.6).

11. Config knobs (Tunables)

```
pub struct Tunables {
    // --- loops ---
    pub tick_period: Duration,           // loop C
        cadence;           default 200 ms [01-DP §4]
    pub idle_poll:
        Duration,           // loop B busy-wait when active_tx
        None; default 10 ms
    // --- handshake / liveness ---
    pub handshake_timeout: Duration,     //
        Handshaking→Reconnecting; default 5 s (derived)
    pub keepalive_timeout: Duration,     // PeerTimeout
        threshold; default 25 s (WG REKEY heuristic, UNVERIFIED
        exact)
    // --- backoff (§7.2) ---
    pub backoff_base: Duration,          // default 250 ms
    pub backoff_cap: Duration,           // default 30 s
    pub backoff_factor: f64,             // default 2.0
    pub backoff_jitter: f64,             // default 1.0 (full
        decorrelated)
    // --- status (§4.3) ---
    pub rtt_emit_delta_ms: u32,          // re-emit Connected if
        rtt drifts >; default 5 ms
    pub status_channel_cap: usize,       // broadcast ring;
        default 64
    // --- shutdown (§2.4) ---
    pub shutdown_grace:
        Duration,           // join timeout; default 2 s
    // --- reconcile (§6.5) ---
    pub reconcile_reconnect_base:
        Duration, // stream re-subscribe backoff; default 1 s
}
```

Every numeric default above tagged "derived" is the **spec recommendation**, to be calibrated against the Phase-0 gates (G1 throughput, < 3 s roam, < 1 s reconcile) with captured benchmark evidence before being frozen (§11.4.6 — a number is a guess until measured) [04_P0 §8].

keepalive_timeout = 25 s mirrors WireGuard's REKEY-after-no-recv window but the exact boringtun-surfaced value is UNVERIFIED and is a Phase-0 measurement item.

12. Edge cases (each maps to a test point, §14)

#	Edge case	Required behaviour
E1	TUN delivers a packet before any WG session exists	loop A <code>WgVerdict::Nothing</code> → silent drop; app retransmits; no crash (§3.1).
E2	<code>active_tx</code> swapped mid-send (re-dial races loop A)	atomic Arc swap; old or new handle, never torn (§7.3).
E3	broadcast receiver wedged > 64 events behind	Lagged returned; receiver skips-to-latest; orchestrator unaffected (§4.2/§4.6).
E4	Reconcile source dies while Connected	loops keep forwarding (fail-static I3); reconciler re-subscribes with backoff (§6.5).
E5	Revoke (peers_removed includes active gateway)	tear session, close kill-switch, <code>Down{auth-failed}</code> ; revoke effective < 1 s (§6.4).
E6	Roam (Wi-Fi→cellular) mid-Connected	<code>RoamDetected</code> interrupts backoff, immediate re-dial at remembered rung, < 3 s recovery (§7.5).
E7	All rungs fail their budget	<code>Reconnecting→Down{ladder-exhausted}</code> only after backoff cap; kill-switch stays closed (§7.4, O-I11).
E8	Pinned transport fails	<code>Down{pinned-transport-failed}</code> , NO escalation (§7.4, L-I7).
E9	Oversize on transport swap (smaller MTU)	lower inner WG MTU, do not truncate (§6.6, I2).
E10	Process panic in a loop	other loops + kill-switch unaffected (closed rules persist O-I11); supervisor restarts core (<code>start()</code> resets ladder, L-I8).
E11	Connector packet fails the verdict map	<code>forward_into_lan</code> drops + counts; never reaches LAN (§9.2, I6).
E12	tick after runtime starvation	<code>MissedTickBehavior::Skip</code> — one catch-up tick, no storm (§3.3, §12.6 host-safety).

13. Security & performance

13.1 Security considerations

- **No plaintext on the status path (I1/I5).** TunnelStatus carries only a transport *kind* label + aggregate RTT; never endpoints, SSIDs, packets, or per-connection timestamps (§4.5). A code-review gate (§11.4.142) must reject any `status_tx.send` payload that embeds an endpoint or per-flow datum.
- **Fail-closed everywhere (O-I11).** Crash, exhaustion, revoke, and host-fatal all leave the kill-switch closed and DNS blocked (§8.4); only a deliberate user-stop opens it (§8.6).
- **Default-deny on the connector (I6).** Decapsulated packets pass the verdict map *before* the LAN (§9.2); a stale/absent policy drops (fail-closed) [01-DP §7].
- **Single state authority (O-I10).** One writer of `OrchState` → no firewall race, no "connected-but-leaking" window (§5.3). This is the structural defence against the classic VPN kill-switch TOCTOU bug.
- **Revocation latency (SEC).** `device.revoked` → verdict-map drop / session tear in < 1 s, no restart (§6.4) [04_ARCH §4.4, SYNTHESIS §7].
- **Host-session safety (constitution §12).** Loop C uses `MissedTickBehavior::Skip` and the loops never spin-busy except the bounded `idle_poll`; no unbounded memory growth (the broadcast ring is fixed, peers are bounded by the map). Reconcile and re-dial backoffs are capped so a flapping gateway cannot induce a CPU storm.

13.2 Performance budget

Budget	Target	Source / rationale
Through-tunnel throughput (plain-udp)	$\geq 80\%$ bare link	G1 gate [04_P0 §8] — orchestrator overhead (mutex + loop dispatch) must be negligible vs WG crypto.
Roam / iface-flap recovery	< 3 s	[04_P0 §8] → §7.5 immediate re-dial.
Reconcile convergence (map delta → applied)	< 1 s (p99)	[04_ARCH §4.4] → §6 in-place apply.
Status emit latency (transition → broadcast)	< 1 ms (UNVERIFIED, measure) bounded, no per-packet alloc	side effects then emit (§4.3); dominated by the firewall syscall, not the channel.

Budget	Target	Source / rationale
Per-packet orchestrator overhead		reused scratch buffers (§3.1/§3.2/§3.3); Bytes ref-count, not copy, on the hot path.
Mutex hold time (wg/active_tx)	µs-scale; never across .await	§3.4 — synchronous helix-wg calls only.
Memory	bounded by peers × Tunn + fixed channels	I5/§12 host-safety; no per-connection durable state.

Each "UNVERIFIED" / "measure" budget is a Phase-0 benchmark item with captured CSV evidence before it is treated as fact (§11.4.6) [04_P0 §8].

14. Test points (§11.4.169)

Per §11.4.169 every workable item declares its required test types from the §11.4.27 / [07-WBS §1] closed vocabulary; every PASS cites rock-solid captured evidence (§11.4.5/.69/.107), never config-only. The escalation/reconnect **event trace** on the TunnelStatus stream IS the captured evidence for the state-machine claims (O-I13 analogue of [01-LAD I9]).

Test id	§11.4.169 types	What it proves	Rig / captured evidence	Source
OS-1	UNIT	transition() table is total + correct: every (state,event) pair → expected OrchState/None	table-driven unit test over the §5.2 matrix; mocks allowed (unit only, §11.4.27)	§5
OS-2	UNIT	diff() produces the exact MapDelta for added/removed/changed peers + dns/policy flags	golden desired/actual RouteMap pairs	§6.3
OS-3	UNIT	Backoff::next_delay() is monotonic-capped + jittered within [base, cap]; reset() zeroes	property test, fixed RNG seed (deterministic, §11.4.50)	§7.2
OS-4	INT	three loops move real WG datagrams tun[?]transport on the netns rig; curl reaches LAN host	netns + boringtun + plain-udp; tcpdump + curl 200	[04_P0 §8], §3
OS-5	E2E+REC			

Test id	§11.4.169 types	What it proves	Rig / captured evidence	Source
		full status arc Connecting→Handshaking→Connected{plain-udp, rtt} driven only by the core stream	Flutter-Linux toggle (G5); window-scoped MP4 + vision verdict (§11.4.159)	[04_PO §4
OS-6	E2E+FA	kill-switch closes on drop — no plaintext egress while Reconnecting/Down	block gateway mid-Connected; tcpdump underlay shows ZERO leaked app packets	AC7 [0 §11.2],
OS-7	SEC	DNS-leak guard: off-tunnel :53 dropped when Connected; system DNS restored on user-stop	dig @off-tunnel-resolver blocked; capture; revert verified	[04_AR §5.6], §
OS-8	E2E	reconcile restart-free: edit map.json → peer reachable, no process restart, OrchState unchanged	file-watch delta; curl newly-reachable host; pid unchanged (G6)	[04_PO §6
OS-9	E2E+PERF	revoke < 1 s: remove active peer → session torn, Down{auth-failed}, kill-switch closed	timestamped status trace; < 1 s revoke→drop	AC6 [0 §11.2],
OS-10	CHAOS	roam recovery < 3 s: flap the client iface → RoamDetected → re-dial → Connected	netns iface flap; status trace timing; recovery < 3 s	[04_PO §7.5
OS-11	CHAOS	mid-send transport swap + reconcile-source SIGKILL: loops keep forwarding (fail-static)	kill the file-watch task; traffic continues; reconciler re-subscribes	I3, §6.5
OS-12	STRESS	sustained re-dial churn (gateway flap N×) does not leak, deadlock, or grow memory unbounded	≥100 flap iters; RSS bounded; no leaked packet; no deadlock	§11.4.8

Test id	§11.4.169 types	What it proves	Rig / captured evidence	Source
OS-13	INT	connector posture: gateway→connector→LAN forward + reverse path, ACL-gated (default-deny)	netns LAN; allowed host reachable, unauthorized DROPPED + counted	§9, [04 §4.5]
OS-14	BENCH	orchestrator overhead: through-tunnel throughput ≥ 80% bare link (plain-udp)	iperf3 via bench.sh; CSV	G1 [04 §8], §1
OS-15	CHAL	a helix_qa/challenges bank entry scores PASS only on the captured status-arc + leak-free capture	HelixQA autonomous session, evidence-gated (§11.4.27/.107)	[SYNT §8]

Each meta-test ships a paired §1.1 mutation (e.g. flip `apply_side_effects` to emit `Connected` *before* opening the kill-switch → OS-6 must FAIL; weaken the verdict gate in §9.2 → OS-13 must FAIL) so the gates provably cannot bluff [04_PO §8, §11.4.5].

15. What is explicitly out of scope here

The ladder *selection algorithm* (per-network memory, regional prior, manual pin internals — [01-LAD]); the Transport trait + per-transport framing [01-DP §3]; WireGuard crypto + Tunn timers [01-DP §4]; `helix-route` ACL→verdict-map compilation [01-DP §6–§7]; DAITA machines [01-DP §9]; the `WatchNetworkMap` protobuf that delivers the map (doc 03); the FFI bridge that re-emits `TunnelStatus` to Dart (doc 05); platform `iface-watchers` + `TUN-fd` handoff beyond the `TunSpec` contract (doc 06). This document specifies only the **orchestrator** — the conductor that wires those frozen contracts into a running, leak-safe, self-converging tunnel on all three roles.

End of nano-detail specification — Volume 2 (Data Plane), orchestrator-and-state.md. Pairs with transport-selection-ladder.md (the §7.4 ladder this orchestrator drives), transport-trait.md (the §3 Transport contract), wireguard-core.md (the §3/§4 WgPeer it pumps), and routing-and-addressing.md (the §6 RouteMap/verdict map it reconciles). The Orchestrator struct (§2.1), the TunnelStatus enum (§4.1), the OrchState machine (§5), and the Mode::Connector posture (§9) are the frozen Phase-0 surviving contracts [04_PO §14]; their implementations may evolve, their shapes may not.