

UDP-over-TCP Transport

UDP-over-TCP Transport

Revision: 2 Last modified: 2026-07-04T12:00:00Z

Rev 2 (enterprise-hardening audit, 2026-07-04):
independently reviewed against the parent gap-analysis checklist. This document already contains the deepest resource-exhaustion treatment in Volume 2 — §9 E8 (slow-loris) / E12 (huge-LEN buffer-cap) edge cases and §12 S5 (reassembly-buffer DoS bound by `max_acc_bytes`) — no addition made here to avoid duplicating already-thorough content. No consistency issue found against transport-trait.md or 01-data-plane.md.

Master technical specification — Volume 2 (Data Plane), nano-detail document. Deepens the *udp-over-tcp — last resort when all UDP is blocked* section (§3.6) of [01-DP §3.6]. This is a SPEC — it describes the helix-transport::udp_over_tcp module to be built, **not** the shipping product. Source evidence cited inline by id: [01-DP §N] = final/01-data-plane.md; [04_ARCH §N] = 04_VPN_CLD/HelixVPN-Architecture-Refined.md; [04_P2 §N] = the Phase-2 refined doc; [SYNTHESIS §N] = the cross-document synthesis; [11-DR §N] = final/11-deep-research-appendix.md; [10-QA §N] = final/10-testing-acceptance-and-qa.md. External protocol constants not part of this evidence base (WireGuard message sizes, TCP/TLS header sizes) are tagged [WG-PROTO] / [TCP] / [TLS] and, where their exact value is not load-bearing, annotated honestly per constitution §11.4.6. Anything not grounded in the evidence base is marked UNVERIFIED.

0. Scope & contract summary

udp-over-tcp (UoT) is the **lowest rung of the transport escalation ladder** [01-DP §5.3]: it length-prefixes already-encrypted WireGuard datagrams over a **single TCP connection**, matching Mullvad's udp2tcp / the Rust udp-over-tcp crate [01-DP

§3.6, 04_ARCH §3.2, 04_P2 §1.2, 11-DR A4]. It exists for exactly one purpose: **keep a tunnel possible on networks where all UDP is blocked** — the failure mode where plain-udp, lwo, masque-h3 (QUIC/UDP), and shadowsocks-over-QUIC all fail because the path drops UDP entirely [11-DR A3 "Hard UDP block"]. It pays a **head-of-line-blocking (HoL) penalty** by construction (TCP imposes an ordered, reliable byte stream over WG's unreliable datagrams) and is therefore ranked **last** in every ladder order — preferred only when nothing faster connects [01-DP §3.6, §0.1 I2].

This document owns, to nano-detail:

1. The UdpOverTcp type and its full Transport implementation (§4).
2. The on-the-wire framing — the 2-byte big-endian length prefix + opaque WG datagram, the optional TLS wrap, byte-exact layouts (§5).
3. The connection lifecycle + framer/deframer state machines (§6).
4. The HoL-blocking tradeoff — why it is unavoidable, what it costs, how the orchestrator bounds its damage (§7).
5. Error taxonomy, config knobs, edge cases, security, performance budget (§8–§12).
6. The concrete test points, tied to the §11.4.169 test-type taxonomy [10-QA §1] (§13).

Invariants inherited unchanged [01-DP §0.1]:

#	Invariant	How UoT honours it
I1	Transport never sees plaintext.	UoT frames an opaque WG datagram; it never inspects, parses, or decrypts the payload — the length prefix is the only byte it adds.
I2	Transport carries unreliable datagrams, not an ordered stream.	UoT violates the spirit of I2 at the wire (TCP is ordered+reliable) but preserves the contract at the trait boundary : send/recv still exchange whole WG datagrams. WG's own loss/replay handling runs unchanged on top; the cost is HoL latency, not a semantic break (§7). This is the documented, accepted tradeoff of the last rung.
I4	One transport crate, three consumers.	The same <code>udp_over_tcp.rs</code> runs on client (frames out), edge (deframes in), and connector; framing is symmetric.
I5	No-logging by construction.	Only aggregate counters (bytes, frames, reconnects) — never per-datagram durable state.

File: `helix-transport/src/udp_over_tcp.rs` (the unit this doc specifies)

Trait: helix-transport/src/lib.rs::Transport (the contract;
frozen Phase-0 interface) [01-DP §3.1]
Config: helix-transport/src/lib.rs::TransportConfig::UdpOverTcp
{ peer, tls_sni } [01-DP §3.1]
Errors: helix-transport/src/error.rs::TransportError [01-DP §3.1]
Phase: 2 (full transport set) [01-DP §12, 04_P2 §1.2]; the
`Transport` trait it implements is frozen Phase-0.

1. When udp-over-tcp is selected — selection conditions

UoT is **ladder rung N (last)** in any restricted-network
TransportPolicy.order; on an unrestricted network it is never
reached [01-DP §5.3]. The orchestrator selects it in this closed set
of cases:

#	Condition	Source
C1	All UDP blocked — plain-udp, lwo, masquerade-h3, shadowsocks (QUIC-carried) each exceeded their per-rung failure budget; the ladder has escalated to its final rung.	[01-DP §5.3 step 3], [11-DR A3]
C2	Per-network memory hit — a prior successful UoT connect on this SSID / gateway fingerprint is remembered and retried first, so the user does not pay the full escalation latency again on the same hostile network.	[01-DP §5.3 step 4], [04_P2 §1.4]
C3	Regional prior — the coordinator pushed a TransportPolicy.order that ends in (or, for a known UDP-blackholing region, <i>starts higher up but still terminates at</i>) udp-over-tcp.	[01-DP §5.3 step 5], [04_P2 §1.4]
C4	Manual pin — user sets TransportPolicy.pin = UdpOverTcp (rare; deliberate "make it work on this captive/corporate network" choice).	[01-DP §5.3 step 1], [04_ARCH §3.2]

UoT is **never** the default order[0]. The default unrestricted order is [plain-udp]; the restricted prior is [plain-udp, lwo, masquerade-h3, shadowsocks, udp-over-tcp] [01-DP §5.3]. Rule: the orchestrator **MUST** exhaust every UDP-capable rung before UoT, because UoT's HoL penalty makes it the worst-performing carrier (§7).

UNVERIFIED (TLS-443 DPI residual): [11-DR A3] notes that on a path with **SNI/TLS DPI on TCP 443**, a naive TLS-wrapped UoT may still be fingerprinted unless it presents a real-site TLS fingerprint (VLESS+Reality / ShadowTLS

class). HelixVPN's MVP UoT ships the *plain* and *opportunistic-TLS-camouflage* variants only (§5.4); a Reality-grade TLS-mimic UoT is a **Phase-3 hardening item**, not specified here. Where TLS-443 DPI also blocks UoT, the honest outcome is EndpointBlocked → no tunnel (the residual-censorship boundary), not a silent degraded PASS.

2. Reference implementations & build stance

Reference	What it gives us	Stance
Mullvad udp-over-tcp Rust crate	The canonical 2-byte-length-prefix framing this module reproduces; Mullvad's udp2tcp obfuscation mode is literally this crate.	Reuse the framing design; reimplement in-crate so UoT shares helix-transport's Transport trait, error taxonomy, and HealthCell (I4) rather than bolting on a foreign socket loop. [11-DR A4], [04_P2 §1.2]
dndx/phantun (fake-TCP, 12-byte overhead, user-mode TCP, 100% safe Rust)	An alternative <i>fake-TCP</i> approach (forges TCP segments over a raw/AF_PACKET socket; no real kernel TCP, no HoL because it does not implement TCP retransmit).	Phase-3 option , not MVP. Phantun needs CAP_NET_RAW / raw-socket privileges and a kernel that lets it inject TCP segments — heavier deploy surface. Recorded as a documented alternative; MVP UoT uses a real kernel TCP socket (tokio::net::TcpStream) for portability. [11-DR A3, A4]

Decision (build stance): MVP `udp_over_tcp.rs` = **real kernel TCP** (tokio::net::TcpStream), Mullvad-style 2-byte length framing, optional opportunistic TLS wrap. Fake-TCP (Phantun-style) is a Phase-3 escape hatch for environments where even real-TCP UoT is throttled but raw-socket injection is available. This keeps the MVP last-rung **portable and unprivileged** at the cost of real-TCP HoL (§7) — the accepted tradeoff for "keep a tunnel *possible*" [01-DP §3.6].

3. Module surface (types this document specifies)

```
// helix-transport/src/udp_over_tcp.rs
use async_trait::async_trait;
use bytes::{Bytes, BytesMut, Buf, BufMut};
use std::net::SocketAddr;
use std::sync::Arc;
use tokio::io::{AsyncReadExt, AsyncWriteExt};
use tokio::sync::Mutex;

use crate::error::TransportError;
use crate::health::HealthCell;           // shared with plain-udp/
    masque [01-DP §3.2]
use crate::{Transport, TransportHealth};

/// Wire constant: the length prefix is a 2-byte big-endian
/// (network order) u16,
/// giving a 0..=65535 byte datagram budget – strictly larger than
/// any WG datagram
/// at any effective_mtu (WG max is well under 1500). [01-DP §3.6,
/// §10]
pub const UOT_LEN_PREFIX_BYTES: usize = 2;
pub const UOT_MAX_DATAGRAM: usize = u16::MAX as usize; // 65535

/// The carrier sub-mode: plain TCP, or TCP wrapped in
/// opportunistic TLS to look like HTTPS.
#[derive(Clone, Debug)]
pub enum UotWrap {
    Plain,
    /// Opportunistic TLS camouflage; `sni` is the ClientHello SNI
    /// presented to blend with HTTPS.
    /// NOT a security boundary – WG already provides
    /// confidentiality (I1). [01-DP §3.6 "Optional TLS wrap"]
    Tls { sni: String },
}

/// Live UoT transport. ONE TCP connection; `send`/`recv` exchange
/// whole WG datagrams.
pub struct UdpOverTcp {
    /// Write half behind a mutex: TCP writes must be atomic per-
    /// frame (no interleave). (§6.4)
    wr: Arc<Mutex<WriteHalf>>,
    /// Read half + deframer state behind a mutex: recv() is the
    /// sole reader. (§6.5)
    rd: Arc<Mutex<ReadHalf>>,
    peer: SocketAddr,
    kind: &'static str,           // "udp-over-tcp" or "udp-over-tcp-
    tls"
    health: HealthCell,
}
```

```

/// Internal write side (either plain TcpStream write half or a
    TLS stream write half).
struct WriteHalf { inner: TcpOrTlsWrite }
/// Internal read side carrying the partial-frame reassembly
    buffer (§6.5).
struct ReadHalf { inner: TcpOrTlsRead, acc: BytesMut /*
    unconsumed bytes across reads */ }

/// Sum type so the same framing code drives plain TCP and TLS-
    wrapped TCP without generics leaking
/// into the Transport object. (`TcpOrTls*` wrap tokio TcpStream
    halves or tokio-rustls TlsStream halves.)
enum TcpOrTlsWrite { Tcp(/* tokio write half */), Tls(/* tokio-
    rustls write half */) }
enum TcpOrTlsRead { Tcp(/* tokio read half */), Tls(/* tokio-
    rustls read half */) }

```

UdpOverTcp is constructed only by dial() (§4.1); it is never directly instantiated by callers. The Arc<Mutex<..>> split (separate write/read locks) lets send and recv run concurrently (full-duplex) while serialising multiple sends against each other so frames never interleave on the wire (§6.4).

4. The Transport impl (nano-detail)

4.1 dial() — establishing the carrier

dial(TransportConfig::UdpOverTcp { peer, tls_sni }) constructs a live UdpOverTcp within the bounded dial timeout; on timeout/refusal it returns DialTimeout / HandshakeFailed, which the orchestrator treats as "this rung failed" — but since UoT is the **last** rung, a failure here means the ladder is exhausted and the orchestrator emits Down { reason } rather than escalating [01-DP §5.3, §3.1].

```

/// helix-transport/src/lib.rs::dial() dispatch arm for UdpOverTcp
pub(crate) async fn dial_udp_over_tcp(
    peer: SocketAddr,
    tls_sni: Option<String>,
    dial_timeout: std::time::Duration,    // from TransportConfig
                                           defaults (§11)
) -> Result<UdpOverTcp, TransportError> {
    // 1. TCP connect with a hard timeout (UDP-blocking firewalls
        often *also* RST :51820/tcp,
    //    so a fast, bounded connect is the common "blocked"
        signal). [11-DR A3]
    let tcp = tokio::time::timeout(dial_timeout,
    tokio::net::TcpStream::connect(peer))
    .await
    .map_err(|_| TransportError::DialTimeout)?    //
    timed out → ladder exhausted
    .map_err(|e| match e.kind() {

```

```

        std::io::ErrorKind::ConnectionRefused
        | std::io::ErrorKind::ConnectionReset =>
TransportError::EndpointBlocked, // RST = DPI/firewall
_ => TransportError::Io(e),
    })?;

// 2. Disable Nagle: WG datagrams are latency-sensitive; Nagle
//    would coalesce small frames
//    and add up to ~40 ms of artificial delay on top of the
//    HoL cost (§7). [TCP]
tcp.set_nodelay(true).map_err(TransportError::Io)?;

// 3. Optional TLS wrap (opportunistic HTTPS camouflage; NOT a
//    security boundary). (§5.4)
let (wr, rd, kind) = match tls_sni {
    None => split_plain(tcp),
    Some(sni) => {
        let tls = tls_client_handshake(tcp, &sni,
            dial_timeout) // tokio-rustls; ALPN "h2"/"http/1.1"
            .await
            .map_err(|e|
TransportError::HandshakeFailed(format!("uot-tls:
{e}")))?;
        split_tls(tls)
    }
};

Ok(UdpOverTcp {
    wr: Arc::new(Mutex::new(WriteHalf { inner: wr })),
    rd: Arc::new(Mutex::new(ReadHalf { inner: rd, acc:
BytesMut::with_capacity(8 * 1024) })),
    peer,
    kind,
    health: HealthCell::new(),
})
}

```

Notes:

- **No application handshake of its own.** UoT carries the WG handshake as its first framed datagrams; the carrier is "up" the moment the TCP (and optional TLS) connection is established. The *tunnel* is up when WG completes its Noise IK handshake **over** UoT — that liveness is the orchestrator's concern via TunnelStatus, not UoT's [01-DP §5.2].
- **EndpointBlocked vs DialTimeout:** a TCP **RST** (refused/reset) → EndpointBlocked (active DPI/firewall verdict); a connect **timeout** → DialTimeout (blackhole). Both are terminal for the last rung; the distinction feeds the aggregate censorship telemetry ("UoT EndpointBlocked in region R" vs "UoT timeout in region R", §13) [01-DP §5.3 step 6].

4.2 send — frame one WG datagram

```
#[async_trait]
impl Transport for UdpOverTcp {
    async fn send(&self, dg: Bytes) -> Result<(), TransportError>
    {
        // Hard cap: WG datagram must fit the 2-byte length prefix
        // (always true at any MTU,
        // but assert rather than truncate – I1 forbids mangling
        // the opaque payload). (§5.2)
        if dg.len() > UOT_MAX_DATAGRAM {
            return Err(TransportError::Oversize(dg.len()));
        }
        let mut frame =
            BytesMut::with_capacity(UOT_LEN_PREFIX_BYTES + dg.len());
        frame.put_u16(dg.len() as u16); // big-endian length
        // prefix [01-DP §3.6]
        frame.put_slice(&dg);          // opaque WG datagram,
        // byte-for-byte (I1)

        // One mutex-guarded write_all: the whole frame lands
        // contiguously; concurrent sends
        // serialise here so two frames never interleave on the
        // byte stream (§6.4).
        let mut wr = self.wr.lock().await;
        wr.inner.write_all(&frame).await.map_err(|e| {
            self.health.mark_send_error();
            map_io_to_transport(e) // BrokenPipe/Reset →
            // Closed/EndpointBlocked (§8)
        })?;
        //
        NOTE: no flush() needed – TCP_NODELAY is set, write_all
        hands bytes to the kernel.
        Ok(())
    }
}
```

4.3 recv — deframe the next WG datagram

```
async fn recv(&self) -> Result<Bytes, TransportError> {
    let mut rd = self.rd.lock().await;
    // Pull whole frames out of the reassembly accumulator;
    // read more bytes only when a
    // complete frame is not yet buffered. Cancel-safe: state
    // lives in `rd.acc`, not on the
    // stack, so a dropped recv() future (e.g. in select!)
    // loses no bytes. (§6.5)
    loop {
        if let Some(dg) = try_take_frame(&mut rd.acc)? { //
            (§6.6)
                self.health.mark_recv();
                return Ok(dg);
            }
        // Need more bytes: read into a scratch buffer, append
        // to acc.
    }
}
```



```

        let mut scratch = [0u8; 16 * 1024];
        let n = rd.inner.read(&mut
scratch).await.map_err(map_io_to_transport)?;
        if n == 0 {
            return Err(TransportError::Closed); // peer
FIN / half-close (§6.3)
        }
        rd.acc.put_slice(&scratch[..n]);
    }
}

```

4.4 kind, effective_mtu, health, close

```

fn kind(&self) -> &'static str {
    self.kind } // "udp-over-tcp" | "udp-over-tcp-tls"

fn effective_mtu(&self) -> u16 {
    // Inner WG MTU budget left after TCP/TLS + 2-byte length
    // prefix overhead. (§10)
    // Plain UoT ~1380; TLS-wrapped is lower (TLS record
    // overhead) – both reported here.
    match self.kind {
        "udp-over-tcp"      => 1380, // TCP MSS – 2-byte
        length prefix [01-DP §10]
        "udp-over-tcp-tls" => 1360, // additionally TLS
        record header+tag [TLS] (measure & tune, §10)
        _                    => 1380,
    }
}

fn health(&self) -> TransportHealth { self.health.snapshot() }

async fn close(&self) -> Result<(), TransportError> {
    // Graceful: flush + shutdown(Write) so the peer sees a
    // clean FIN, then drop. Idempotent. (§6.3)
    let mut wr = self.wr.lock().await;
    let _ = wr.inner.flush().await;
    let _ = wr.inner.shutdown().await; // ignore errors –
    // close is best-effort, idempotent
    Ok(())
}
}


```

effective_mtu() values are **derived** [01-DP §10 marks UoT ~1380 "derived"] and MUST be re-measured against the path in Phase 2 (record real MSS via the connected socket) rather than trusted as constants — the §11.4.6 honest stance: the table value is a planning floor, not a measured fact, until a BENCH/MEM run confirms it (§13).

5. Wire format (byte-exact)

5.1 The frame

UoT puts a sequence of **length-prefixed frames** onto the TCP byte stream. Each frame:



- **LEN**: u16, **big-endian / network byte order**, = the length L of the WG datagram that follows (NOT including the 2 prefix bytes). Range 0..=65535; in practice L ≈ 32..1420 [WG-PROTO] [01-DP §3.6, §10].
- **Payload**: exactly L bytes of **opaque** WG datagram — UoT never inspects it (I1).
- Frames are **back-to-back** on the stream with no separator, no checksum, no sequence number: TCP already provides ordering, reliability, and integrity, so UoT adds only the length prefix needed to recover datagram boundaries from the byte stream (the sole job of UoT framing) [01-DP §3.6].

5.2 Why 2 bytes / big-endian / no checksum

Choice	Rationale	Source
2-byte length	u16 covers 0..=65535, strictly larger than any WG datagram at any effective_mtu (max WG payload << 1500) [WG-PROTO]; a 1-byte prefix could not. Matches Mullvad udp-over-tcp.	[11-DR A4], [01-DP §3.6]
Big-endian	Network byte order is the wire convention; symmetric framer/deframer on both ends agree by spec, not by host endianness.	[TCP convention]
No checksum / no seq	TCP delivers an ordered, error-checked, reliable stream; re-adding a checksum or sequence number would be redundant overhead. WG's own AEAD already authenticates the payload.	[01-DP §3.6, I1]
Oversize is a hard error	If a caller ever hands a datagram > 65535 (impossible at real MTU but defended), send returns TransportError::Oversize — UoT	[01-DP §3.1, §10 rule 3]

Choice	Rationale	Source
	MUST NOT truncate or split the opaque payload (I1).	

5.3 Worked example (handshake-init datagram)

A WireGuard handshake-initiation datagram is 148 bytes on the wire [WG-PROTO; value cross-checked in 11-DR A3 "fixed 148-byte handshake-init"]. UoT frames it as:

```

00 94 | <148 opaque WG bytes>
└──┘
  └─ 0x0094 = 148 (big-endian u16)
total on TCP stream = 2 + 148 = 150 bytes

```

The deframer reads 0x0094 → expects 148 payload bytes → emits exactly the original 148-byte WG datagram up to recv. Byte-for-byte round-trip; no mangling (I1).

5.4 The optional TLS wrap (UotWrap::Tls)

When `tls_sni` is `Some(sni)`, the entire UoT frame stream rides **inside a TLS record stream** so a passive observer on TCP :443 sees an HTTPS-looking handshake + encrypted records, not a bare length-prefixed stream [01-DP §3.6 "Optional TLS wrap (`tls_sni`) to look like HTTPS"].

```

L1: IP
L2: TCP (:443/tcp typically)
    └─ TLS records (ClientHello SNI = <sni>, ALPN h2/http1.1)
← camouflage layer
    └─ UoT frames: [u16 LEN][WG datagram] [u16 LEN][WG
datagram] ...
                        └─ opaque WG datagrams (I1)

```

- **Camouflage only, NOT a security boundary.** WG already provides confidentiality + integrity (I1); the TLS layer exists to make the flow resemble HTTPS to DPI, not to protect the payload. Its keys are unrelated to WG keys.
- **SNI** is presented to blend with real HTTPS to the gateway's :443/tcp listener; the edge's :443 listener already serves a believable decoy site to non-tunnel TLS (the masquerade property described for `masque-h3`) [01-DP §3.3 "Masquerade"], so a probe that completes the TLS handshake but sends no valid framed WG handshake gets the decoy, not a UoT signature.
- **ALPN** SHOULD advertise h2 / http/1.1 to match a browser; the inner bytes are UoT frames regardless of the advertised protocol (the ALPN is camouflage, the gateway routes by listener config). UNVERIFIED: exact ALPN/JA3 mimicry quality is

a Phase-3 DPI-hardening concern (§1 TLS-443 DPI note); MVP ships a standard tokio-rustls ClientHello.

6. State machines & lifecycle

6.1 Carrier lifecycle (the `UdpOverTcp` object)

[MERMAID RENDER FAILED — block #1]

6.2 Send-path framing FSM (per datagram)

[MERMAID RENDER FAILED — block #2]

6.3 Close & half-close

- **Graceful close (`close()`):** flush write buffer → `shutdown(Write)` (sends TCP FIN) → drop. **Idempotent** — second `close()` is a no-op (locks + best-effort flush/shutdown ignore already-closed errors). The peer's `recv` then observes `read == 0` and returns `Closed`.
- **Peer FIN during `recv`:** `read()` returns 0 → `recv` returns `TransportError::Closed`. The orchestrator treats this as a transport drop → Reconnecting (re-dial UoT) or, if re-dial also fails, Down [01-DP §5.2].
- **No silent half-open:** UoT MUST NOT keep sending into a write-closed socket pretending it is healthy; the first send after a peer RST surfaces `BrokenPipe/Reset` → `Closed/EndpointBlocked` via `map_io_to_transport` (§8). (Honesty over a fake-up tunnel — §11.4.6.)

6.4 Write atomicity (why the write half is mutex-guarded)

A WG datagram's frame (`[len][payload]`) MUST land **contiguously** on the TCP stream; if two send calls interleaved their `write_all`s, the deframer on the far end would read a corrupted length prefix. The `Arc<Mutex<WriteHalf>>` serialises all sends: each holds the lock across its single `write_all(&frame)`, so frames are atomic with respect to one another. `recv` uses a **separate** lock (`Arc<Mutex<ReadHalf>>`), so the connection is full-duplex (a send never blocks a `recv`). This is the data-race / deadlock surface the RACE test type targets (§13).

6.5 Receive-path reassembly (the partial-frame problem)

TCP delivers a **byte stream**, not datagrams: one `read()` may return half a frame, several frames, or a frame split across two reads. The deframer keeps an accumulator `acc: BytesMut` across `recv` calls and is **cancel-safe** (the accumulator lives in `ReadHalf`, not on the `recv` future's stack, so dropping a `recv` in a `select!` loses no buffered bytes — the I2 cancel-safety contract [01-DP §3.1]).

[MERMAID RENDER FAILED — block #3]

6.6 `try_take_frame` (the pure deframer — fully testable in isolation)

```
/// Pull exactly one complete frame out of `acc` if present;
/// otherwise return Ok(None).
/// Pure function of the buffer → the `UNIT`-testable core of the
/// deframer (§13).
fn try_take_frame(acc: &mut BytesMut) -> Result<Option<Bytes>,
    TransportError> {
    if acc.len() < UOT_LEN_PREFIX_BYTES {
        return Ok(None); // not even a
        // length prefix yet
    }
    // Peek the big-endian u16 length WITHOUT consuming it.
    let len = u16::from_be_bytes([acc[0], acc[1]]) as usize;
    let total = UOT_LEN_PREFIX_BYTES + len;
    if acc.len() < total {
        return Ok(None); // payload not
        // fully arrived
    }
    acc.advance(UOT_LEN_PREFIX_BYTES); // drop the
    // prefix
    let dg = acc.split_to(len).freeze(); // take exactly
    // `len` payload bytes
    Ok(Some(dg)) // remaining
    // bytes (next frame, partial) stay in acc
}
```

Edge cases this function MUST handle (each a UNIT test, §13):

1. `acc` empty → `None`.
2. `acc` = 1 byte → `None` (header incomplete).
3. `acc` = `[0x00, 0x94]` only → `None` (header complete, 0 payload bytes arrived).
4. `acc` = exactly one full frame → `Some(dg)`, `acc` emptied.
5. `acc` = 1.5 frames → `Some(first)`, second frame's partial bytes retained.
6. `acc` = 3 back-to-back frames → three successive `Some` across three calls, no read in between.

7. `LEN == 0` (a zero-length datagram) $\rightarrow$ `Some(empty Bytes)`. **Policy:** a 0-length frame is a protocol anomaly (WG never sends empty datagrams [WG-PROTO]); the *transport* returns it faithfully, and *helix-wg* drops it — UoT does not editorialise the opaque payload (I1). A flood of 0-length frames is a DoS vector handled by the read-rate guard (§9 G4, §12).

7. The head-of-line-blocking tradeoff (the defining property)

7.1 Why HoL is unavoidable here

WG datagrams are **independent and loss-tolerant** (I2): if datagram k is lost, datagram $k+1$ is still useful, and WG's own state machine tolerates loss/reorder/replay. Carrying them over **reliable, ordered TCP** inverts that: if the TCP segment carrying datagram k is lost, TCP **holds back every later byte** (including datagrams $k+1$, $k+2$, ... already received) until k is retransmitted and ACKed. That stall is **head-of-line blocking** — the structural cost of running a datagram protocol over a stream protocol [01-DP §0.1 I2, §3.6, 04_P2 §1.2].

[MERMAID RENDER FAILED — block #4]

7.2 What it costs & how the orchestrator bounds the damage

Effect	Magnitude	Mitigation in spec
Added latency under loss	Each loss event stalls delivery by ≈ 1 TCP RTO (or fast-retransmit RTT). On a 5% loss path this is frequent — exactly the scenario where masquerading3/QUIC keeps higher goodput [01-DP §3.3 "Loss resilience", 04_P0 §5.3].	The ladder prefers QUIC carriers precisely because of this; UoT runs only when UDP (hence QUIC) is impossible (§1). The ladder ranks UoT last [01-DP §5.3].
Throughput ceiling	TCP congestion control + HoL caps goodput below plain-udp/QUIC on lossy/throttled links.	Accepted: "exists purely to keep a tunnel <i>possible</i> " [01-DP §3.6]. <code>effective_mtu</code> and a BENCH run quantify it (§13).

Effect	Magnitude	Mitigation in spec
WG retransmit interaction	WG also retransmits its handshake; over reliable TCP, WG retransmits are rarely needed (TCP already delivered them) — mostly benign, occasionally redundant work.	No special handling — WG timers run unchanged (helix-wg::tick) [01-DP §4].
TCP-in-TCP meltdown risk	If the <i>inner</i> tunnelled traffic is itself TCP (the common case), you have TCP-over-WG-over-TCP; under loss the two congestion controllers can interact pathologically.	UNVERIFIED (classic "TCP meltdown" lore; not measured in this evidence base). Documented as a known risk; UoT is last-resort so the exposure window is minimised. A STRESS/PERF run under netem loss characterises it (§13) rather than assuming it.

Design stance: UoT does **not** try to fix HoL (that would mean reimplementing QUIC). It **accepts** HoL as the price of working on a UDP-blackholed network, and the *system* bounds the damage by only ever reaching UoT as the last rung. This is the honest, no-bluff framing: UoT is the transport that **works when nothing else can**, explicitly at the cost of being the slowest [01-DP §3.6].

8. Error taxonomy (mapping IO → TransportError)

UoT reuses the shared TransportError enum [01-DP §3.1]; it adds no new variants. The mapping from `std::io::Error` (and TLS error) to TransportError is the load-bearing contract:

```
fn map_io_to_transport(e: std::io::Error) -> TransportError {
    use std::io::ErrorKind::*;
    match e.kind() {
        BrokenPipe | ConnectionAborted =>
            TransportError::Closed,      // local/peer closed write
        ConnectionReset =>
            TransportError::EndpointBlocked, // RST mid-session =
            active block
        UnexpectedEof =>
            TransportError::Closed,      // peer FIN mid-frame
```

```

        TimedOut                                     =>
        TransportError::DialTimeout,    // (read/connect timeout)
                                         =>
        -TransportError::Io(e),
    }
}

```

TransportError	UoT trigger	Orchestrator response [01-DP §5.3]
DialTimeout	TCP connect (or TLS handshake) exceeded dial_timeout.	Last rung → Down { reason } (ladder exhausted).
HandshakeFailed(String)	TLS handshake error (UotWrap::Tls).	Last rung → Down.
EndpointBlocked	TCP RST on connect or mid-session reset — active firewall/ DPI verdict.	Down; record "UoT EndpointBlocked region R" (aggregate, I5).
Closed	Peer FIN / broken pipe / EOF mid-frame.	Reconnecting → re-dial UoT (C2); persistent → Down.
Oversize(usize)	send got a datagram > 65535 (defensive; impossible at real MTU).	Bug-class; surfaced, never truncated (I1).
Io(std::io::Error)	Any other IO error.	Reconnecting; logged as aggregate counter only (I5).

§11.4.1 FAIL-bluff guard: a UoT failure MUST be a *genuine* transport failure, never a script artifact — e.g. a deframer panic on a malformed length prefix is a §11.4.1 violation; the deframer returns Closed/IO, it does not panic (the RACE/CHAOS tests assert no panic, §13).

9. Edge cases (closed enumeration)

#	Edge case	Required behaviour	Test
E1	Partial frame across reads (half a length	Accumulate in acc; emit only on a complete frame (§6.5).	UNIT (§6.6 cases 1–6)

#	Edge case	Required behaviour	Test
	prefix, then the rest).		
E2	Multiple frames in one read().	Emit each on successive recv calls without an extra read (§6.6 case 6).	UNIT
E3	LEN == 0 frame.	Return empty Bytes; helix-wg drops it; do not editorialise (I1).	UNIT (§6.6 case 7)
E4	send of an oversize datagram.	Err(0versize); never truncate/split (I1).	UNIT
E5	Peer RST mid-session.	recv/send → EndpointBlocked/Closed; orchestrator re-dials or Down.	CHAOS
E6	TCP connect succeeds but WG handshake never completes (port-forward to a black hole).	Carrier "up" but tunnel never Connected; the orchestrator's WG-handshake timeout fires → escalate/Down — UoT itself reports healthy TCP, honestly (it is carrying bytes).	E2E (negative)
E7	TLS handshake to a non-UoT :443 (decoy site answered).	TLS completes, but no valid framed WG handshake returns → WG handshake timeout → Down (the masquerade working <i>against</i> a misdirected client).	SEC
E8	Slow-loris-style trickle (1 byte/sec) feeding acc.	Read-idle + per-frame assembly timeout bounds buffering (§12); surface Closed on timeout, do not grow acc unbounded.	STRESS, DDOS (Phase 2)
E9	Concurrent sends from multiple tasks.	Serialised by the write mutex; frames never interleave (§6.4).	RACE, CONC
E10	recv future dropped in a select! then re-entered.	Cancel-safe: acc retains buffered bytes; no datagram lost (§6.5).	UNIT, RACE
E11	NAT/firewall idle-timeout kills an idle TCP connection silently.	WG keepalives (helix-wg::tick) generate periodic framed traffic [01-DP §4]; a dead	CHAOS

#	Edge case	Required behaviour	Test
E12	Maliciously huge LEN (e.g. 0xFFFF) with no payload following.	connection surfaces as Closed on next send. try_take_frame returns None and waits; the per-frame assembly timeout (§12) caps the wait → Closed; acc is bounded by UOT_MAX_DATAGRAM so memory cannot blow up.	SEC, STRESS

10. MTU & overhead budget

The inner WG MTU over UoT = `min(effective_mtu(), path-MTU-discovered)` [01-DP §10 rule 1].

Layer	Overhead	Note
L1 IP	20 (v4) / 40 (v6) [IP]	path-dependent
L2 TCP	20 base (+ options) [TCP]	TCP header
UoT length prefix	2	the only byte UoT adds [01-DP §3.6]
(TLS wrap, if <code>UotWrap::Tls</code>)	TLS record header 5 + AEAD tag 16 ≈ 21/record [TLS]	camouflage variant only
Reported effect <code>ive_mtu()</code>	plain 1380 , TLS 1360	both derived [01-DP §10]; measure & tune in Phase 2 (§13)

Rules (inherited [01-DP §10]):

1. Inner WG MTU = `min(transport.effective_mtu(), path-MTU)`.
2. DAITA padding (if active) is added *after* WG encrypt and counts against the transport budget, not the inner MTU [01-DP §10 rule 2, §9].
3. Oversize is a hard error — the orchestrator lowers the inner MTU rather than letting UoT truncate (I1) [01-DP §10 rule 3].
4. The `effective_mtu()` constants are **planning floors**, not measured facts (§11.4.6); a BENCH/MEM run replaces them with the real connected-socket MSS (§13).

11. Config knobs

TransportConfig::UdpOverTcp { peer, tls_sni } is the public, NetworkMap-resolved config [01-DP §3.1]. Internal tunables (defaults; coordinator may override per-network via the policy push):

Knob	Type	Default	Meaning	Source
peer	SocketAddr	from NetworkMap	gateway TCP endpoint (typically :443/tcp for TLS variant, or :51820/tcp).	[01-DP §3.1]
tls_sni	Option<String>	None	Some(sni) enables the opportunistic TLS wrap (§5.4).	[01-DP §3.1]
dial_timeout	Duration	5 s UNVERIFIED (planning default; not pinned in evidence base)	bound on TCP+TLS connect before DialTimeout.	derived
tcp_nodelay	bool	true (always)	disable Nagle (§4.1).	[TCP]
read_idle_timeout	Duration	30 s UNVERIFIED	no bytes (incl. WG keepalive) for this long → Closed (§9 E11).	derived
frame_assembly_timeout	Duration	10 s UNVERIFIED	a partial frame must complete within this; else Closed (§9 E12 slow-loris guard).	derived
max_acc_bytes	usize	UOT_MAX_DATAGRAM (65535)	reassemble buffer hard cap; a single frame can never exceed it (§9 E12).	spec

All Duration defaults are tagged UNVERIFIED because the evidence base does not pin them; they are planning values to be calibrated against the netns rig in Phase 2, **not** invented facts (§11.4.6). The per-network failure budget that governs *when the ladder escalates to/past UoT* lives in `TransportPolicy.budget`, not here [01-DP §5.3].

12. Security considerations

#	Concern	Spec response	Source
S1	Transport must not see plaintext (I1).	UoT frames an opaque WG datagram; it never parses/decrypts. The 2-byte prefix is the only added byte.	[01-DP §0.1 I1]
S2	TLS wrap is camouflage, not security.	WG provides confidentiality/integrity; the TLS layer only resembles HTTPS to DPI. Its compromise does not expose payload.	[01-DP §3.6, I1]
S3	DPI fingerprinting of a bare length-prefixed stream.	Plain UoT on a non-443 port is fingerprintable; the TLS variant + edge decoy-site masquerade raises the bar. Reality-grade mimicry deferred to Phase 3 (§1 UNVERIFIED note).	[11-DR A3], [01-DP §3.3]
S4	Active probing of the :443/tcp listener.	The edge serves a believable decoy to any TLS client that does not produce a valid framed WG handshake (§9 E7), returning nothing useful to a prober.	[01-DP §3.3 masquerade], [11-DR A3 active-probing]
S5	Reassembly-buffer DoS (huge LEN, slow-loris).	<code>max_acc_bytes</code> caps memory at one max datagram; <code>frame_assembly_timeout + read_idle_timeout</code> bound trickle attacks; surface Closed, never OOM (§9 E8/E12).	spec (§11)
S6	No-logging by construction (I5).	Only aggregate counters (bytes, frames, reconnects,	[01-DP §0.1 I5],

#	Concern	Spec response	Source
		EndpointBlocked-per-region); no per-datagram/per-connection durable state.	[SYNTHESIS §7]
S7	Deframer robustness (no panic on malformed input).	try_take_frame is total over arbitrary byte input — it returns None/Closed, never panics (a panic would be a §11.4.1 FAIL-bluff + a DoS).	§6.6, §11.4.1
S8	Secret handling.	UoT holds no WG keys; TLS keys (camouflage) are ephemeral, never logged (§11.4.10).	[§11.4.10]

13. Test points (§11.4.169 taxonomy [10-QA §1])

Per §11.4.169 the udp_over_tcp.rs workable item declares the closed set of test types below; any warranted-but-absent type is an explicit NOT_APPLICABLE: <reason> / §11.4.3 SKIP, never an omission [10-QA §1]. Four-layer coverage per §11.4.4(b) (pre-build gate + post-build + runtime-on-clean-target + paired §1.1 mutation) and §11.4.50 determinism (N=3 / N=10) apply to every PASS.

Code	Scope for UoT	Captured-evidence shape (§11.4.5/.69/.107)
UNIT	try_take_frame (§6.6 cases 1–7), send frame builder, map_io_to_transport, effective_mtu per wrap. Mocks allowed only here (§11.4.27). Real	cargo test pass log; property test: deframe(frame(dg)) == dg for random dg (round-trip identity, I1).
INT	tokio::net::TcpStream loopback (+ tokio-rustls TLS variant) carrying real WG datagrams between two helix-wg instances; no mocks.	testcontainers / loopback log; pcap of the framed stream.
E2E	netns rig with all UDP dropped (nft drop udp); ladder	curl HTTP 200 body hash from overlay netns + pcap showing TCP-only carriage +

Code	Scope for UoT	Captured-evidence shape (§11.4.5/.69/.107)
FA	escalates to UoT; curl reaches an authorized LAN host through the real tunnel [04_P0 §3 rig]. The above E2E self-driving, re-runnable -count=3, zero human-in-loop (§11.4.98).	StatusReport.transport=="udp-over-tcp". make qa one-shot exit 0, identical evidence hashes ×3 (§11.4.50).
SEC	Decoy-site response to a TLS probe that sends no valid WG handshake (E7); deframer no-panic on fuzzed input (S7); no plaintext on wire (I1) via pcap.	pcap (no WG signature in clear), fuzz corpus pass, decoy-response capture.
STRESS	≥100 iters / ≥30 s sustained framed throughput; ≥10 parallel sends (write-mutex serialisation); slow-loris + huge-LEN buffer-cap (E8/E12) [§11.4.85].	stress_chaos.sh ab_stress_* latency.json; acc bytes bounded ≤ 65535 (counter).
CHAOS	Mid-session TCP RST / FIN / iface flap; recovery re-dials UoT and restores reachability (E5/E11) [§11.4.85].	ab_chaos_* recovery_trace.log; TunnelStatus event trace Connected→Reconnecting→Connected.
CONC	Concurrent multi-task send + recv full-duplex; no interleaved frames, no lost update (E9/E10).	Go/Rust harness pass; ordered-frame assertion.
RACE	cargo +nightly loom / tsan on the write-lock / read-lock + acc cancel-safety (§6.4/§6.5).	loom report 0 findings; no panic.
MEM	acc reassembly buffer never exceeds max_acc_bytes; no leak over a soak; (iOS	/proc/Instruments RSS flat over soak; buffer high-water counter.

Code	Scope for UoT	Captured-evidence shape (§11.4.5/.69/.107)
BENCH	path inherits the general MEM ceiling — UoT itself is allocation-light). UoT goodput + p99 framing latency vs plain-udp baseline, under netem loss 0/1/5% — quantifies the HoL penalty (§7) so the ladder cost model is real.	bench.sh CSV [04_P0 §8]; goodput-vs-loss curve showing UoT degrading faster than QUIC (the documented tradeoff).
PERF	latency p50/p99 vs SLO budget; TCP-in-TCP interaction characterised under loss (§7.2 UNVERIFIED risk → measured, not assumed).	histogram metrics; loss-sweep report.
REC	window-scoped MP4 of the Console/Access showing Connected { transport: "udp-over-tcp" } on a UDP-blocked network (§11.4.154/.155) + media-validation verdict (§11.4.163).	panoptic capture → vision_engine PASS verdict.
CHAL/ HQA	A challenges/helix_qa bank entry scoring PASS only on the E2E captured evidence above (§11.4.27/.107).	bank result.json PASS citing the pcap + curl artifacts.
DDOS	handshake-flood / huge-LEN flood against the edge :443/tcp UoT listener.	NOT_APPLICABLE: single-node-selfhost (Phase 2) [10-QA §1, 04_P1 §11] — parked, re-arms with HA.
SCALE	N simulated UoT clients holding TCP carriers.	NOT_APPLICABLE: single-node-selfhost (Phase 2) [10-QA §1]; partial soak only at MVP.
UI/UX	covered by REC + the shared client-core status walkthrough	flutter golden + flow walkthrough vision verdict.

Code	Scope for UoT	Captured-evidence shape (§11.4.5/.69/.107)
	(transport label surfaces in the UI).	

Anti-bluff acceptance gate (the one that matters): the headline UoT claim — "*a tunnel comes up when all UDP is blocked*" — is proven by the E2E/FA row: nft drops **all** UDP, the ladder escalates to UoT, and curl reaches a real LAN host through the real tunnel, with a pcap proving TCP-only carriage and **zero** UDP egress. Config-only / "the code path exists" PASS is forbidden (§11.4 / §11.4.69) — the captured curl 200 + pcap is the evidence [01-DP §12.1, 04_P2 §1.2 acceptance].

14. Build checklist & cross-references

Item	Owns	Phase	Gate
helix-transport/src/udp_over_tcp.rs	UdpOverTcp, framing, dial_udp_over_tcp, deframer	2	— (Phase-2 transport set) [01-DP §12]
TransportConfig::UdpOverTcp arm in lib.rs::dial()	dispatch	2	—
helix-core/src/ladder.rs last-rung ordering	UoT ranked last in restricted order	1→2	[01-DP §5.3]

Cross-references: the Transport trait + dial() ladder [01-DP §3.1]; the escalation sequence [01-DP §5.3, §5.4]; the MTU budget [01-DP §10]; plain-udp (the baseline UoT is measured against, §7.2) [01-DP §3.2 / transport-plain-udp.md]; masque-h3 (the QUIC carrier the ladder prefers *before* UoT) [01-DP §3.3 / transport-masque-quic.md]; shadowsocks (the sibling TCP-riding carrier with the same HoL caveat) [01-DP §3.5]; the §11.4.169 taxonomy [10-QA §1].

End of UDP-over-TCP transport nano-detail spec (Volume 2, Data Plane). Deepens [01-DP §3.6]. The Transport trait it implements is a frozen Phase-0 contract [01-DP §3.1]; this module is a Phase-2 deliverable [01-DP §12, 04_P2 §1.2]. All UNVERIFIED tags mark facts not grounded in the evidence base (TLS-443 DPI residual, TCP-in-TCP meltdown magnitude, the timeout defaults) and are to be resolved by measurement, not assumption (§11.4.6).