

coordinator service (the brain)

coordinator service (the brain)

Revision: 2 Last modified: 2026-06-26T12:00:00Z

Master technical specification — Volume 3 (Control Plane, Go), document `svc-coordinator`. Scope: the **internal/coordinator** package — HelixVPN's brain. It holds a per-tenant in-memory topology graph (hydrated from Postgres on boot, kept fresh by events), owns every open `WatchNetworkMap` stream, and on each event computes the **minimal affected agent set** and pushes a `MapDelta` to exactly those agents — convergence $p99 < 1$ s. This document *deepens* the coordinator part of [02_CP §6] to nano-detail and is implementation-ready. It is a SPEC (describe the implementation; do not build the product). Evidence cited inline by id: [02_CP §N] (pass-1 control-plane spec `final/02-control-plane.md`), [04_P1 §N] (`HelixVPN-Phase1-MVP.md`), [04_ARCH §N] (`HelixVPN-Architecture-Refined.md`), [research-go_cp §N] (`v09-research/research-go_cp.md`), [SYNTHESIS §N]. Unproven facts are marked UNVERIFIED per §11.4.6.

0. What this package owns, and what it does NOT

The coordinator is the **only** stateful-in-RAM, stateless-on-disk module. It is the runtime projection of Postgres truth + the live event stream into per-agent network maps.

It OWNS	It does NOT own
The per-tenant in-memory Graph (devices, connectors, prefixes, groups, compiled policy, presence).	Any durable table. R2: coordinator MUST NOT create/migrate a Postgres table [02_CP §1.2-R2, 04_P1 §4].
The registry of open <code>WatchNetworkMap</code> streams (subRegistry).	Identity/enrollment (owned by identity), IPAM, PKI issuance — read through their interfaces [02_CP §1.3].
The fan-out engine: event → minimal affected set → <code>MapDelta</code> enqueue → <code>stream.Send</code> .	Policy <i>compilation</i> (owned by policy; coordinator consumes <code>CompiledPolicy</code>) [02_CP §7].

It OWNS	It does NOT own
The monotonic per-tenant map version counter.	The event <i>bus transport</i> (owned by events; coordinator is a consumer) [02_CP §5].
The <code>helix_reconcile_seconds</code> SLO histogram + back-pressure.	The Rust data path (edge reconciles WG peers from the delta — doc 01) [C1].

Governing invariants inherited verbatim from [02_CP §0.1]: **C1** Go never in the packet path (fail-static); **C2** Postgres = truth, Redis = ephemeral; **C4** default-deny need-to-know (a node only ever learns peers a compiled rule grants); **C5** push-don't-poll, snapshot-then-deltas, $p99 < 1\text{ s}$; **C7** package boundary == future service boundary; **C8** RLS at the DB. The coordinator is the concrete embodiment of **C4** + **C5**.

The proto package is **helix.coordinator.v1** (unified across the document set); generated Go import alias `coordinatorv1` [02_CP §4 — note: the §4 `.proto` option `go_package` path retains the historical `helix/coordinator/v1` directory but the wire package is `helix.coordinator.v1`; this document uses the package name as canonical].

1. The in-memory topology graph

1.1 Position in the system

flowchart LR

```

    subgraph PG[Postgres truth C2]
        T[(devices / connectors /advertised_prefixes /<br/>>groups / policies)]
    end
    subgraph BUS[Redis Streams D3]
        E[(events:devices · routes ·<br/>policy · presence · gateway)]
    end
    subgraph COORD[internal/coordinator the brain]
        H[hydrator<br/>boot snapshot]
        G[(per-tenant Graph<br/>RWMutex-guarded)]
        A[applier<br/>event to graph mutation]
        F[fan-out engine<br/>minimal affected set]
        S[subRegistry<br/>open WatchNetworkMap streams]
    end
    EDGE[Rust edge / agents<br/>reconcile WG peers C1]

    T -- boot hydrate --> H --> G
    E -- XReadGroup --> A -- mutate --> G
    A -- changed node set --> F
    G -- buildMap/diff --> F
    F -- enqueue MapDelta --> S
    S -- stream.Send --> EDGE

```

1.2 Graph data structures (Go, nano-detail)

The graph is per-tenant. The coordinator holds a top-level map[TenantID]*tenantState, each tenantState independently locked so one tenant's churn never blocks another (per-tenant sharding is the concurrency unit) [04_P1 §4].

```
// internal/coordinator/graph.go
package coordinator

import (
    "net/netip"
    "sync"
    "time"

    "github.com/google/uuid"
)

type (
    TenantID = uuid.UUID
    DeviceID = uuid.UUID
)

// DeviceKind mirrors the DB enum device_kind and the proto enum
// DeviceKind [02_CP §2.2/§4].
type DeviceKind uint8

const (
    KindUnspecified DeviceKind = iota
    KindClient
    KindConnector
)

// PresenceState is the coordinator's view of an agent's liveness
// (Redis-fed, ephemeral C2).
type PresenceState uint8

const (
    PresenceUnknown PresenceState = iota // never seen since boot
    PresenceOnline                       // an open stream OR a
        // fresh heartbeat
    PresenceOffline // stream closed AND
        // heartbeat TTL expired
)

// node is one device (client or connector) in the graph.
type node struct {
    ID          DeviceID
    Kind        DeviceKind
    WGPubKey    [32]byte // Curve25519 public key only (C6: private never
        // registered)
    OverlayIP   netip.Addr
    OS          string
}
```

```

    Presence    PresenceState
    LastRTTMS   uint32        // from ReportStatus; health only, NOT
                             traffic (C3)
    LastSeen    time.Time     // coarse; refreshed by heartbeat, not
                             per-packet (C3)
    Revoked     bool          // belt-and-suspenders mirror of
                             devices.revoked_at
    siteID      uint32        // connectors only: 4via6 site id (D4);
                             0 for clients
}

// connectorAdv holds a connector's advertised reachable CIDRs.
type connectorAdv struct {
    connectorID DeviceID
    prefixes    []netip.Prefix // from advertised_prefixes WHERE
                             enabled
}

// tenantState is the per-tenant graph + version + the policy
// projection.
type tenantState struct {
    mu sync.RWMutex

    tenantID TenantID
    version  uint64 // monotonic map version; bumped on every
                  accepted mutation (C5)

    nodes      map[DeviceID]*node
    connectors  map[DeviceID]*connectorAdv // subset of nodes
                             with kind=connector
    gateway     GatewayInfo                // current gateway
                             endpoint for this tenant
    dns         DnsConfig                  // tenant DNS config
    transport   TransportPolicy            // escalation ladder [02_CP
    §4]

    // compiled policy projection – the need-to-know truth (C4).
    // Replaced wholesale on
    // policy.compiled; never mutated field-by-field (compilation
    // is pure+versioned §7 [02_CP]).
    policy *CompiledPolicy

    // versionedDeltas is a bounded ring of recent deltas keyed by
    // the version they produced,
    // enabling cheap known_version reconnect catch-up (§3.2).
    // Capacity DELTA_RING_CAP (§6.4).
    versionedDeltas *deltaRing
}

// GatewayInfo / DnsConfig / TransportPolicy mirror the proto
// messages 1:1 [02_CP §4].
type GatewayInfo struct {
    Endpoint  string
    WGPubKey  [32]byte

```

```

    MasqueSNI string
}
type DnsConfig struct {
    Resolvers []string
    Search     []string
}
type TransportPolicy struct {
    Order          []string // e.g. ["plain-udp", "lwo", "masque-
        h3"]
    AllowUserOverride bool
}

// CompiledPolicy is produced by internal/policy.Compiler.Compile
// and consumed read-only here
// [02_CP §1.3/§7]. The coordinator NEVER recomputes it; it diffs
// successive versions.
// Canonical definition (full field set) is owned by svc-policy.md
// §3 – this is the
// same struct, reproduced read-only. Do NOT let the field set
// drift from svc-policy.md.
type CompiledPolicy struct {
    Version      int64
    SpecHash     [32]byte // sha256
                of canonical spec (audit/dedup)
    VisibleTo    map[DeviceID]map[DeviceID]struct{} // need-
                to-know peer set (C4)
    AllowedIPs   map[DeviceID]map[DeviceID][]netip.Prefix // coarse
                WG AllowedIPs (ARTIFACT 1b)
    Verdicts     map[DeviceID]map[DeviceID][]PortRule // fine
                port-level map (ARTIFACT 2; edge-enforced)
    Via6         map[DeviceID]map[DeviceID][]Via6Route // 4via6
                mappings (D4)
    ExitNodes    map[DeviceID]struct{}
}
type Via6Route struct {
    IPv4CIDR netip.Prefix
    Via6Prefix netip.Prefix
}

```

Architectural note (§11.4.6): CompiledPolicy carries the full field set above (canonical in svc-policy.md §3), but the coordinator's buildMap only ships coarse AllowedIPs + Via6 over the wire — the proto Peer carries allowed_ips + via6 but **not** port verdicts [02_CP §4 Peer]. The fine Verdicts map is enforced at the **edge** (nftables/eBPF, doc 01); whether it reaches the edge via a *separate* coordinator stream field or out-of-band is a doc-01/doc-03 decision, NOT resolved here — marked pending, not fabricated.

1.3 Coordinator top-level type & lifecycle

```

// internal/coordinator/coordinator.go
type Coordinator struct {
    mu      sync.RWMutex
    tenants map[TenantID]*tenantState
}

```

```

registry registry.Registry // §1.3 [02_CP]; boot hydrate +
    presence reads
policy policy.Compiler // read CompiledPolicy
bus events.Bus // consume events (§4)
subs *subRegistry // open streams (§3)
metrics *metrics // helix_reconcile_seconds etc.
    (§7)
clock func() time.Time // injectable for deterministic
    tests (§8)

hostID string // consumer name for XReadGroup/XAutoClaim (§4)
    [research-go_cp §5]
}

// Boot lifecycle: NewCoordinator → Hydrate(ctx) → Run(ctx) [04_P1
    §4].
func NewCoordinator(deps Deps) *Coordinator { /* wire deps;
    clock=time.Now */ }

// Hydrate rebuilds every tenant graph from Postgres truth BEFORE
    consuming any event, so the
// graph is a faithful projection and the first delta is computed
    against a correct baseline (C2).
func (c *Coordinator) Hydrate(ctx context.Context) error

// Run starts the event-consumer loop (§4) + the XAutoClaim DLQ
    sweeper + the presence-TTL reaper.
// It blocks until ctx is cancelled; returns ctx.Err(). Idempotent
    reactions make replay safe.
func (c *Coordinator) Run(ctx context.Context) error

```

1.4 Hydration (boot) — exact sequence

```

sequenceDiagram
    autonumber
    participant Run as helixd boot
    participant Coord as coordinator.Hydrate
    participant Reg as registry
    participant Pol as policy.Compiler
    participant PG as Postgres (RLS)
    Run->>Coord: Hydrate(ctx)
    Coord->>Reg: TenantsWithDevices(ctx)
    loop each tenant t (WithTenant scope C8)
        Coord->>Reg: DevicesForTenant(t)
        Reg->>PG: SELECT devices/connectors/advertised_prefixes
        Coord->>Pol: ActiveCompiledPolicy(t)
        Pol->>PG: SELECT policies WHERE active
        Coord->>Coord: build tenantState{nodes, connectors, policy,
version:=active.Version}
    end
    Coord-->>Run: nil (ready) | error (abort boot – fail-closed)

```

Rules: (1) hydration runs inside `store.WithTenant` per tenant (RLS floor, C8) [02_CP §2.3]; (2) the initial version for a tenant is the **active policy version** — so a reconnecting agent's `known_version` is comparable across coordinator restarts (the version namespace is the policy-version namespace, not a process-local counter — this makes restarts transparent to agents, §3.2); (3) if `ActiveCompiledPolicy` returns none, the tenant graph is built default-deny (every node sees zero peers, C4) — never fail-open; (4) a hydration error for ANY tenant aborts boot (fail-closed, §11.4.6 — never serve a partially-true graph).

2. Map computation (per node)

`buildMap(node)` and `buildDelta(old, new)` are pure functions of (graph, policy) — given the same graph + policy version they are deterministic (a property test asserts this, §8) [02_CP §6.2].

2.1 buildMap — full snapshot for one node

```
// internal/coordinator/build.go (caller holds ts.mu.RLock)
func (ts *tenantState) buildMap(self *node)
    *coordinatorv1.NetworkMap {
    reachable := ts.policy.VisibleTo[self.ID]           // need-
        to-know filter (C4)
    peers := make([]*coordinatorv1.Peer, 0, len(reachable))
    for pid := range reachable {
        p := ts.nodes[pid]
        if p == nil || p.Revoked {
            continue //
            never advertise a revoked/absent peer
        }
        peers = append(peers, &coordinatorv1.Peer{
            DeviceId:    p.ID.String(),
            WgPubkey:     p.WGPubKey[:],
            AllowedIps:
            cidrsToStrings(ts.policy.AllowedIPs[self.ID][pid]), //
            compiled CIDs (C4)
            Endpoint:
            ts.relayEndpoint(p), // via gateway in MVP
            (hub-and-spoke)
            IsConnector: p.Kind == KindConnector,
            Via6:         via6ToProto(ts.policy.Via6[self.ID]
            [pid]), // 4via6 (D4)
        })
    }
    sortPeersByID(peers) //
        STABLE order ⇒ deterministic snapshot (§8)
    return &coordinatorv1.NetworkMap{
        Self:      &coordinatorv1.Node{DeviceId:
            self.ID.String(), OverlayIp: self.OverlayIP.String()},
        Gateway:   ts.gateway.proto(),
        Peers:     peers,
        Dns:       ts.dns.proto(),
    }
```

```

        Transport: ts.transport.proto(),
    }
}

```

relayEndpoint always returns the gateway relay endpoint in the MVP (hub-and-spoke; direct P2P is Phase 2 — Peer.endpoint gains NAT-traversal candidates without reshaping) [02_CP §6.2/§14].

2.2 buildDelta — minimal diff between two graph states for one node

The delta is computed against the node's **last-sent map** (the coordinator tracks per-stream lastVersion; for a fresh recompute it diffs the old vs new peer sets). The diff is by DeviceID:

```

// returns nil if the node's view is byte-unchanged (NO delta sent
//   – quiet for unaffected nodes)
func diffPeers(old, new map[DeviceID]*coordinatorv1.Peer)
    *coordinatorv1.MapDelta {
    var d coordinatorv1.MapDelta
    for id, np := range new {
        if op, ok := old[id]; !ok || !proto.Equal(op, np) {
            d.UpsertPeers = append(d.UpsertPeers, np) // added or
            changed
        }
    }
    for id := range old {
        if _, ok := new[id]; !ok {
            d.RemovePeerIds = append(d.RemovePeerIds,
                id.String()) // peer no longer visible (C4)
        }
    }
    if len(d.UpsertPeers) == 0 && len(d.RemovePeerIds) == 0 &&
        d.Transport == nil && d.Dns == nil {
        return nil // unchanged – this is how the minimal-
            affected-set stays minimal (§5)
    }
    sortUpserts(&d) // deterministic wire order (§8)
    return &d
}

```

Edge rule: transport/dns fields are populated in the delta **only if changed** [02_CP §4 MapDelta]. A node whose diffPeers returns nil AND whose transport/dns are unchanged receives **no** MapUpdate for that event (it stays quiet) — this is the literal mechanism of "never broadcast a full resync on a small change" [02_CP §6.4].

3. Open-stream registry & the WatchNetworkMap loop

3.1 subRegistry — the open-stream index

```
// internal/coordinator/subs.go
type subscription struct {
    tenantID    TenantID
    deviceID    DeviceID
    lastVersion uint64                // last version
    // this stream has converged to
    ch          chan *coordinatorv1.MapUpdate // bounded send
    // queue (cap SEND_QUEUE_CAP, §6.4)
    cancel      context.CancelFunc           // closes the
    // stream on revoke/back-pressure
    mu          sync.Mutex
}

type subRegistry struct {
    mu sync.RWMutex
    // tenant -> device -> set of subs (a device MAY hold >1
    // stream: reconnect overlap, multi-NIC)
    byTenant map[TenantID]map[DeviceID]map[*subscription]struct{}
}

func (r *subRegistry) add(s
    *subscription)                // registers +
    // flips presence online
func (r *subRegistry) remove(s
    *subscription)                // deregisters +
    // maybe-offline
func (r *subRegistry) forTenant(t TenantID)
    []*subscription              // fan-out enumeration
func (r *subRegistry) forDevice(t TenantID, d DeviceID)
    []*subscription
```

A single device may legitimately hold more than one stream momentarily (a reconnect lands before the old stream's TCP FIN; a multi-homed device). Fan-out enqueues to **every** matching subscription; presence flips offline only when the device's **last** subscription closes AND its heartbeat TTL has lapsed (§4.4).

3.2 WatchNetworkMap — the streaming handler (server side)

```
// internal/coordinator/watch.go
func (c *Coordinator) WatchNetworkMap(
    ctx context.Context,
    req *connect.Request[coordinatorv1.WatchRequest],
    stream *connect.ServerStream[coordinatorv1.MapUpdate],
) error {
    dev, err :=
        authDevice(ctx) // mTLS device cert -> device_certs ->
        // devices (§8.2 [02_CP])
```

```

if err != nil {
    return connect.NewError(connect.CodeUnauthenticated,
        err) // E-AUTH (§4 err taxonomy)
}
if dev.Revoked {
    return connect.NewError(connect.CodePermissionDenied,
        errRevoked) // E-REVOKED
}
if req.Msg.DeviceId != "" && req.Msg.DeviceId !=
    dev.ID.String() {
    return connect.NewError(connect.CodePermissionDenied,
        errDeviceMismatch) // E-IMPERSONATE
}

ts := c.tenant(dev.TenantID)
if ts == nil {
    return connect.NewError(connect.CodeFailedPrecondition,
        errNoTenant) // E-NOTENANT
}

sctx, cancel := context.WithCancel(ctx)
sub := newSubscription(dev.TenantID, dev.ID, cancel)
c.subs.add(sub) // registers + presence ONLINE +
                publishes device.online (§4.4)
defer func() {
    c.subs.remove(sub) // deregister + maybe presence OFFLINE
    + device.offline
    cancel()
}()

// (1) snapshot OR resume-from-known_version (cheap reconnect,
// C5) [02_CP §6.3, 04_P1 §3]
ts.mu.RLock()
cur := ts.version
switch {
case req.Msg.KnownVersion == 0 || req.Msg.KnownVersion > cur:
    snap := ts.buildMap(ts.nodes[dev.ID])
    ts.mu.RUnlock()
    if err := stream.Send(&coordinatorv1.MapUpdate{
        Version: cur, Body:
        &coordinatorv1.MapUpdate_Snapshot{Snapshot: snap},
    }); err != nil {
        return err // E-STREAM-SEND
    }
default: // 0 < known_version <= cur : replay catch-up deltas
    if the ring still has them
    deltas, complete :=
    ts.versionedDeltas.since(req.Msg.KnownVersion)
    snap := (*coordinatorv1.NetworkMap)(nil)
    if !complete {
        snap = ts.buildMap(ts.nodes[dev.ID]) // ring evicted
        the gap → fall back to snapshot
    }
    ts.mu.RUnlock()

```

```

    if snap != nil {
        if err := stream.Send(snapshotUpdate(cur, snap));
        err != nil { return err }
    } else {
        for _, d := range deltas {
            if err := stream.Send(d); err != nil { return
err }
        }
    }
}
sub.setLastVersion(cur)

// (2) live delta loop – events via sub.ch, keepalive on a
// ticker (C5)
ka := time.NewTicker(KEEPALIVE_INTERVAL) // 20 s [02_CP §6.3]
defer ka.Stop()
for {
    select {
    case <-sctx.Done():
        return sctx.Err() //
        client disconnect OR forced-close
    case up, ok := <-sub.ch:
        if !ok {
            return connect.NewError(connect.CodeAborted,
errBackpressure) // E-BACKPRESSURE
        }
        if err := stream.Send(up); err != nil { return err }
        sub.setLastVersion(up.Version)
    case <-ka.C:
        if err := stream.Send(keepAlive()); err != nil {
return err } // proves liveness
    }
}
}

```

Resume semantics (the known_version contract, C5):

- $\text{known_version} == 0 \Rightarrow$ full snapshot (first connect).
- $\text{known_version} > \text{current} \Rightarrow$ the client claims a version the server never issued (a stale/forked client, or a coordinator that lost+rebuilt state) $\Rightarrow$ full snapshot (safe re-baseline).
- $0 < \text{known_version} \leq \text{current} \Rightarrow$ replay the deltaRing entries ($\text{known_version}, \text{current}$]; if the ring evicted any entry in that span (gap), fall back to a full snapshot. This bounds reconnect cost while never serving an incomplete view.

3.3 Stream lifecycle state machine

stateDiagram-v2

[*] --> Authenticating

Authenticating --> Rejected: mTLS invalid / revoked /

impersonation

Authenticating --> Snapshotting: authDevice ok

Snapshotting --> Streaming: snapshot/catch-up sent

Streaming --> Streaming: MapDelta sent / KeepAlive sent
Streaming --> Draining: ctx cancelled (client FIN)
Streaming --> ForcedClose: revoke event OR back-pressure
overflow
ForcedClose --> Draining
Draining --> [*]: subRegistry.remove + presence maybe-OFFLINE
Rejected --> [*]

4. Event consumption & the reaction table (event → recompute)

The coordinator is a durable consumer-group member on the §5 [02_CP] Redis Streams. The consumer loop is **XAutoClaim-stalled-first, then XReadGroup-new** [research-go_cp §5].

4.1 Consumer loop (Go)

```
// internal/coordinator/consume.go
func (c *Coordinator) consume(ctx context.Context) error {
    streams := []string{"events:devices", "events:routes",
        "events:policy",
        "events:presence", "events:gateway"}
    for {
        if ctx.Err() != nil { return ctx.Err() }
        // 1. reclaim stalled PEL entries from crashed peers (no-
        // work-loss §11.4.147)
        for _, s := range streams {
            claimed, _, _ := c.bus.AutoClaim(ctx, s,
                "coordinator", c.hostID,
                EVENT_CLAIM_MIN_IDLE /*30s*/, 64)
            for _, env := range claimed { c.handleEnvelope(ctx,
                s, env) }
        }
        // 2. read new entries (blocking)
        batch, err := c.bus.ReadGroup(ctx, "coordinator",
            c.hostID, streams, 64, 5*time.Second)
        if err != nil { continue } // transient; loop reattempts
        (§11.4.6 no silent give-up)
        for _, env := range batch { c.handleEnvelope(ctx,
            env.Stream, env) }
    }
}

func (c *Coordinator) handleEnvelope(ctx context.Context, stream
    string, env events.Envelope) {
    t0 := c.clock()
    defer func() {
        if r := recover(); r != nil {
            c.metrics.applyPanics.Inc(); /* env stays in PEL →
            reclaimed */ return }
    }()
}
```

```

changed, err := c.apply(env) // mutate graph,
    return affected node set (§4.3)
if err != nil {
    if env.DeliveryCount >= EVENT_RETRY_CAP /*5*/ {
        c.bus.DeadLetter(ctx, stream, env) // route to
        events:<s>:dlq + ACK (§5.4 [02_CP])
        c.metrics.dlqTotal.Inc()
    }
    return // leave un-ACKed → XAutoClaim retries (at-least-
        once)
}
c.fanout(env.TenantID, changed) // enqueue MapDelta
    to minimal set (§5)
c.bus.Ack(ctx, stream, "coordinator", env.ID)
c.metrics.reconcile.Observe(c.clock().Sub(t0).Seconds()) //
    the < 1 s SLO timer (§7)
}

```

Idempotency (load-bearing): every reaction recomputes the affected nodes' maps from *current* graph state, so a replayed/duplicate event is harmless [02_CP §5.4]. apply MUST be a pure graph-state transition — never an increment that double-applies on redelivery.

4.2 Reaction table (event → graph mutation → recompute scope)

Concrete and exhaustive for the MVP event taxonomy [02_CP §5.3]. "Recompute scope" is the set of nodes whose map MIGHT change — the fan-out then diffs each and sends only real deltas (§2.2).

Event type	Payload	Graph mutation
device.enrolled	{device_id, kind, overlay_ip, wg_pubkey}	add node; if conn add connectorAd
device.online	{device_id}	set Presence=On refresh LastSeen
device.offline	{device_id}	set Presence=Of
device.revoked	{device_id}	set node.Revoke drop from nodes/ connectors
connector.attached	{device_id, site}	set node.siteID connectorAdv

Event type	Payload	Graph mutation
connector.prefixes.changed	{connector_id,cidrs[]}	replace connectorAdv.p
route.conflict.detected	{cidr,connector_ids[]}	annotate (no map
policy.compiled	{version}	load new CompiledPolicy ts.policy, vers
gateway.failover	{from,to}	replace ts.gatew

policy.updated (vs policy.compiled) is consumed by the policy module to trigger compile, NOT by the coordinator [02_CP §5.3]; the coordinator reacts only to policy.compiled (the compiled artifact is ready). Listing it here would be a fabricated reaction — omitted by design.

4.3 apply() — policy.compiled is the only tenant-wide recompute

```
// the heaviest reaction; all others touch O(small) nodes. Diff
// old vs new projection so a
// tenant-wide recompile still produces a MINIMAL affected set
// (§5) [02_CP §6.4].
func (c *Coordinator) applyPolicyCompiled(env events.Envelope)
    (changed map[DeviceID]struct{}, err error) {
    v := env.Payload["version"].(int64)
    fresh, err := c.policy.CompiledAt(env.TenantID, v) // read-
        only [02_CP §1.3]
    if err != nil { return nil, err } // → retry/
        DLQ (§4.1)
    ts := c.tenant(env.TenantID)
    ts.mu.Lock()
    defer ts.mu.Unlock()
    if v <= int64(ts.version) { return nil, nil } // stale/
        replayed compile – idempotent no-op
    old := ts.policy
    ts.policy = fresh
    ts.version = uint64(v)
    changed = projectionDiff(old, fresh) // nodes
        whose VisibleTo/AllowedIPs/Via6 differ
```

```

    ts.versionedDeltas.recordBoundary(uint64(v))           // ring
    boundary for known_version resume
    return changed, nil
}

```

4.4 Presence model (Redis-fed, ephemeral C2)

Presence is the union of two signals: (a) **an open WatchNetworkMap stream** (authoritative online), and (b) **a heartbeat** (ReportStatus/events:presence) with a TTL. A node is Online if it has ≥ 1 open stream OR a non-expired heartbeat; Offline once both lapse. A background **presence reaper** scans heartbeat TTLs every PRESENCE_REAP_INTERVAL and emits device.offline for expired-and-streamless nodes. Losing Redis loses presence but **no durable state** (C2) — on Redis recovery, open streams re-assert online; the graph itself was never in Redis.

4.5 Revoke fast-path (< 1 s, the security-critical reaction)

device.revoked is the one reaction with a dual action: (1) the standard fan-out (remove_peer_ids delta to every node that saw the device — C4); (2) subs.forDevice(t, revoked) → cancel() each, **force-closing the revoked device's own streams** so it stops receiving maps immediately. The edge separately removes the kernel WG peer [02_CP §9.3]. Revocation latency target == convergence SLO (§7). This is authDevice-redundant defense in depth (the stream handler also checks dev.Revoked at open time, §3.2).

5. Fan-out: minimal affected set → MapDelta

```

// internal/coordinator/fanout.go
func (c *Coordinator) fanout(t TenantID, changed
    map[DeviceID]struct{}) {
    if len(changed) == 0 { return }           // nothing
    affected → zero sends (§2.2)
    ts := c.tenant(t)
    for did := range changed {
        for _, sub := range c.subs.forDevice(t, did) {
            ts.mu.RLock()
            newMap := ts.buildMap(ts.nodes[did])
            ver := ts.version
            ts.mu.RUnlock()
            delta := diffAgainstLastSent(sub, newMap) // nil if
            this stream's view is unchanged
            if delta == nil { continue }
            up := &coordinatorv1.MapUpdate{Version: ver,
                Body: &coordinatorv1.MapUpdate_Delta{Delta:
            delta}}
            select {
            case sub.ch <- up:                //
            enqueue (non-blocking)
            default:                          // queue
            FULL → back-pressure (§6.4)
                c.metrics.backpressureDrops.Inc()
            }
        }
    }
}

```

```

        sub.cancel() // drop
        slow consumer → reconnect-with-snapshot
    }
}
ts.recordDeltaForResume(t) // append
    to deltaRing for known_version (§3.2)
}

```

Back-pressure (the 24 h-soak memory SLO, §7): each stream's `sub.ch` is bounded (`SEND_QUEUE_CAP`, §6.4). A consumer too slow to drain its queue is **dropped** (`cancel()`), not buffered unboundedly — it reconnects and gets a fresh snapshot. This converts "slow consumer" from a memory leak into a bounded reconnect (back-pressure, not growth) [02_CP §6.4].

6. Concurrency model, ordering, and bounds

6.1 Lock hierarchy (deadlock-free by construction)

1. `Coordinator.mu` (RWMutex) — guards the tenants map only; held briefly to fetch a `*tenantState`.
2. `tenantState.mu` (RWMutex) — guards one tenant's graph + version + policy + `deltaRing`.
3. `subRegistry.mu` (RWMutex) — guards the open-stream index.
4. `subscription.mu` — guards `lastVersion`.

Ordering rule: never acquire a higher-numbered lock then a lower one. The event applier holds `tenantState.mu` (write) but enqueues to `sub.ch` **without** holding it (the channel send is non-blocking, §5). The stream handler holds only `tenantState.mu.RLock` for snapshot reads. No path holds two `tenantState` locks (tenants are independent). This is the §11.4 concurrency discipline made mechanical.

6.2 Event ordering guarantee

Within one tenant, events are applied in **Redis-Stream-ID order per stream** (`XReadGroup` returns ordered IDs). Cross-stream ordering (e.g. `events:devices` vs `events:policy`) is NOT totally ordered — but reactions are idempotent and recompute from current state, so a `policy.compiled` that lands before its `device.enrolled` simply produces a smaller delta now and the correct one when the device event applies (eventual convergence, bounded by the SLO).

UNVERIFIED (§11.4.6): strict cross-stream causal ordering is NOT guaranteed by Redis Streams and is not required for correctness here; if a future invariant needs it, a single merged stream or a vector clock is the Phase-2 mechanism — not asserted as present.

6.3 Version monotonicity

`tenantState.version` is bumped only under `tenantState.mu.Lock`, and only forward. For `policy.compiled` the version IS the policy version (§4.3); for non-policy mutations (`device.online`, `prefixes.changed`) the version is incremented by 1 so every wire `MapUpdate` carries a strictly-increasing version an agent can use for `known_version` resume. **UNVERIFIED:** whether non-policy mutations should share the policy-version namespace or use a separate monotonic counter is a doc-03 wire-contract decision; this spec uses a single monotonic `tenantState.version` (incremented for any accepted mutation) as the simplest correct choice and flags the alternative.

6.4 Tunable bounds (named constants — calibrated on the project's own load, §11.4.6)

Constant	Default	Meaning
KEEPALIVE_INTERVAL	20 s	keepalive ticker on each stream [02_CP §6.3]
SEND_QUEUE_CAP	256	per-stream bounded delta queue; overflow $\Rightarrow$ drop+reconnect (§5)
DELTA_RING_CAP	1024	per-tenant recent-delta ring for <code>known_version</code> resume (§3.2)
EVENT_CLAIM_MIN_IDLE	30 s	XAutoClaim reclaim threshold [research-go_cp §5]
EVENT_RETRY_CAP	5	per-event delivery-count ceiling before DLQ [02_CP §5.4]
PRESENCE_REAP_INTERVAL	10 s	heartbeat-TTL reaper cadence (§4.4)
HEARTBEAT_TTL	45 s	online-via-heartbeat expiry (§4.4)

Defaults are starting points; the SLO suite (§7) re-calibrates them on captured load — hardcoding from literature without measuring is the §11.4.6 anti-pattern this note forbids.

7. The < 1 s convergence SLO & metrics

7.1 The reconciliation path the SLO measures

sequenceDiagram

autonumber

participant Pol as policy

participant Bus as Redis Streams

participant Coord as coordinator

participant Edge as Rust edge

Pol->>Bus: XADD events:policy {policy.compiled, v}

Note over Coord: t0 = event receive (handleEnvelope)

Bus-->>Coord: XReadGroup delivers

Coord->>Coord: apply $\rightarrow$ projectionDiff $\rightarrow$ minimal changed set

```

loop each affected open stream
  Coord->>Edge: stream.Send(MapDelta)
end
Note over Coord: t1 = stream.Send returns ; observe(t1-t0)
Edge-->>Edge: reconcile WG peers + verdict map (doc 01)
Note over Pol,Edge: SLO: p99(t1-t0) < 1 s, ZERO restarts (C5)

```

7.2 Metrics (Prometheus, owned by telemetry, fed by coordinator) [02_CP §10.2]

```

helix_reconcile_seconds    histogram // event-receive →
                           stream.Send; p99 < 1 s is the gate
helix_open_streams         gauge       // global aggregate – NO
                           tenant label (C3 cardinality)
helix_fanout_affected_nodes histogram // size of the minimal
                           affected set per event (should be small)
helix_backpressure_drops_total counter // slow-consumer drops (§5)
helix_events_dlq_total     counter    // poison events routed to
                           DLQ (§4.1)
helix_presence_online      gauge       // global aggregate – NO
                           tenant label (C3 cardinality)
helix_graph_nodes          gauge       // global aggregate – NO
                           tenant label (C3 cardinality)
process_resident_memory_bytes gauge    // 24 h-soak slope ≈ 0 (no
                           leak)

```

Reconciled (§11.4.35, 2026-06-26): the /metrics gauges are **global aggregates with no {tenant} label** — a tenant_id label is forbidden on /metrics (unbounded cardinality + leaks tenant population) per [svc-telemetry.md §3.1]'s C3 guard. Per-tenant breakdowns are NOT exposed on /metrics; they live behind the authenticated /v1/stats API [svc-telemetry.md §5.2], never the public scrape endpoint.

7.3 SLO acceptance numbers [02_CP §10.2, 04_P1 §10]

Metric	Target	Measurement
event → delta-on-wire	p99 < 1 s	helix_reconcile_seconds p99
device revoke → stream force-closed	< 1 s	revoke-event → sub.cancel() timer (§4.5)
coordinator memory @ 10k streams	bounded, slope≈0 over 24 h	process_resident_memory_bytes
minimal-affected-set size (small change)	≪ tenant size	helix_fanout_affected_nodes
catch-up reconnect (within ring)	no full snapshot	ring-hit ratio metric

8. Error taxonomy & edge cases

8.1 Error taxonomy (Connect codes + internal)

ID	Connect code	Cause	Handling
E-AUTH	Unauthenticated	mTLS cert invalid/expired	reject stream; agent re-enrolls/ renews cert (§9 [02_CP])
E-REVOKED	PermissionDenied	device revoked	reject at open; force-close if mid-stream (§4.5)
E-IMPERSONATE	PermissionDenied	WatchRequest.device_id ≠ cert identity	reject (anti-spoof)
E-NOTENANT	FailedPrecondition	tenant graph not hydrated	reject; client retries after boot
E-BACKPRESSURE	Aborted	send queue overflow (slow consumer)	close stream → client reconnect-with-snapshot (§5)
E-STREAM-SEND	(propagated)	transport write failed (client gone)	subs.remove + presence maybe-offline
E-APPLY-RETRY	n/a (internal)	transient graph-apply error	leave un-ACKed → XAutoClaim retry (§4.1)
E-APPLY-POISON	n/a (internal)	delivery-count ≥ cap	DLQ + alert; never spin (§4.1)

A streaming response is always HTTP 200 even on error in Connect/gRPC — errors live in the trailer [research-go_cp §2]; the handler returns a `*connect.Error` and the client reads the trailer status, NOT the HTTP code.

8.2 Edge cases (each is a concrete test point, §9)

1. **Reconnect storm:** 10k agents reconnect simultaneously (gateway flap). Each sends a snapshot; bounded by `SEND_QUEUE_CAP` per stream + per-tenant `tenantState.mu.RLock` snapshot reads. No global lock is held during Send.
2. **Revoke-while-streaming:** revoke lands while the device's `WatchNetworkMap` is open ⇒ §4.5 force-close within the SLO; the device's last map is purged from every peer.

3. **Stale known_version (> current):** client claims a future version (forked/rolled-back coordinator) $\Rightarrow$ safe full snapshot (§3.2).
 4. **Ring gap on resume:** known_version older than DELTA_RING_CAP retained $\Rightarrow$ full snapshot, not a partial-and-wrong delta.
 5. **Duplicate/replayed event (at-least-once):** idempotent apply recomputes from current state $\Rightarrow$ diffPeers returns nil $\Rightarrow$ zero spurious sends (§4.1).
 6. **Policy compile out-of-order ($v \leq \text{current}$):** no-op (§4.3) — never regresses the graph.
 7. **Empty/default-deny tenant:** node with no VisibleTo entries $\Rightarrow$ snapshot has zero peers (C4), never fail-open.
 8. **Connector advertising overlapping CIDR (route.conflict.detected):** coordinator does NOT push a delta; the conflict surfaces to Console; 4via6 site disambiguation resolves it (D4) [02_CP §7.3].
 9. **Multi-stream device (reconnect overlap):** fan-out hits all subs; presence offline only on the last close (§3.1/§4.4).
 10. **Coordinator restart mid-stream:** agents' streams break; on reconnect they resume by known_version against the rebuilt-from-Postgres graph (version namespace = policy version, so the version is meaningful across restarts, §1.4).
-

9. Test strategy (anti-bluff, §11.4.169 comprehensive test-type coverage)

Per §11.4.169 (mandatory comprehensive test-type coverage) every coordinator behaviour maps to ≥ 1 concrete test of the appropriate type, each producing captured evidence per §11.4.5/§11.4.69 — no metadata-only PASS. Tests compose with §11.4.27 (no fakes beyond unit), §11.4.85 (stress + chaos), §11.4.50 (determinism), §11.4.107 (liveness — here the "frame" is the delta-on-wire). Integration infra (Postgres + Redis) is booted on-demand via vasic-digital/containers (§11.4.76), never ad-hoc docker run.

Test type (§11.4.169)	Concrete coordinator test point	Captured evidence
Unit	buildMap/diffPeers determinism: same (graph, policy) $\Rightarrow$ byte-identical proto (proto.Equal), N=10 (§11.4.50).	golden proto fixtures + diff hashes
Unit	projectionDiff minimality: a 1-rule change touches only the granted nodes — assert len(changed) exact.	per-case affected-set assertions
Unit	known_version resume: ring-hit $\Rightarrow$ deltas; ring-gap $\Rightarrow$ snapshot; future-version $\Rightarrow$ snapshot.	decision-table fixtures
Integration	enroll $\rightarrow$ advertise $\rightarrow$ policy.compiled $\rightarrow$ WatchNetworkMap: assert exact delta stream content + version monotonicity, real PG+Redis.	captured MapUpdate transcript
Integration	RLS floor: coordinator hydration as helix_app (FORCE RLS) cannot read tenant B (C8).	cross-tenant denial log
E2E / full-automation	full self-host slice: deny-unauthorized peer never appears in any node's snapshot (C4).	recorded peer-set per node

Test type (§11.4.169)	Concrete coordinator test point	Captured evidence
Performance	helix_reconcile_seconds p99 < 1 s @ 10k streams, 1 policy flip.	histogram capture (§7.3)
Soak (stress)	24 h: N agents hold streams, policies flap; assert process_resident_memory_bytes slope ≈ 0 + SLO held (§11.4.85).	24 h memory + p99 time-series
Stress	reconnect storm (10k simultaneous) — no deadlock, bounded queues (§8.2-1).	latency distribution + drop count
Chaos	kill a coordinator consumer mid-apply $\Rightarrow$ PEL reclaim via XAutoClaim, zero lost events (§11.4.147).	before/after event-ID ledger
Chaos	drop Redis mid-stream $\Rightarrow$ presence degrades, graph intact, recovers on reconnect (C2, §4.4).	recovery trace
Security	revoke-while-streaming forces close < 1 s; impersonation (device_id $\neq$ cert) rejected (§4.5/§8.1).	revoke-to-close timer + reject log
Meta-test (§1.1)	mutate diffPeers to always-emit-full-map $\Rightarrow$ helix_fanout_affected_nodes minimality test FAILs; restore $\Rightarrow$ passes.	paired RED/GREEN
Meta-test (§1.1)	mutate buildMap to skip the VisibleTo filter $\Rightarrow$ C4 deny-unauthorized E2E FAILs (need-to-know guard).	paired RED/GREEN

The §1.1 paired mutations are the anti-bluff floor: an analyzer (the test) that passes its own golden-bad fixture is itself a bluff (§11.4.107(10)) — each coordinator gate ships a mutation that **MUST** flip it RED.

10. Phase-2 forward seams (additive, not a rewrite) [02_CP §14, 04_P1 §12]

The coordinator's seams extend without reshaping: (a) `events.Bus` swaps Redis Streams $\rightarrow$ NATS JetStream for multi-region fan-out (the consumer loop is bus-agnostic, §4.1) [02_CP §1.3-D3]; (b) `relayEndpoint` gains direct-P2P NAT-traversal candidates so `Peer.endpoint` stops always relaying through the gateway [02_CP §6.2]; (c) coordinators become **stateless and horizontally scalable** — the graph is rebuilt from Postgres + events on boot (R2 makes this true today), so K8s replicas need no sticky state [02_CP §11.3]; (d) the `version` namespace + `deltaRing` already support cross-restart resume, which generalises to cross-instance resume under a shared bus. Phase 1 drew the seams (R1–R4, the Bus/Compiler/Registry/PKI interfaces) so Phase 2 is additive.

Sources

[02_CP] docs/research/mvp/final/02-control-plane.md — pass-1 control-plane spec (coordinator §6, proto §4, events §5, policy §7, SLO §10, invariants §0.1). · [04_P1] docs/research/mvp/04_VPN_CLD/HelixVPN-Phase1-MVP.md (coordinator, WatchNetworkMap, convergence, SLOs). · [04_ARCH] docs/research/mvp/04_VPN_CLD/HelixVPN-Architecture-Refined.md §1–§7 (control/data separation, push-don't-poll, need-to-know, ULA/48+4via6, no-logging). · [research-go_cp] v09-research/research-go_cp.md (Connect-RPC server-streaming + trailer-error semantics §2, RLS SET LOCAL §4, Redis Streams XAutoClaim/DLQ §5 — latest versions, accessed 2026-06-25). · [SYNTHESIS] v09-research/_SYNTHESIS.md §§1–9 (stack floor, D3/D4 decisions, ecosystem + constitution bindings).