

Disaster recovery (RTO/RPO, backup & region-failover runbook)

Disaster recovery (RTO/RPO, backup & region-failover runbook)

Revision: 1 Last modified: 2026-06-25T12:00:00Z

Master technical specification — Volume 6 (Deployment, Tooling & Operations), nano-detail document `disaster-recovery.md`. **This document CLOSES source-coverage ledger gap G1** (`99-source-coverage-ledger.md` — `[03_ZAI]`'s named DR posture: RTO/RPO budget, KMS-encrypted backups, Terraform region-failover runbook, "disaster recovery as a first-class operational section" — was scattered, never consolidated). Scope: the **RTO/RPO budget** (stated as **TARGETS**, §11.4.6), exactly **what is backed up and what deliberately is not** (Postgres = truth, backed up; Redis/NATS = ephemeral, NOT backed up — and why that is *safe* under the no-logging promise), backup cadence + restore runbook, the step-by-step region-failover runbook, the §9.2 absolute-data-safety discipline, and concrete backup/restore + verification commands. This is a SPEC (describe the procedure; do not run the product). Evidence cited inline: `[05_§N]`, `[svc-coordinator §N]`, `[svc-events §N]`, `[svc-telemetry §N]`, `[ha-and-multiregion §N]`, `[kubernetes §N]`, `[03_ZAI]` (the DR source), `[ledger G1]`. Unproven facts are marked UNVERIFIED per §11.4.6.

Table of contents

- 0. DR thesis — back up truth, regenerate the rest
- 1. RTO / RPO budget (TARGETS, not guarantees)
- 2. What is backed up — and what deliberately is NOT
- 3. Backup cadence & mechanism (Postgres + secrets)
- 4. The §9.2 absolute-data-safety discipline applied to DR

- [5. Restore runbook \(single-store recovery\)](#)
 - [6. Region-failover runbook \(whole-region loss\)](#)
 - [7. Restore verification \(the anti-bluff gate\)](#)
 - [8. DR drills & captured evidence](#)
 - [9. UNVERIFIED register](#)
 - [Sources verified](#)
-

0. DR thesis — back up truth, regenerate the rest

HelixVPN's DR design is shaped by the same architectural decision that makes HA cheap ([ha-and-multiregion §0]): the system has **exactly one durable source of truth — Postgres (C2)** — and *everything else is a regenerable projection of it*. From [svc-events §7.4]:

"the bus is a fast-path accelerator over Postgres truth, never a second source of truth."

So DR reduces to a sharp, auditable statement:

1. **Back up Postgres** (truth) + **the secrets** needed to bring a stack up. That is the entire durable surface.
2. **Do NOT back up Redis/NATS** — they are ephemeral by construction; restoring them would restore *stale* presence/event state, and (critically) the no-logging promise means there is nothing in them worth persisting. The coordinator graph is rebuilt from Postgres on boot [svc-coordinator §1.4]; the outbox re-relays in-flight events [svc-events §7.4].
3. **Regenerate the deploy substrate** (quadlets/compose/k8s) from the in-code spec via `helixvpctl deploy [05 §6.3(B)]` — the manifests are a §11.4.77 build derivative, re-emitted, not restored.

The DR posture is therefore: *one thing to back up (Postgres), one thing to protect (the secrets), everything else regenerated*. This is both simpler and **safer** than a "back up everything" posture — there is no risk of restoring a stale connection log because there is no connection log.

1. RTO / RPO budget (TARGETS, not guarantees)

Per §11.4.6 these are stated explicitly as **TARGETS** the design aims at, to be *measured* in DR drills (§8) — never asserted as achieved. They consolidate [03_ZAI]'s "RTO < 15 min" framing that [ledger G1] found scattered.

Scenario	RTO target (time to service)	RPO target (data loss window)	Basis
helixd replica loss	seconds (reconnect)	0	stateless reconnect [svc-coordinator §3.2]; no data lost
Postgres primary failover (HA standby)	< 60 s (promotion)	0 for sync- committed	sync standby remote_write [ha-and- multiregion §3]
Redis/NATS loss	seconds– minutes (recovery)	0 durable (ephemeral)	rehydrate from PG + outbox re-relay [svc- events §7.4]
Postgres total loss → restore from backup	< 15 min ([03_ZAI] target)	≤ backup- cadence (see §3)	restore latest backup + WAL replay
Whole- region loss → failover to standby region	< 15 min	0 (sync) or ≤ async-lag	promote cross-region standby [ha-and- multiregion §1] + repoint geoDNS [ha- and-multiregion §5]


RPO precision (the honest number, §11.4.6). RPO is 0 for any transaction that committed against a **synchronous** standby ([ha-and-multiregion §3]). RPO for a *backup-only* restore is bounded by **the most recent recoverable point** = last base backup + last shipped WAL segment (with continuous WAL archiving, RPO ≈ the WAL archive interval, typically seconds–minutes; see §3). The 15-min figure is RTO (*time to restore*), distinct from RPO (*data lost*) — conflating them is the bluff this note forbids. The data-plane impact during all of the above is **none for live tunnels** (C1 fail-static, [ha-and-multiregion §3]): DR pauses *control* operations, never drops *forwarding*.

2. What is backed up – and what deliberately is NOT

Asset	Backed up?	Mechanism	Why
Postgres (tenants, devices, connectors, prefixes, policies, certs metadata, audit_events, event_outbox)	YES	base backup + continuous WAL archive (§3)	the only durable truth (C2); losing it loses the system
CA private key / PKI root (doc 04)	YES (separately, stricter)	KMS-encrypted, offline/HSM-class storage	catastrophic if lost (every device cert chains to it); catastrophic if leaked (§4)
Stack secrets (DB password, etc.)	YES (referenced)	external secret store (ESO/ Sealed-Secrets, [kubernetes §2])	needed to bring a stack up; §11.4.10 never plaintext-in-git
Redis stream contents	NO	—	ephemeral (C2); restoring = stale events; outbox re-relays in-flight [svc-events §7.4]
Redis presence keys	NO	—	ephemeral by design — this IS the no-logging mechanism [svc-telemetry §5]; persisting them would <i>create</i> a connection log
NATS JetStream state (Phase 2)	NO (replicated, not backed up)	cluster replication [ha-and-multiregion §4]	same ephemeral logic; HA via replication, not backup
Coordinator in-mem graph	NO	rebuilt on boot [svc-	a projection of PG; not durable

Asset	Backed up?	Mechanism	Why
		coordinator §1.4]	
Deploy manifests (quadlet/compose/ k8s)	NO (regenerated)	helixvpctl deploy [05 §6.3B]	§11.4.77 build derivative — re- emitted from in- code spec

The privacy-positive DR property (the thing that makes "don't back up Redis" *safe, not reckless*). Redis holds presence (who-is-online-now) and in-flight events. Presence is deliberately non-durable — "online/offline state lives only in Redis with a TTL, so the durable store never accumulates a connection log" [svc-telemetry §0/§5]. **Backing Redis up would manufacture exactly the durable connection log the no-logging promise (C3) forbids.** So the DR design's refusal to back up Redis is not a gap — it is the no-logging invariant extended into the backup layer. A DR plan that snapshotted Redis would be a §11.4 / C3 *violation*, not a safety improvement.



3. Backup cadence & mechanism (Postgres + secrets)

Postgres = base backup + continuous WAL archiving (point-in-time recovery, PITR). This is the mechanism that makes RPO $\approx$ WAL-archive-interval rather than $\approx$ backup-interval.

Layer	Cadence (target)	Mechanism	Retention
Base backup	every 24 h	pg_basebackup / CNPG scheduled backup $\rightarrow$ object store	30 days (UNVERIFIED — operator policy)
WAL archive	continuous (seconds)	archive_command $\rightarrow$ object store (S3/GCS)	until next base backup + grace
Logical dump (belt-and-suspenders)	every 24 h	pg_dump --format=custom (portable, cross-version)	7 days
CA / PKI root	on rotation only	KMS-encrypted export to offline storage	indefinite (the root)

```
# deploy/scripts/dr/pg-backup.sh (illustrative – KMS-encrypted, §11.4.10/§9.2)
set -euo pipefail
TS="$(date -u +%Y%m%dT%H%M%SZ)"
DEST="s3://helix-dr-backups/pg/${TS}"

# 1) consistent base backup (CNPG does this via its Backup CR; raw form shown for portability)
pg_basebackup -h helix-pg-rw -U helix_owner -D - -Ft -X fetch -z \
| aws s3 cp - "${DEST}/base.tar.gz" --sse aws:kms --sse-kms-key-id "${HELIX_KMS_KEY}" # [03_ZAI] KMS-encrypted

# 2) portable logical dump (cross-version restore safety net)
pg_dump -h helix-pg-rw -U helix_owner -d helix --format=custom \
| aws s3 cp - "${DEST}/helix.dump" --sse aws:kms --sse-kms-key-id "${HELIX_KMS_KEY}"

# 3) record the recoverable point for the §7 verification gate
echo "{\"ts\":\"${TS}\",\"lsn\":\"$(psql -h helix-pg-rw -U helix_owner -tAc 'SELECT pg_current_wal_lsn()')\"}" \
| aws s3 cp - "${DEST}/MANIFEST.json"
```

WAL archiving is configured on the cluster (CNPG barmanObjectStore, or raw archive_command):

```
# CNPG continuous WAL archive (the RPO≈seconds mechanism)
spec:
  backup:
    barmanObjectStore:
      destinationPath: "s3://helix-dr-backups/pg"
      encryption: aws:kms # [03_ZAI] KMS-encrypted backups
      wal: { compression: gzip }
      data: { compression: gzip }
      retentionPolicy: "30d"
```

Secrets cadence. Stack secrets are referenced from an external store ([kubernetes §2]); the store's own backup is the operator's responsibility. The CA root is the one secret whose loss is *unrecoverable* (every device cert chains to it) — it is exported KMS-encrypted on rotation only and stored offline/HSM-class (doc 04). Its handling is the highest-stakes part of DR and is governed by §9.2 (§4) + §11.4.10.

4. The §9.2 absolute-data-safety discipline applied to DR

Every DR operation is a destructive-or-high-stakes operation; §9.2 (absolute data safety, zero risk) applies in full:

1. **Backup-first, always.** Before *any* restore overwrites a live store, take a fresh backup of the current (even if corrupt) state — a corrupt-but-present DB may still be more recoverable than a botched restore. Hardlinked/cheap snapshot first (§9.2 step 1).
2. **Record metadata.** Capture the pre-restore HEAD/LSN, the target recovery point, the expected post-restore row counts / schema version, and the recoverable-point MANIFEST. json (§3).
3. **Define expected post-op state.** The §7 verification gate's expected values, written down *before* the restore runs.
4. **Run the restore** — never with `--force` that skips integrity checks; never auto-retry a failed restore onto the same volume.
5. **Post-op gate (§7).** Schema present, RLS intact, schemalint green against the restored DB (no-logging runtime signature §11.4.108), row-count sanity, a live WatchNetworkMap snapshot matches expectation. If ANY check fails → do NOT promote the restored DB; investigate.
6. **CA-key handling is operator-gated (§11.4.101/.122).** Restoring or rotating the CA root is irreversible + catastrophic-blast-radius → it is NEVER autonomous; it goes through the operator with the exact commands surfaced (§11.4.66), never decided by the agent.
7. **Audit trail.** Every DR action recorded under docs/qa/<run-id>/ (§11.4.83) + a docs/changelogs/ DR-event section.

No-force-push analogue for DR (§11.4.113 spirit). Just as history-rewrite is never autonomous, a *destructive restore over a live primary* is never autonomous — it is the §11.4.101 block-only-when (irreversible + high-blast-radius) case, surfaced to the operator. The agent prepares the restore (fetches the backup, validates the manifest, stages the commands) but the "overwrite the live primary" step is operator-confirmed.

5. Restore runbook (single-store recovery)

The common case: Postgres is lost/corrupt, the region is intact. Restore from backup + WAL.

```
# deploy/scripts/dr/pg-restore.sh (illustrative – §9.2
disciplined)
set -euo pipefail
SRC="s3://helix-dr-backups/pg/${RESTORE_TS:?set RESTORE_TS}"
TARGET_PITR="${TARGET_PITR:-latest}" # PITR target time, or
'latest' for end-of-WAL

# §9.2 STEP 1 – backup-first: snapshot whatever is currently there
(even if corrupt)
deploy/scripts/dr/pg-backup.sh || echo "WARN: pre-restore backup
of current state failed – record + decide (§9.2)"

# §9.2 STEP 2/3 – record metadata + expected state (written to the
run dir)
mkdir -p "qa-results/dr/${RESTORE_TS}"
aws s3 cp "${SRC}/MANIFEST.json" "qa-results/dr/${RESTORE_TS}/
recoverable-point.json"

# §9.2 STEP 4 – restore to a FRESH volume (never overwrite the
live primary in place; operator-gated promote)
# CNPG: create a Cluster with bootstrap.recovery pointing at the
barman store + recoveryTarget.
# raw form:
aws s3 cp "${SRC}/base.tar.gz" - | tar -xzf - -C "$
{RESTORE_DATA_DIR:?}"
cat > "${RESTORE_DATA_DIR}/recovery.signal" <<'EOF'
EOF
# recovery.conf / postgresql.auto.conf: restore_command from the
WAL archive + recovery_target_time

# CNPG recovery cluster (the recommended form – restores to a NEW
Cluster, not in-place)
apiVersion: postgresql.cnpg.io/v1
kind: Cluster
metadata: { name: helix-pg-restored, namespace: helixvpn }
spec:
  instances: 3
  bootstrap:
```



```


recovery:
  source: helix-pg-backup
  recoveryTarget: { targetTime: "2026-06-25T11:59:00Z" } #
                    PITR; omit for end-of-WAL
externalClusters:
  - name: helix-pg-backup
    barmanObjectStore: { destinationPath: "s3://helix-dr-
                        backups/pg", encryption: aws:kms }

```

After the restored cluster is healthy and **§7-verified**, the operator repoints helix-pg-rw (the DATABASE_URL host, [kubernetes §2]) to it. The coordinators then rehydrate from the restored truth on their next boot/reconnect [svc-coordinator §1.4] — Redis is untouched (it rebuilds itself), so no Redis restore is needed (§2). Live tunnels never dropped (C1).

6. Region-failover runbook (whole-region loss)

The catastrophic case: region A (primary) is gone. Promote region B's cross-region standby and re-home clients. This consolidates [03_ZAI]'s "Terraform-driven region-failover" [ledger G1] with the HA topology [ha-and-multiregion §1].



Step-by-step:

- 1. Confirm region loss, not a partition.** A transient partition must NOT trigger a promote (split-brain risk, [ha-and-multiregion §7]). The fencing/consensus check (§7) gates this — and per §4, the promote itself is operator-confirmed (irreversible + high-blast-radius).
- 2. Promote the cross-region standby** (region B) to primary. RPO = 0 if it was synchronous, ≤ async-lag otherwise [ha-and-multiregion §3, §1].
- 3. Repoint DATABASE_URL** to region B's primary; helixd replicas in B reconnect and rehydrate from the (now-promoted) truth [svc-coordinator §1.4]; the outbox drains any staged events [svc-events §7.4].
- 4. Re-home gateway selection** — drop region A's edge IPs from geoDNS / withdraw its anycast route [ha-and-multiregion §5]; clients resolve to region B and re-handshake (the re-handshake budget is the data-plane ladder's, NOT the < 1 s control SLO — [ha-and-multiregion §6]).

5. **Re-establish HA in the surviving region** — bring up new standbys in region B (and a new region C if available) so the post-failover topology is not single-point-of-failure.
6. **Verify (§7)** at every step; record evidence (§8).

Terraform's role (§3_ZAI). The region's infra (gateway VPS, edges, PG nodes) is Terraform-provisioned [05 §2 deploy/terraform]; a region-failover may include terraform apply against the surviving/standby region to scale it up. The *infra* re-provision is Terraform; the *data* promote is the PG operator; the *traffic* re-home is geoDNS/anycast. Three distinct mechanisms, sequenced above.

7. Restore verification (the anti-bluff gate)

A restore is **not done when pg_restore exits 0** — it is done when the restored truth is proven correct and serving (§11.4.108 runtime signature, §11.4.123 rock-solid proof). The gate, run before any promote:

```
# deploy/scripts/dr/verify-restore.sh (illustrative – the §9.2
step-5 post-op gate)
set -euo pipefail
PG="${RESTORED_PG_URL:?}"

# 1) schema + migration version present
psql "$PG" -tAc "SELECT max(version) FROM
schema_migrations;" # matches expected (§9.2 step 3)

# 2) RLS intact (C8) – helix_app cannot cross tenants
psql "$PG" -tAc "SELECT relrowsecurity AND relforcerowsecurity
FROM pg_class WHERE relname='devices';" # expect t

# 3) no-logging runtime signature (§11.4.108) – schemalint green
against the RESTORED db
go run ./tools/schemalint --dsn
"$PG" # exit 0 = C3 intact
[svc-telemetry §7]

# 4) row-count sanity vs the recoverable-point manifest (no silent
truncation)
psql "$PG" -tAc "SELECT count(*) FROM
tenants;" # ≥ expected
psql "$PG" -tAc "SELECT count(*) FROM devices WHERE revoked_at IS
NULL;"

# 5) live proof: point a throwaway helixd at the restored db, open
a WatchNetworkMap, assert a
# sane snapshot (the device sees its authorized peers – C4) –
booted via containers (§11.4.76)
deploy/scripts/dr/smoke-watch.sh --dsn "$PG" --expect-peers-for
"$KNOWN_DEVICE"
```

Check	Proves
migration version	restored to the expected schema, no partial restore
RLS forced	C8 multi-tenant isolation survived the restore
schemalint green vs restored DB	no-logging C3 intact (a restore can't smuggle a traffic table)
row-count vs manifest	no silent data loss/truncation
live WatchNetworkMap snapshot	the restored truth actually <i>serves</i> a correct network map (the user-visible proof)

A restore that passes (1)–(4) but fails (5) is NOT recovered — the durable bytes are back but the system doesn't work; only the live snapshot proves end-to-end recovery (§11.4.108 RUNTIME → USER-VISIBLE layers).

8. DR drills & captured evidence

DR is unproven until *rehearsed* (§11.4.85 chaos / §11.4.123 rock-solid proof). The drills:

Drill	Procedure	Captured evidence (§11.4.5/§11.4.69)
PG restore drill	restore a base backup + WAL to a fresh cluster; run §7 gate; measure RTO	restore transcript + §7 gate output + measured RTO vs the 15-min target
PITR drill	restore to a target time <i>before</i> a known mutation; assert that mutation is absent	before/after row diff at the PITR boundary
Region-failover drill	tabletop → live: promote standby, repoint, re-home geoDNS; measure end-to-end RTO	the §6 sequence with timestamps; client re-handshake recording
"Redis is NOT restored" assertion	kill Redis, restore PG only, confirm presence/events rebuild without any Redis backup	rebuild trace [svc-events §7.4] — proves the §2 design
No-logging-survives-restore	restore PG, run schemalint against it	green lint = C3 intact; paired mutation planting a traffic table in the backup → restore gate FAILs

The drill IS the spec's proof. The RTO/RPO numbers in §1 are TARGETS (§11.4.6) until a drill *measures* them; a DR plan with no rehearsed restore is a §11.4 bluff (untested recovery is no recovery). Drills are scheduled (e.g. quarterly) and their evidence committed under docs/qa/<run-id>/dr/. The first drill's measured RTO replaces the target as the *stated capability* — never the design estimate (§11.4.6).

9. UNVERIFIED register

#	UNVERIFIED item	Why / status
U1	The 15-min RTO target (§1)	[03_ZAI]'s figure; a TARGET until a §8 drill measures it — never asserted as achieved
U2	Backup retention windows (30d base / 7d dump) (§3)	operator policy, illustrative defaults
U3	KMS provider + key management exact shape (§3/§4)	[03_ZAI] says KMS-encrypted; the specific KMS (AWS/GCP/HSM) is operator-chosen
U4	CA-root recovery mechanism (§2/§4)	doc 04 owns PKI; DR only states it is the unrecoverable-if-lost secret, KMS-encrypted, operator-gated
U5	Cross-region async standby lag → backup-restore RPO (§1)	depends on link/topology; sync standby gives RPO=0, async ≤ measured lag
U6	Terraform region-failover automation depth (§6)	[03_ZAI]/[05] reference Terraform for infra; the exact apply automation is fleet-specific

Sources verified

- 99-source-coverage-ledger.md **G1** (the gap this document closes) + [03_ZAI] (KMS-encrypted automated backups, Terraform-driven region-failover, RTO < 15 min, GitOps — the DR source whose scattered pieces are consolidated here) — [ledger G1].
- 05-repo-layout-tooling-and-helix-ecosystem.md §2 (deploy/terraform, deploy/grafana), §6.3(B) (deploy regeneration via containers/helixvpctl deploy), §7 (substrates) — [05].

- `v03-control-plane/svc-coordinator.md` §1.4 (boot rehydration from Postgres — why DR needs only PG), §3.2 (stateless reconnect) — [svc-coordinator].
- `v03-control-plane/svc-events.md` §5.2 (outbox), §7.4 (cold-boot rehydration / Redis-loss recoverability — the basis for "don't back up Redis") — [svc-events].
- `v03-control-plane/svc-telemetry.md` §0/§5 (ephemeral Redis presence operationalizes no-logging — why backing it up would CREATE a connection log), §7 (schemalint no-logging runtime signature used in the §7 restore gate) — [svc-telemetry].
- `v06-deploy/ha-and-multiregion.md` §1 (multi-region topology), §3 (Postgres HA / sync standby / RPO=0), §4 (NATS replicated not backed up), §5 (gateway re-home), §6 (failover convergence semantics), §7 (split-brain fencing — gating the §6 promote) — [ha-and-multiregion] (sibling).
- `v06-deploy/kubernetes.md` §2 (helix_app DSN + external secret store), §3 (CloudNativePG / Patroni HA + backup CRs) — [kubernetes] (sibling).

Constitution bindings applied: §11.4.44 (revision header), §11.4.6 (no-guessing — RTO/RPO stated as TARGETS, RTO-vs-RPO distinction made explicit, UNVERIFIED register §9), §9.2 (absolute data-safety DR discipline — §4 backup-first/verify/operator-gate), §11.4.10/.30 (KMS-encrypted, never-plaintext secrets), §11.4.101/.122 (CA-root + destructive-restore operator-gated, never autonomous), §11.4.108/.123 (restore verification = runtime signature + rock-solid proof, §7), §11.4.85 (DR drills = chaos rehearsal, §8), §11.4.83 (evidence under docs/qa/). Closes ledger gap G1.