

High availability & multi-region (Phase-2 scale-out)

High availability & multi-region (Phase-2 scale-out)

Revision: 1 Last modified: 2026-06-25T12:00:00Z

Master technical specification — Volume 6 (Deployment, Tooling & Operations), nano-detail document `ha-and-multiregion.md`. Scope: how HelixVPN goes from the single-node Phase-1 deploy to a **horizontally-scalable, multi-region, highly-available** topology — and how the **< 1 s convergence SLO survives a failover**. It deepens the HA hooks already designed into the control plane: stateless coordinators [`svc-coordinator §10`], the bus-agnostic event seam [`svc-events §10`], Postgres-as-truth / Redis-as-ephemeral (C2). It is **additive, not a rewrite** — every Phase-2 mechanism plugs into a seam Phase-1 already drew. This is a SPEC (describe the implementation; do not build the product). Evidence cited inline: [`05 §N`] (`05-repo-layout-tooling-and-helix-ecosystem.md`), [`svc-coordinator §N`], [`svc-events §N`], [`svc-telemetry §N`], [`kubernetes §N`] (sibling), [`research-podman_k8s`]. The disaster-recovery runbook is its own sibling `disaster-recovery.md` (closes ledger gap G1); this document owns the *steady-state* HA topology and the *failover convergence* story. Unproven facts are marked UNVERIFIED per §11.4.6.

Table of contents

- 0. The HA thesis — stateless coordinators make it cheap
- 1. Reference multi-region topology
- 2. Stateless coordinators behind a load balancer
- 3. Postgres HA — Patroni / synchronous replication
- 4. Event bus HA — Redis replication → NATS JetStream

- [5. Gateway selection — anycast / geoDNS](#)
 - [6. Failover semantics — the convergence-during-failover story](#)
 - [7. Split-brain handling](#)
 - [8. SLOs under failover & anti-bluff evidence](#)
 - [9. UNVERIFIED register](#)
 - [Sources verified](#)
-

0. The HA thesis — stateless coordinators make it cheap


HelixVPN's HA story is unusually cheap because of one architectural decision made in Phase 1: the coordinator is **stateless on disk** [svc-coordinator §0/§10(c)]. Its entire runtime state — the per-tenant topology graph + open WatchNetworkMap streams — is a *projection* of Postgres truth (C2) plus the live event stream, rebuilt from scratch on every boot [svc-coordinator §1.4].

Consequences that define this document:

1. **Coordinators are cattle.** Any helixd replica can serve any agent's stream; killing one loses no state; scaling is a Deployment replicas bump + HPA [kubernetes §5].
2. **The hard state is exactly two stores.** Postgres (truth — MUST be HA) and Redis/NATS (ephemeral — loss is recoverable, HA is a latency/availability nicety, not a correctness need, [svc-events §7.4]).
3. **Failover is a stream-reconnect, not a resync.** An agent whose coordinator dies reconnects to another and resumes by known_version against the version-namespace that survives restarts (version namespace == policy-version namespace, [svc-coordinator §1.4/§3.2]).

So HA = (a) make Postgres HA, (b) front stateless coordinators with a load balancer, (c) pick a gateway-selection mechanism, (d) optionally make the bus HA. Each is below.

1. Reference multi-region topology



Design rules of the reference topology:

- **Single Postgres writer (C2).** There is exactly ONE primary at a time (region A here). All helixd replicas — in every region — write to it. Multi-master Postgres is NOT used (it breaks the single-truth invariant). Cross-region write latency is acceptable because control-plane writes are low-rate (enroll/revoke/policy-activate), not packet-path [svc-coordinator §0 C1].
- **Regional read scale (Phase-2 optimization).** Coordinator boot-hydration reads MAY come from a regional read-replica ([kubernetes §3], helix-pg-ro); UNVERIFIED whether the MVP routes hydration to replicas — Phase-1 reads the primary.
- **Per-region edges.** Each region runs its own helix-edge DaemonSet so a client connects to the nearest gateway; the coordinator's MapDelta carries the right GatewayInfo.endpoint per node [svc-coordinator §2.1].
- **One logical event bus.** Phase-2 NATS JetStream clusters across regions behind the same events.Bus interface (D3, §4).

2. Stateless coordinators behind a load balancer


The control-plane API (helixd:8443, Connect-RPC + REST) sits behind a regional L7 LB / Ingress. Because coordinators are stateless [svc-coordinator §10(c)]:

- **No sticky sessions required.** An agent's long-lived WatchNetworkMap stream is pinned to one pod for the life of the TCP/HTTP-2 connection (normal LB behaviour for a streaming RPC), but on reconnect it may land on any pod and resume by known_version [svc-coordinator §3.2].
- **Each replica is a full consumer.** Every helixd joins the event-bus consumer groups with a distinct consumer name (pod name = hostID, [svc-events §3.1]); work is partitioned across replicas by the bus (Redis Streams / NATS), not duplicated. The outbox relay uses FOR

UPDATE SKIP LOCKED so every replica safely runs the relay loop [svc-events §5.2].

- **Fan-out reaches the right replica.** An event affecting device D is reacted by whichever replica owns D's open stream — but ALL replicas consume the same streams and recompute, and only the replica holding D's WatchNetworkMap subscription has a subRegistry entry to send to, so a replica with no affected streams simply produces a zero-send no-op [svc-coordinator §5].

UNVERIFIED — cross-replica fan-out. In the single-binary MVP, one coordinator owns every stream, so "the replica that consumed the event also holds the affected stream" is trivially true. In multi-replica Phase 2, an event consumed by replica X may affect a stream held by replica Y. The seam that makes this work is the **shared bus**: every replica consumes every stream and recomputes its *own* held streams [svc-coordinator §10(d)]. This is asserted as the design intent, not as validated behaviour — it is flagged for the Phase-2 multi-replica integration test (§8) and MUST NOT be claimed working until that test captures evidence.




3. Postgres HA — Patroni / synchronous replication

Postgres is the only store whose HA is a *correctness* requirement (C2). The recommended posture (from [kubernetes §3] D-K8S-PG-HA option a, deepened here):

- **Topology:** 1 primary + ≥ 2 standbys. At least one standby is **synchronous** (`synchronous_commit = remote_write` or stricter) so an acknowledged control-plane write survives a primary loss with **RPO = 0** for committed transactions; additional async standbys (possibly cross-region) serve read scale + DR [03_ZAI] (DR detail in disaster-recovery.md §3).
- **Failover:** Patroni / CloudNativePG promotes a healthy standby on primary failure, fences the old primary, and re-points the `helix-pg-rw` Service to the new primary [kubernetes §3].

- **What helixd sees:** during the promotion window, writes fail (the rw endpoint is briefly unavailable). This is **fail-static-correct** — existing tunnels keep forwarding (C1, the edge has the WG peers programmed; the Go control plane is not in the packet path), and /readyz goes NOT-READY so the LB stops sending new API traffic until the new primary is up [svc-telemetry §8.2].
- **Recovery:** when the new primary is rw, coordinators reconnect, the outbox relay drains any staged-but-unrelayed events, and the graph (already in RAM, never lost) keeps serving — only *new* mutations were paused, never the live topology [svc-events §7.4].



The HA payoff stated plainly: a Postgres failover pauses *control* operations (you can't enroll/revoke/activate-policy for a few seconds) but NEVER drops *data-plane* traffic, because the Go control plane is structurally out of the packet path (C1). This is the single most important HA property and it is free — it falls out of the Phase-1 control/data separation.

4. Event bus HA — Redis replication → NATS JetStream

The bus is ephemeral (C2): losing it loses no durable state [svc-events §7.4], so its HA is about *availability* + *latency*, not correctness. Two stages:

Phase-1 (Redis): a Redis replica (StatefulSet replicas or Redis Sentinel) reduces the recovery window, but even a total Redis loss is recoverable — the coordinator rehydrates from Postgres and the outbox re-relays [svc-events §7.4]. HA Redis is OPTIONAL in Phase 1.

Phase-2 (NATS JetStream): the D3 swap [svc-events §10]. NATS JetStream gives clustered, replicated, durable streams across regions behind the **same events.Bus interface** — "the consumer loop is bus-agnostic" [svc-coordinator §10(a)], so this is a transport change, not a rewrite. JetStream consumer groups map onto the same (stream, group, consumer=hostID) topology [svc-events §3.1].

```
// internal/events/natsbus.go – Phase-2 impl behind the SAME
// events.Bus interface (§2.1 [svc-events])
// NOTHING above this interface knows Redis vs NATS is live [svc-
// events §10].
type NatsBus struct {
    js      nats.JetStreamContext
    hostID  string
    // Publish → js.Publish to the stream subject; Subscribe →
    // durable pull consumer per group;
    // Ack/Nack → msg.Ack()/msg.Nak(); XAUTOCLAIM reclaim →
    // JetStream redelivery + MaxDeliver (DLQ).
}
```

UNVERIFIED — NATS JetStream multi-region replication factor.


The exact JetStream cluster sizing (R3/R5, per-region placement, mirror vs cluster) is a Phase-2 capacity decision, not resolved here. What IS asserted: the Bus interface is the seam, the envelope/taxonomy are transport-neutral [svc-events §2.2], and the at-least-once + idempotent-reaction contract [svc-events §6.2] holds identically on JetStream (whose redelivery + MaxDeliver→DLQ map onto the Redis XAUTOCLAIM→DLQ machinery). The replication topology is flagged for Phase-2 design.

5. Gateway selection — anycast / geoDNS

A client must reach the **nearest healthy** gateway (edge). Two mechanisms, recommendation explicit (§11.4.66):

Decision D-GW-SELECT (options + recommendation). (a) **geoDNS** — gw.helix.example resolves to the nearest region's edge IP(s) by client geo; health-checked DNS removes a dead region from rotation. Simplest; DNS-TTL-bound failover. (b) **Anycast** — a single advertised IP announced (BGP) from every region's edge; routing delivers the client to the topologically-nearest PoP; a dead PoP withdraws the route. Fastest failover, requires BGP/anycast capability. (c) **Client-side from the network map** — GatewayInfo.endpoint in the MapDelta [svc-coordinator §2.1] carries a per-node gateway; the client could be handed a region-specific endpoint by the coordinator at enroll time. **Recommended:** (a) **geoDNS** for the Phase-2 MVP (no BGP dependency); (b) **anycast** for latency-critical fleets that have the network capability; (c) as a complement — the coordinator already knows each node's region, so it can bias GatewayInfo regardless of (a)/(b).

The coordinator's role is consistent across all three: it sets `GatewayInfo.endpoint/MasqueSNI` per node [svc-coordinator §1.2/§2.1], and on a regional gateway failover it emits `gateway.failover {from,to}` [svc-events §3.4], which recomputes every affected node's `GatewayInfo` and pushes a `GatewayInfo-only delta` [svc-coordinator §4.2].




6. Failover semantics — the convergence-during-failover story

This is the heart of the document: **what does the < 1 s convergence SLO mean while something is failing over?** The answer differs by what fails, and each case is governed by an existing invariant — none is invented here.

Failure	Data-plane effect	Control-plane effect	Convergence semantics
One helix replica dies	none (C1 — not in packet path)	agents on that pod's streams reconnect to another replica	resume by <code>known_version</code> ; deltas, not full snapshot [svc-coordinator §3.2]; SLO clock restarts on reconnect, p99<1s for subsequent events
Postgres primary failover (§3)	none (tunnels keep forwarding)	writes paused seconds; /readyz NOT-READY	<i>new</i> mutations queue/fail-and-retry; the already-pushed topology is intact in every coordinator's RAM; convergence resumes once the new primary is <i>rw</i>
Redis/NATS blip (§4)	none	event delivery paused; presence goes stale	outbox stages writes durably (C2); on bus recovery the relay drains and convergence catches up [svc-events §7.4]; presence

Failure	Data-plane effect	Control-plane effect	Convergence semantics
			reaper reconciles [svc-telemetry §5.4]
Regional gateway (edge) loss (§5)	clients on that gateway drop; reconnect to next-nearest	coordinator emits <code>gateway.failover</code> ; re-points <code>GatewayInfo</code>	clients re-handshake against the new gateway; the WG/MASQUE re-handshake is the client's reconnection state machine (data-plane spec), bounded by the transport ladder, NOT by the <1s control SLO
Whole region loss	clients fail over to surviving region's gateways	surviving-region coordinators keep serving from RAM + the (cross-region) primary or its promoted standby	combines the rows above; DR runbook in <code>disaster-recovery.md</code> §6

The crucial honesty (§11.4.6): the < 1 s SLO measures **event-receive** → **MapDelta-on-wire** at a *healthy* coordinator [svc-coordinator §7.1, svc-telemetry §6.1]. It is NOT a promise that a client mid-failover re-handshakes its tunnel in < 1 s — that is the data-plane reconnection ladder's budget, a different and larger number. Conflating the two would be a bluff. What the HA design DOES guarantee: a failover **pauses** convergence (briefly) and never **corrupts** it — when the failed component returns, the graph is rebuilt-or-intact and deltas resume; no node ever converges to a *wrong* map because of a failover (idempotent recompute-from-state, [svc-events §6.2]).



7. Split-brain handling

Split-brain is the classic HA hazard. HelixVPN's exposure is bounded by the single-writer invariant:

- **Postgres split-brain (two primaries).** Prevented by the HA operator's fencing: Patroni/CNPG uses a distributed-consensus lock (etcd/ Kubernetes lease) so at most one node holds the leader key; a partitioned old primary is demoted/fenced before a standby is promoted [research-podman_k8s]. If fencing genuinely fails (UNVERIFIED edge case), the sync-standby's `remote_write` requirement means an unfenced old primary cannot acknowledge writes the new primary doesn't have — the damage is bounded, not silent.
- **Coordinator "split-brain" is a non-issue.** Two `helixd` replicas serving the same tenant is the *normal* steady state, not a fault — they are stateless, both read the same Postgres truth, both consume the same bus, and both recompute idempotently [svc-events §6.2]. The presence layer is explicitly split-brain-safe: "two `helixd` ... both write same key — idempotent (SET is last-writer-wins; SADD is a set) — harmless" [svc-telemetry §5.4].
- **Bus split-brain (NATS partition).** JetStream's RAFT-based clustering elects a single stream leader per partition; a minority partition cannot accept publishes (loses availability, not consistency). Outbox staging means a publish that can't land is retried, not lost [svc-events §5.2]. Exact behaviour is **UNVERIFIED** pending the Phase-2 NATS topology (§4).

The unifying principle: **every stateful component has exactly one writer/leader at a time, chosen by a consensus lock; the stateless components are safe to run N-wide by construction.** Split-brain is therefore a *bounded-availability* event, never a *silent-corruption* event.

8. SLOs under failover & anti-bluff evidence

SLO / claim	Target	Captured-evidence proof (§11.4.5/§11.4.69/§11.4.85)
helixd replica loss → agent resumes	deltas not snapshot; p99<1s post- reconnect	chaos test: kill a replica mid- stream, capture the reconnect known_version resume transcript [svc-coordinator §8.2–10]
PG primary failover → no tunnel drop	data plane unaffected (C1)	chaos test: trigger CNPG switchover under load; packet capture shows tunnels keep forwarding; /readyz flaps, / healthz stays green [svc- telemetry §8.2]
bus loss → no lost events	RPO=0 for committed mutations	chaos test: kill Redis/NATS mid- flight; assert outbox re-relays every staged event on recovery; before/after event-ID ledger [svc- events §7.4, §11.4.147]
multi-replica cross-stream fan- out (§2, UNVERIFIED)	event on X reaches stream on Y	Phase-2 integration test: 2 replicas, event consumed by X affects a stream held by Y → Y sends the delta; captured MapUpdate transcript
gateway failover → clients re-home	clients reach surviving gateway	chaos test: drop a region's edge; geoDNS health-check drops it; clients re-handshake on the next region's gateway; recorded handshake
split-brain presence safe (§7)	two helixd, same key, no corruption	unit/integration: concurrent MarkOnline from two writers → consistent presence set [svc- telemetry §5.4]

Integration infra (PG + Redis/NATS, multi-replica) booted on-demand via containers (§11.4.76), never ad-hoc. Chaos tests are mandatory per §11.4.85 (process-death, network-fault, state-corruption injection with categorised recovery). A "HA works" claim with no failover chaos evidence is a §11.4 bluff.

9. UNVERIFIED register

#	UNVERIFIED item	Why / status
U1	Multi-replica cross-stream fan-out (§2)	trivially true single-binary; Phase-2 multi-replica behaviour is design-intent, flagged for the §8 test, NOT claimed working
U2	Hydration reads routed to regional read-replicas (§1/§3)	Phase-2 optimization; Phase-1 reads the primary
U3	NATS JetStream replication factor / multi-region placement (§4)	Phase-2 capacity decision; only the Bus-seam neutrality is asserted
U4	geoDNS vs anycast final pick (§5)	D-GW-SELECT options; recommendation given, not validated against a real fleet
U5	Postgres fencing under a genuine partition (§7)	bounded by sync-standby remote_write; exact operator fencing behaviour is operator-specific
U6	Whether the <1s SLO and the client re-handshake budget are ever conflated in tooling (§6)	the doc asserts they are distinct; a test must prove the tooling reports them separately

Sources verified

- 05-repo-layout-tooling-and-helix-ecosystem.md §7.3 (K8s edge/coordinator/PG shapes the HA topology refines), §7.4 (multi-region-fleet → K8s recommendation), §12 (decision discipline) — [05].
- v03-control-plane/svc-coordinator.md §0 (C1/C2/C5 + stateless-on-disk), §1.4 (boot hydration, version-namespace survives restart), §2.1 (GatewayInfo/relayEndpoint), §3.2 (known_version resume), §4.2 (gateway.failover reaction), §5 (fan-out), §7.1 (the SLO measurement point), §8.2-10 (coordinator-restart-mid-stream edge case), §10 (Phase-2 stateless/bus-agnostic seams) — [svc-coordinator].
- v03-control-plane/svc-events.md §2.2 (transport-neutral envelope), §3.1 (consumer-group partitioning per replica), §5.2 (SKIP LOCKED relay safe on every replica), §6.2 (idempotent recompute-from-state), §7.4 (cold-boot rehydration / Redis-loss correctness), §10 (NATS JetStream swap behind Bus) — [svc-events].
- v03-control-plane/svc-telemetry.md §5.4 (presence split-brain safety, drift reconciliation), §6.1 (convergence measured at a healthy

coordinator), §8.2 (liveness/readiness fail-static semantics) — [svc-telemetry].

- v06-deploy/kubernetes.md §3 (D-K8S-PG-HA / CloudNativePG), §5 (stateless Deployment + HPA), §6 (edge DaemonSet) — [kubernetes] (sibling).
- 99-source-coverage-ledger.md G1 + [03_ZAI] (KMS backups, RTO, Terraform DR — the multi-region resilience framing this document operationalizes; DR runbook lives in disaster-recovery.md).
- [research-podman_k8s] — Patroni/CNPG failover + consensus-lock fencing, NATS JetStream clustering, UDP gateway ingress per region; cross-referenced against latest official docs per §11.4.99.

Constitution bindings applied: §11.4.44 (revision header), §11.4.6 (no-guessing — the <1s-SLO vs re-handshake-budget distinction made explicit, UNVERIFIED register in §9), §11.4.66 (D-GW-SELECT, D-K8S-PG-HA referenced as options+recommendation), §11.4.85 (failover chaos evidence mandatory in §8), §11.4.147 (no-work-loss outbox re-relay), §11.4.76 (containers submodule for HA integration infra).