

Kubernetes deployment (fleet / Phase-2 HA substrate)

Kubernetes deployment (fleet / Phase-2 HA substrate)

Revision: 1 Last modified: 2026-06-25T12:00:00Z

Master technical specification — Volume 6 (Deployment, Tooling & Operations), nano-detail document `kubernetes.md`. Scope: the **Kubernetes substrate** for HelixVPN — plain manifests (Helm-less, kustomize-overlaid) that run the *same* OCI image set as the canonical Podman-quadlet substrate ([05 §7], 05-repo-layout-tooling-and-helix-ecosystem.md). It deepens [05 §7.3] (the K8s sketch + D-K8S-EDGE-INGRESS) to implementation-ready nano-detail: a stateless helixd Deployment + HPA, a Patroni/StatefulSet Postgres, a Redis (→ NATS JetStream Phase-2) backing store, a privileged helix-edge DaemonSet, the NetworkPolicy default-deny mesh, ConfigMap/Secret wiring, and the load-bearing **gateway-needs-kernel-net** caveat. K8s is an **equivalence target**, not the canonical substrate (Podman quadlets are canonical per §11.4.76/.161); it is recommended only for multi-region fleets [05 §7.4]. This is a SPEC (describe the implementation; do not build the product). Evidence cited inline: [05 §N] overview, [svc-coordinator §N], [svc-telemetry §N], [svc-events §N], [research-podman_k8s] (the Podman-Quadlet / Kubernetes-for-WireGuard deep-research thread, verified against latest official sources per §11.4.99). Unproven facts are marked UNVERIFIED per §11.4.6 — never fabricated.

Table of contents

- 0. Why K8s exists in this spec, and what it must NOT regress
- 1. Namespace, topology & the one image set
- 2. ConfigMap & Secret wiring (the env contract)
- 3. Postgres — StatefulSet / Patroni (the only durable truth)
- 4. Redis / NATS — the ephemeral event + presence backing store
- 5. helixd — the stateless control-plane Deployment + HPA
- 6. helix-edge — the privileged data-plane DaemonSet (the kernel-net caveat)
- 7. Services, ingress & the UDP problem
- 8. NetworkPolicy default-deny mesh
- 9. Probes, PodDisruptionBudgets & rollout discipline

- 10. Generation, parity & anti-bluff evidence
- 11. UNVERIFIED register
- Sources verified

0. Why K8s exists in this spec, and what it must NOT regress

The K8s substrate exists for **one** operator profile: the multi-region fleet that already runs Kubernetes and wants HelixVPN coordinators to scale horizontally behind a load balancer [05 §7.4]. It is NOT the canonical path (Podman quadlets are, §11.4.76/.161) and NOT recommended for homelab/single-VPS buyers. The non-negotiable invariants K8s must preserve, inherited from the control-plane spec [svc-coordinator §0]:

#	Invariant	K8s consequence
C1	Go never in the packet path (fail-static)	the edge (Rust) carries packets, not helixd; a helixd rollout never drops a live tunnel
C2	Postgres = truth; Redis = ephemeral	Postgres is the only StatefulSet whose PersistentVolume is load-bearing; Redis/NATS PVs are disposable
C5	push-don't-poll, convergence p99 < 1 s	helixd replicas are stateless — each rebuilds its graph from Postgres + the event bus on boot [svc-coordinator §10(c)], so K8s may scale/restart them freely
C8	RLS at the DB	helixd connects as the non-superuser helix_app role (FORCE RLS), never the owner

The load-bearing fact that makes K8s viable: the coordinator is *stateless on disk* [svc-coordinator §0] — "the graph is rebuilt from Postgres + events on boot (R2 makes this true today), so K8s replicas need no sticky state" [svc-coordinator §10(c)]. This is why helixd is a Deployment (cattle), Postgres is a StatefulSet (pet), and HPA is safe.

1. Namespace, topology & the one image set

All three substrates run the same OCI images [05 §7]: helixd (control plane), helix-edge (data plane), stock postgres:16, stock redis:7 (→ nats:2 JetStream in Phase 2). Images are tagged <PREFIX>-<version> to match the release-prefix mandate (helix_vpn-..., §11.4.151, [05 §9]).

flowchart TB

```

subgraph ns[namespace: helixvpn]
  subgraph cp[control plane – cattle]
    D[Deployment helixd<br/>replicas≥2 + HPA<br/>stateless C5]
  end
  subgraph dp[data plane – node-local]
```

```

    DS[DaemonSet helix-edge<br>hostNetwork + NET_ADMIN<br>UDP/
443 + UDP/51820]
    end
    subgraph state[durable + ephemeral state]
        PG[(StatefulSet helix-pg<br>Patroni HA - TRUTH C2)]
        RD[(StatefulSet helix-redis<br>ephemeral C2 → NATS P2)]
    end
    CM[ConfigMap helixd-config]
    SE[Secret helix-db / helix-redis]
end
D -- helix_app RLS --> PG
D -- streams+presence --> RD
DS -. reconciles WG peers from MapDelta .-> D
CM --> D
SE --> D
SE --> PG

```

kubectl apply -k deploy/k8s/overlays/<env> is the entry point; the base lives at deploy/k8s/ and is generated by helixvpncctl deploy kube through the containers submodule primitives so the three substrates derive from one in-code spec, not three hand-maintained files (§11.4.81 cross-platform-parity, [05 §6.3(B)]).

2. ConfigMap & Secret wiring (the env contract)

The env contract is identical across substrates [05 §7.1/§7.2]: DATABASE_URL (non-superuser helix_app → RLS enforced, [svc-coordinator §0]), REDIS_URL, and a podman/k8s Secret for the DB password. No secret is ever inlined in a manifest (§11.4.10).

```

# deploy/k8s/base/configmap.yaml
apiVersion: v1
kind: ConfigMap
metadata: { name: helixd-config, namespace: helixvpn }
data:
  HELIX_RELEASE_PREFIX: "helix_vpn" # §11.4.151 -
    matches image tag
  DATABASE_URL: "postgres://helix_app@helix-pg-rw:5432/helix?
    sslmode=require"
  REDIS_URL: "redis://helix-redis:6379"
  METRICS_BIND: ":9090" # /metrics on
    the internal port, NOT 443 [svc-telemetry §8.3]
  LOG_LEVEL: "info"

```

```

---
# deploy/k8s/base/secret.example.yaml (RENDERED from a k8s
    Secret backend, never committed plaintext §11.4.10/.30)
apiVersion: v1
kind: Secret
metadata: { name: helix-db, namespace: helixvpn }
type: Opaque

```

```
stringData:
  password: "<from external secret store – Sealed Secrets /
    External Secrets Operator>"
```

Secret sourcing (UNVERIFIED choice). The MVP does NOT mandate a specific secret backend. Recommended: External Secrets Operator or Sealed Secrets so the deploy/k8s tree carries *references*, never plaintext (the §11.4.77 regeneration mechanism = "re-pull from the cluster secret store"). The literal stringData above is an example shape, not a committed value.

helix_app is the **non-superuser** request-path role with FORCE ROW LEVEL SECURITY [svc-telemetry §4.1, data-model-ddl]; the owner role helix_owner runs migrations only and is never the DATABASE_URL principal for helixd.

3. Postgres — StatefulSet / Patroni (the only durable truth)

Postgres is the single source of truth (C2). In a real fleet it MUST be HA. Two options, recommendation explicit per §11.4.66:

Decision D-K8S-PG-HA (options + recommendation). (a) **Patroni operator** (e.g. the Postgres Operator / CloudNativePG) managing a primary + ≥ 2 replicas with automatic failover + streaming replication. (b) **External managed Postgres** (RDS/Cloud SQL/Crunchy) referenced by DATABASE_URL, Postgres entirely outside the cluster. (c) **Single-replica StatefulSet** (the [05 §7.3] sketch) — DEV/SINGLE-REGION ONLY, no HA. **Recommended:** (a) **Patroni/CloudNativePG** for self-managed fleets (keeps the failover story in-cluster, see ha-and-multiregion.md §3); (b) for teams that already run managed PG. (c) is explicitly NOT production-grade — it is the single-node parity build only.

```
# deploy/k8s/base/pg-cnpg.yaml – CloudNativePG cluster
  (recommended HA, option a)
apiVersion: postgresql.cnpg.io/v1
kind: Cluster
metadata: { name: helix-pg, namespace: helixvpn }
spec:
  instances: 3                                # 1 primary + 2
    hot-standby (sync replication)
  storage: { size: 20Gi }
  postgresql:
    parameters:
      max_connections: "200"
      wal_level: "replica"
    pg_hba:
      - "hostssl helix all all scram-sha-256"      #
        sslmode=require honoured (DATABASE_URL §2)
  bootstrap:
    initdb:
      database: helix
```

```

owner: helix_owner                                # owner runs
migrations
postInitSQL:
- "CREATE ROLE helix_app LOGIN;"                  # non-superuser
request role (RLS floor C8)
- "ALTER DATABASE helix SET row_security = on;"
managed:
services:
additional:
- { selectorType: rw, serviceTemplate: { metadata: {
name: helix-pg-rw } } } # primary
- { selectorType: ro, serviceTemplate: { metadata: {
name: helix-pg-ro } } } # read replicas

```

The helix-pg-rw Service (primary) is the DATABASE_URL host (§2). helixd only ever needs the read-write endpoint (all control-plane writes go to the primary; the coordinator's reads are boot-time hydration [svc-coordinator §1.4], small enough to take from the primary in the MVP). Read-replica routing of hydration reads is a Phase-2 optimization (ha-and-multiregion.md §3.4), not asserted here.

The plain single-replica StatefulSet form (option c) is retained verbatim from [05 §7.3] for the single-region parity build:

```

# deploy/k8s/overlays/single-region/pg-statefulset.yaml (option c
- NO HA, dev/parity only)
apiVersion: apps/v1
kind: StatefulSet
metadata: { name: helix-pg, namespace: helixvpn }
spec:
  serviceName: helix-pg
  replicas: 1                                     # → Patroni/
  CNPG for real fleets (D-K8S-PG-HA)
  volumeClaimTemplates:
  - metadata: { name: data }
    spec: { accessModes: ["ReadWriteOnce"], resources: {
requests: { storage: 20Gi } } }
  template:
    spec:
      containers:
      - name: postgres
        image: docker.io/library/postgres:16
        env:
        - { name: POSTGRES_USER, value: helix_owner }
        - { name: POSTGRES_DB, value: helix }
        - name: POSTGRES_PASSWORD
          valueFrom: { secretKeyRef: { name: helix-db, key:
password } }
        volumeMounts: [{ name: data, mountPath: /var/lib/
postgresql/data }]

```

4. Redis / NATS — the ephemeral event + presence backing store

Redis carries the event streams [svc-events §4] and the TTL presence keys [svc-telemetry §5]. **Losing it loses no durable state (C2)** — on recovery the coordinator rehydrates from Postgres and the outbox re-relays [svc-events §7.4]. Therefore its PV is disposable; a single replica is acceptable in the MVP, with the Phase-2 swap to NATS JetStream behind the same events. Bus interface (D3, [svc-events §10]).

```
# deploy/k8s/base/redis.yaml — ephemeral C2 backing store
(Phase-1)
apiVersion: apps/v1
kind: StatefulSet
metadata: { name: helix-redis, namespace: helixvpn }
spec:
  serviceName: helix-redis
  replicas: 1 # ephemeral C2
    - loss is recoverable, no HA needed in P1
  template:
    spec:
      containers:
        - name: redis
          image: docker.io/library/redis:7
          args: ["--notify-keyspace-events", "Ex"] # presence
            expiry → device.offline [svc-telemetry §5.3]
          ports: [{ containerPort: 6379 }]
          securityContext: { runAsNonRoot: true,
            readOnlyRootFilesystem: true, allowPrivilegeEscalation:
            false }
      volumeClaimTemplates:
        - metadata: { name: data }
          spec: { accessModes: ["ReadWriteOnce"], resources: {
            requests: { storage: 2Gi } } }
```

notify-keyspace-events Ex is mandatory, not optional: the presence reaper subscribes to `__keyevent@<db>__:expired` so an expired presence key still emits `device.offline` [svc-telemetry §5.3]; without it a lapsed device stays falsely "online" (the §11.4.144 corpus-hole class). The belt-and-suspenders sweeper [svc-telemetry §5.3] covers the documented best-effort-delivery gap of keyspace notifications.

The Phase-2 NATS JetStream form (durable, clustered, multi-region) is specified in `ha-and-multiregion.md §4`; the Bus interface makes it a transport change, not a rewrite [svc-events §10].

5. helixd — the stateless control-plane Deployment + HPA

helixd is a **Deployment** (not a StatefulSet) precisely because the coordinator is stateless-on-disk [svc-coordinator §10(c)]. Each replica joins the event-bus consumer groups with a distinct consumer name (the pod name = hostID, [svc-events §3.1]) so work is partitioned across replicas, not duplicated. The outbox relay uses FOR UPDATE SKIP LOCKED so every replica can run the relay loop safely [svc-events §5.2].

```
# deploy/k8s/base/helixd-deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata: { name: helixd, namespace: helixvpn }
spec:
  replicas: 2                                # stateless
    coordinators → HA [svc-coordinator §10c]
  strategy: { type: RollingUpdate, rollingUpdate: {
    maxUnavailable: 0, maxSurge: 1 } }
  selector: { matchLabels: { app: helixd } }
  template:
    metadata:
      labels: { app: helixd }
      annotations: { prometheus.io/scrape: "true", prometheus.io/
        port: "9090" }
    spec:
      serviceAccountName: helixd
      containers:
        - name: helixd
          image: ghcr.io/helixdevelopment/
            helixd:helix_vpn-1.0    # §11.4.151 prefixed
          envFrom: [{ configMapRef: { name: helixd-config } }]
          env:
            - name: PGPASSWORD
              valueFrom: { secretKeyRef: { name: helix-db, key:
                password } }
            - name: HELIX_HOST_ID          # = pod name →
              unique consumer name [svc-events §3.1]
              valueFrom: { fieldRef: { fieldPath:
                metadata.name } }
          ports:
            - { name: api,      containerPort:
              8443 }    # agent/app API (Connect + REST)
            - { name: metrics, containerPort: 9090 }    # /metrics
          — internal only [svc-telemetry §8.3]
          readinessProbe:                                # /readyz:
            PG + Redis reachable [svc-telemetry §8.2]
            httpGet: { path: /readyz, port: 9090 }
            periodSeconds: 5

          livenessProbe:                                # /healthz:
            dependency-FREE (DB blip ≠ restart) [svc-telemetry §8.2]
            httpGet: { path: /healthz, port: 9090 }
            periodSeconds: 10
```

```

resources:
  requests: { cpu: "250m", memory: "256Mi" }
  limits:   { cpu: "2",    memory: "1Gi" }    # 24h-soak
slope≈0 [svc-coordinator §7.3]
securityContext:
  readOnlyRootFilesystem: true
  runAsNonRoot: true
  allowPrivilegeEscalation: false
  capabilities: { drop: ["ALL"] }           # control
plane needs ZERO capabilities (C1)

# deploy/k8s/base/helixd-hpa.yaml
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata: { name: helixd, namespace: helixvpn }
spec:
  scaleTargetRef: { apiVersion: apps/v1, kind: Deployment, name:
                    helixd }
  minReplicas: 2
  maxReplicas: 10
  metrics:
    - type: Resource
      resource: { name: cpu, target: { type: Utilization,
                                       averageUtilization: 70 } }
    # Custom-metric option (UNVERIFIED – needs Prometheus
      Adapter): scale on open WatchNetworkMap streams.
    # - type: Pods
    #   pods: { metric: { name: helix_open_watch_streams },
      target: { type: AverageValue, averageValue: "5000" } }

```

Readiness vs liveness split is load-bearing [svc-telemetry §8.2]: /healthz MUST be dependency-free (a Postgres blip must not trigger a pod-restart loop, fail-static C1); /readyz gates *new* traffic only (removes the pod from the Service when PG/Redis are unreachable, but existing WatchNetworkMap streams keep serving from the in-mem graph). Wiring a DB-dependent liveness probe here would be a correctness regression.

HPA on a custom metric is UNVERIFIED. CPU-based HPA is standard and asserted. Scaling on helix_open_watch_streams [svc-telemetry §3.2] requires the Prometheus Adapter and is left commented — recommended for stream-bound workloads but not validated in this writing context.

6. helix-edge — the privileged data-plane DaemonSet (the kernel-net caveat)

The edge is the one privileged workload. It terminates MASQUE/QUIC on UDP/443 and plain-WG on UDP/51820, and programs **kernel** WireGuard peers from the coordinator's MapDelta. This is the **gateway-needs-kernel-net caveat** that shapes the entire K8s story:

The kernel-net caveat (the reason the edge is special). WireGuard and MASQUE are **UDP** and require kernel network capability — at minimum `NET_ADMIN` to create/configure the wg interface and program AllowedIPs, plus the packet path must reach the node's real UDP socket [05 §7.3, research-podman_k8s]. A normal pod behind a ClusterIP/L7 ingress **cannot** carry this: most cloud L7 ingress can't forward UDP, and an extra NAT hop breaks the WG endpoint model. Therefore the edge runs as a **DaemonSet with hostNetwork: true** so UDP/443 lands directly on the node IP (the node IP *is* the gateway IP), with capabilities: `add: [NET_ADMIN]` and everything else dropped. This is D-K8S-EDGE-INGRESS option (a), recommended for self-host-on-k8s parity with the quadlet model [05 §7.3].

```
# deploy/k8s/base/edge-daemonset.yaml — DATA PLANE (privileged-
minimal)
apiVersion: apps/v1
kind: DaemonSet
metadata: { name: helix-edge, namespace: helixvpn }
spec:
  selector: { matchLabels: { app: helix-edge } }
  template:
    metadata: { labels: { app: helix-edge } }
    spec:
      hostNetwork: true                                # UDP/443 on
                                                         the node IP (kernel-net caveat)
      dnsPolicy: ClusterFirstWithHostNet
      nodeSelector: { helixvpn.io/edge: "true" }        # only nodes
                                                         labelled as gateways run the edge
      tolerations: [{ key: helixvpn.io/edge, operator: Exists,
effect: NoSchedule }]
      containers:
        - name: edge
          image: ghcr.io/helixdevelopment/helix-edge:helix_vpn-1.0
          securityContext:
            readOnlyRootFilesystem: true
            allowPrivilegeEscalation: false
            capabilities: { drop: ["ALL"], add: ["NET_ADMIN"] }
          # ONLY NET_ADMIN (doc 04 §)
          ports:
            - { name: masque, containerPort: 443, protocol:
UDP, hostPort: 443 }
            - { name: wg, containerPort: 51820, protocol:
UDP, hostPort: 51820 }
          env:
            - { name: HELIXD_ADDR, value:
"helixd:8443" }    # edge dials the control-plane Service
for its stream
```

UNVERIFIED — kernel WireGuard module availability. `hostNetwork + NET_ADMIN` lets the edge program WG, but the *kernel* wireguard module must be present on the node (or boringtun userspace fallback used, [SYNTHESIS §2], referenced from helix-core). Whether the target node kernels ship the module is cluster-specific and **MUST** be verified per fleet, never assumed (§11.4.6/§11.4.133). The edge image carries the userspace

fallback so a module-less node still functions, at a performance cost — that fallback path is the data-plane spec's concern, not this document's, and its presence is asserted only as the documented mitigation.

Alternative ingress options (record, not silently dropped) per D-K8S-EDGE-INGRESS [05 §7.3]: **(b)** Service type=LoadBalancer with a **UDP-capable** LB (MetalLB on-prem, NLB on AWS) — for managed-cloud fleets where node-IP-as-gateway-IP is undesirable; **(c)** NodePort UDP for homelab-on-k8s. Recommended: **(a)** for self-host parity, **(b)** for managed fleets.

7. Services, ingress & the UDP problem

```
# deploy/k8s/base/services.yaml
apiVersion: v1
kind: Service                                # control-
      plane API (TCP, ClusterIP – internal)
metadata: { name: helixd, namespace: helixvpn }
spec:
  selector: { app: helixd }
  ports: [{ name: api, port: 8443, targetPort: 8443 }]
  type: ClusterIP
---
apiVersion: v1
kind: Service                                # metrics
      (TCP, internal scrape only [svc-telemetry §8.3])
metadata:
  name: helixd-metrics
  namespace: helixvpn
  labels: { app: helixd }
spec:
  selector: { app: helixd }
  ports: [{ name: metrics, port: 9090 }]
  type: ClusterIP
```

The agent/app API (helixd:8443, TCP, Connect-RPC + REST) is a normal ClusterIP and MAY sit behind an L7 ingress for the Console/REST surface. The **edge UDP ingress does NOT use a Service** in option (a) — it binds the host port directly (§6). This asymmetry — TCP control plane behind ingress, UDP data plane on hostNetwork — is the structural consequence of the kernel-net caveat and is the single most important thing to get right when porting to K8s.

flowchart LR

```
  C[VPN client] -- UDP/443 MASQUE --> N[node IP : hostNetwork edge]
  C2[admin browser] -- TCP/443 https --> ING[L7 Ingress] -->
  SVC[Service helixd:8443] --> POD[helixd pod]
  N -- reconcile from MapDelta --> SVC
```

8. NetworkPolicy default-deny mesh

Zero-trust at the pod mesh (C4 spirit applied to the cluster): default-deny, then allow only the exact flows. NetworkPolicy requires a CNI that enforces it (Calico/Cilium) — **UNVERIFIED** for any given cluster; flagged.

```
# deploy/k8s/base/netpol.yaml
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata: { name: default-deny, namespace: helixvpn }
spec: { podSelector: {}, policyTypes: [Ingress,
    Egress] } # deny all, then allow below
---
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata: { name: helixd-egress, namespace: helixvpn }
spec:
  podSelector: { matchLabels: { app: helixd } }
  policyTypes: [Egress]
  egress:
    - to: [{ podSelector: { matchLabels: { cnpg.io/cluster: helix-
        pg } } }] # → Postgres 5432
      ports: [{ protocol: TCP, port: 5432 }]
    - to: [{ podSelector: { matchLabels: { app: helix-
        redis } } }] # → Redis 6379
      ports: [{ protocol: TCP, port: 6379 }]
  ---
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata: { name: helixd-ingress, namespace: helixvpn }
spec:
  podSelector: { matchLabels: { app: helixd } }
  policyTypes: [Ingress]
  ingress:
    - from: [{ podSelector: { matchLabels: { app: helix-
        edge } } }] # edge dials its stream
      ports: [{ protocol: TCP, port: 8443 }]
    - from: [{ namespaceSelector: { matchLabels: { name:
        monitoring } } }] # Prometheus scrape
      ports: [{ protocol: TCP, port: 9090 }]
```

The helix-edge DaemonSet runs on hostNetwork, so NetworkPolicy does **NOT** govern its UDP/443 ingress (host-network pods bypass pod-network policy) — node-level firewalling (the host's nftables, owned by the data-plane/security spec) governs that path instead. This is a real, UNVERIFIED-in-the-general-case boundary worth calling out: pod NetworkPolicy protects the control plane; host firewalling protects the edge.

9. Probes, PodDisruptionBudgets & rollout discipline

```
# deploy/k8s/base/pdb.yaml
apiVersion: policy/v1
kind: PodDisruptionBudget
metadata: { name: helixd, namespace: helixvpn }
spec: { minAvailable: 1, selector: { matchLabels: { app:
    helixd } } }
```

- **Rollout:** maxUnavailable: 0, maxSurge: 1 (§5) so a helixd rollout never drops below the replica count — combined with the stateless property, agents whose stream lands on a draining pod reconnect by known_version against any other replica and resume with deltas, not a full snapshot [svc-coordinator §3.2/§10(d)]. A rollout is therefore a brief stream-reconnect, never a convergence outage.
 - **PDB** keeps ≥ 1 helixd during voluntary disruptions (node drain).
 - **Postgres** (CNPG) carries its own PDB + switchover semantics from the operator; do NOT hand-roll a PDB that fights the operator (a documented anti-pattern, deferred to ha-and-multiregion.md §3).
 - **Edge** is a DaemonSet: it has no PDB (one per node); draining a gateway node drains its tunnels — that failover is the ha-and-multiregion.md §5 anycast/geoDNS story, not a pod concern.
-

10. Generation, parity & anti-bluff evidence

K8s manifests are **generated**, not hand-maintained, from the same in-code spec as quadlets + compose via helixvpctl deploy kube through containers/pkg/compose [05 §6.3(B)] — one source, three renders (§11.4.81). The parity + correctness claims carry captured-evidence proofs (§11.4.5/§11.4.69, extending [05 §11]):

Claim	Captured-evidence proof
quadlet/compose/k8s run the SAME image & env contract (§1)	three substrates boot the gateway; helixvpctl gateway status GREEN on each; recorded per §11.4.159 [05 §11]
helixd is genuinely stateless (§5)	kill a replica mid-stream → agents resume by known_version on another replica, no full resync [svc-coordinator §8.2-10]; before/after stream transcript
edge UDP/443 reaches the node (§6)	from an external client, MASQUE handshake completes against the node IP; packet capture of the QUIC initial on UDP/443
RLS holds under K8s helix_app (§2)	helixd pod cannot read tenant B's rows even with a crafted query (FORCE RLS) [svc-telemetry §10 T-STORE-1]
NetworkPolicy default-deny enforced (§8)	a pod not in an allow rule cannot reach helixd:8443 (probe denied); paired mutation removing the egress rule → PG connection FAILs

Claim	Captured-evidence proof
no /metrics on public 443 (§7)	scan: :9090 reachable only intra-cluster; :443 carries no Prometheus exposition [svc-telemetry §8.3]

Integration tests that *claim* to exercise the K8s substrate MUST actually stand up a cluster (kind/k3d booted on-demand via the containers submodule, §11.4.76) — a manifest-lint-only PASS that never applies the manifests is a §11.4 bluff (§11.4.38 installable-asset evidence).

11. UNVERIFIED register

Per §11.4.6, every claim not validated in this writing context is enumerated here rather than silently asserted:

#	UNVERIFIED item	Why / mitigation
U1	Custom-metric HPA on helix_open_watch_streams (§5)	needs Prometheus Adapter; CPU HPA is the asserted default
U2	Kernel wireguard module present on edge nodes (§6)	cluster-specific; boringtun userspace fallback is the documented mitigation; verify per fleet (§11.4.133)
U3	CNI enforces NetworkPolicy (§8)	requires Calico/Cilium-class CNI; default kube-proxy alone does not enforce it
U4	UDP LoadBalancer availability (§6 option b)	MetalLB/NLB-dependent; option (a) hostNetwork has no such dependency
U5	External secret backend choice (§2)	not mandated; ESO/Sealed-Secrets recommended; example stringData is illustrative only
U6	CloudNativePG vs Patroni-operator exact pick (§3)	both satisfy D-K8S-PG-HA(a); the manifest shows CNPG as the concrete example, not the only valid operator

None of these is fabricated as "done"; each is a real boundary the operator must close per cluster.

Sources verified

- 05-repo-layout-tooling-and-helix-ecosystem.md §7 (three substrates, one image set), §7.3 (the K8s sketch + D-K8S-EDGE-INGRESS UDP ingress decision), §7.4 (substrate recommendation matrix), §6.3(B) (generation through containers/pkg/compose), §9 (release-prefix on images), §11 (anti-bluff evidence plan) — [05].

- `v03-control-plane/svc-coordinator.md` §0 (C1/C2/C5/C8 invariants), §1.4 (boot hydration), §3.2 (known_version resume), §10(c)/(d) (stateless, K8s-replica-safe forward seam), §7.3 (SLO + memory soak) — [svc-coordinator].
- `v03-control-plane/svc-telemetry.md` §3.2 (helix_open_watch_streams etc.), §5.3 (presence expiry → device.offline, notify-keyspace-events Ex), §8.2 (liveness vs readiness probe contract), §8.3 (/metrics internal-only scrape isolation), §4.1 (helix_app FORCE RLS), §10 T-STORE-1 (RLS denial evidence) — [svc-telemetry].
- `v03-control-plane/svc-events.md` §3.1 (consumer name = hostID, per-replica partitioning), §5.2 (FOR UPDATE SKIP LOCKED relay safe on every replica), §7.4 (cold-boot rehydration from Postgres), §10 (NATS JetStream swap behind the Bus interface) — [svc-events].
- `v03-control-plane/data-model-ddl.md` (the helix_app non-superuser role + FORCE RLS the DATABASE_URL relies on) — referenced for the role contract.
- [research-podman_k8s] — Kubernetes UDP ingress for WireGuard (hostNetwork DaemonSet vs UDP LoadBalancer/MetalLB/NLB vs NodePort), NET_ADMIN/hostPort requirements for in-cluster tunnels, podman kube generate parity, cross-referenced against latest official Kubernetes + Podman docs per §11.4.99.

Constitution bindings applied: §11.4.44 (revision header), §11.4.6 (no-guessing — every unproven fact marked UNVERIFIED + the §11 register), §11.4.66 (decisions as options + recommendation — D-K8S-PG-HA, D-K8S-EDGE-INGRESS), §11.4.76/.161 (containers submodule sole orchestration + rootless ethos — K8s is the equivalence target, quadlets canonical), §11.4.81 (one source → per-substrate render), §11.4.10/.30 (no inlined secrets), §11.4.133 (target-hardware safety — kernel-module verification, never assumed), §11.4.38/.5/.69 (installable-asset + captured-evidence anti-bluff proofs in §10).