

Podman Quadlets (rootless — the canonical substrate)

Podman Quadlets (rootless — the canonical substrate)

Revision: 2 Last modified: 2026-06-26T12:00:00Z

Master technical specification — Volume 6 (Deployment, Tooling & Operations), nano-detail document. [RES] research-backed. Scope: the **canonical rootless Podman quadlet deployment** of HelixVPN — .pod/.container/.network/.volume systemd-generator units, rootless-MANDATORY (§11.4.161 — NO sudo, NO rootful Docker), the NET_ADMIN/NET_RAW capability set for the data- plane edge, :443/udp MASQUE ingress, the one-pod model, read-only rootfs + seccomp confinement, and the vasic-digital/containers submodule as the sole orchestration layer (§11.4.76). This deepens [05-overview §7.1] with the FACT-grade quadlet semantics from [research-podman_k8s]. It is a SPEC: unit files are illustrative of the contract, not the shipping deployment (2–3 refinement passes follow).

Evidence cited inline by id: [05_OV §N] = 05-repo-layout-tooling-and-helix-ecosystem.md; [RPK §N] = v09-research/research-podman_k8s.md (verified against latest Podman docs 2026-06-25 per §11.4.99); [API §N] = v03-control-plane/svc-api.md (the env/readyz contract). Unproven facts are marked **UNVERIFIED** per §11.4.6 — never fabricated.

0. What this document owns (and does not)

This document owns the **rootless quadlet substrate**: the search-path/lifecycle rules, the pod + per-container unit files (edge, control, postgres, redis), the capability/security posture, the :443/udp rootless-bind requirement, the network/volume quadlets, the on-demand-infra invariant, and the rootless-networking caveats that bite a VPN edge.

It does **not** own: the helixvpctl deploy quadlet render command (that is helixvpctl.md); the equivalent Compose stack (docker-compose.md); the K8s manifests (kubernetes.md); the env-contract / readyz semantics (canonical in

[API §4.7]); the edge data-plane internals (Volume 2). Quadlets are **canonical** (§11.4.76/.161); Compose and K8s are *equivalents* generated from the same OCI images [05_OV §7].

1. Why rootless quadlets are canonical (§11.4.161)

Podman in **rootless mode** is the MANDATORY container runtime for ALL HelixVPN containerized workloads (§11.4.161): Docker in rootful mode, sudo, or any escalation to root for container management is **FORBIDDEN** unless the target platform has no rootless option AND that constraint is documented per §11.4.112. The vasic-digital/containers submodule is the **sole** orchestration layer (§11.4.76) — no ad-hoc docker/podman commands outside its pkg/boot/pkg/compose/ pkg/health primitives [05_OV §7.1]. Rootless quadlets fit the primary buyer (homelab / single VPS, spine §2): no Docker daemon, no root, helixvpncctl init one-shot, systemctl --user start lifecycle.

flowchart LR

```
    IMG["OCI images: helixd, helix-edge, postgres:16, redis:7"]
    IMG --> Q["Podman quadlets (CANONICAL §11.4.76/.161 – rootless)"]
    IMG -.equivalence.-> C["Docker Compose (docker-compose.md)"]
    IMG -.equivalence.-> K["Kubernetes (kubernetes.md)"]
    Q --> H["homelab / single VPS – systemctl --user, no daemon, no root"]
```

2. Quadlet fundamentals (Podman 5.x, FACT-grade)

Quadlet is a **core part of Podman** (not a separate project); Podman 5.x added podman quadlet list|print|install|rm and the .pod/.build/.image/.artifact unit types beyond the original .container/.network/.volume/.kube. Quadlet .unit files are read by a systemd generator that emits real .service units at boot / daemon-reload [RPK §1].

2.1 Rootless = file location, NOT User= (the hard constraint)

FACT [RPK §1]. "Quadlet units do not support running as a non-root user by defining the User, Group, or DynamicUser systemd options." Rootless is selected purely by **where the file lives**. The rootless search paths (verbatim upstream):

- \$XDG_RUNTIME_DIR/containers/systemd/
- \$XDG_CONFIG_HOME/containers/systemd/ (i.e. ~/.config/containers/systemd/)
- /etc/containers/systemd/users/\${UID} and /etc/containers/systemd/users/
- /usr/share/containers/systemd/users/\${UID} and /usr/share/containers/systemd/users/

To run rootless with boot-time start (no interactive login), `loginctl enable-linger <user>` is **required** [RPK §1]. HelixVPN places its units in `~/.config/containers/systemd/` and enables `linger` for the gateway service user.

Reconciliation with the overview (§11.4.35). [05_OV §7.1] sketched the units under `/etc/containers/systemd/` — that is the **rootful** system path. Because §11.4.161 mandates **rootless**, the canonical HelixVPN path is the `rootless ~/.config/containers/systemd/` (or `/etc/containers/systemd/users/${UID}` for a system-managed service user). This is a visible correction, recorded so the two docs do not silently disagree; the unit *contents* are unchanged.

2.2 Lifecycle

```
# place units in ~/.config/containers/systemd/, then:
systemctl --user daemon-reload
systemctl --user start helixvpn-pod.service      # the .pod
          generates *-pod.service
loginctl enable-linger "$USER"                  # boot-start
          without interactive login (FACT [RPK §1])
```

The generator guarantees the `*-pod.service` starts **before** any member `Pod=helixvpn.pod` container [RPK §1]; quadlet auto-wires the ordering — no hand-written `After=`/`Requires=` between pod and members.

3. The one-pod model (edge + control + postgres + redis)

Containers in a pod **share the network namespace**, so ports are published at the **pod**, not the member container, and members reach each other on `127.0.0.1` [RPK §1]. HelixVPN groups the four services into one pod for single-node self-host [05_OV §7.1, 04_P1 §9.2].

```
flowchart TD
    subgraph POD["helixvpn.pod (shared netns; 127.0.0.1 between members)"]
        EDGE["helix-edge\nNET_ADMIN+NET_RAW · :443/udp · :51820/\nudp\nread-only rootfs · seccomp"]
        HELIXD["helixd\ncontrol plane · DROP ALL\ncaps\nDATABASE_URL=...@127.0.0.1 · :8443"]
        PG["postgres:16\nDROP ALL · ReadOnly · Volume=helix-pgdata"]
        RD["redis:7\nDROP ALL · ReadOnly · ephemeral presence"]
    end
    CLIENT["agents · :443/udp MASQUE / :51820/udp plain-WG"] --> EDGE
    EDGE --> |127.0.0.1| HELIXD
    HELIXD --> |127.0.0.1:5432| PG
    HELIXD --> |127.0.0.1:6379| RD
```

3.1 helixvpn.pod

```
# ~/.config/containers/systemd/helixvpn.pod
[Pod]
PodName=helixvpn
PublishPort=443:443/udp      # MASQUE/QUIC ingress (Volume 2
    transport ladder)
PublishPort=51820:51820/udp  # plain-WG ingress (auto-ladder
    base)
[Install]
WantedBy=default.target
```

3.2 helix-edge.container — the data plane (the ONLY privileged-ish member)

```
# ~/.config/containers/systemd/helix-edge.container – DATA PLANE
[Unit]
Description=HelixVPN data-plane edge
[Container]
Image=ghcr.io/helixdevelopment/helix-edge:helix_vpn-1.0  #
    §11.4.151 prefixed tag
Pod=helixvpn.pod
DropCapability=ALL
AddCapability=NET_ADMIN NET_RAW      # NET_RAW required for WG/
    tunnels (FACT [RPK §2]) – NOT just NET_ADMIN
AddDevice=/dev/net/tun:/dev/net/tun  # for userspace WG
    (boringtun) fallback path (FACT [RPK §2])
ReadOnly=true                        # read-only rootfs
ReadOnlyTmpfs=true                   # rw tmpfs on /dev /dev/
    shm /run /tmp /var/tmp (default true when ReadOnly)
SeccompProfile=/etc/containers/seccomp/helix-edge.json
Notify=true                          # sd_notify readiness
    passthrough
# no SSH, no shell in this image
[Service]
Restart=always
[Install]
WantedBy=default.target
```

FACT [RPK §2] — NET_ADMIN alone is insufficient. WireGuard/tunnels need **NET_RAW** too ("which Docker enables by default but Podman does not" for rootless), the **host WG kernel module loaded outside the container** (a container cannot modprobe rootlessly; kernels ≥ 5.6 ship it built-in), and **/dev/net/tun** for the userspace path. The unit above adds all three; the kernel module is a host prerequisite documented in the operator guide.

Reconciled (§11.4.35, 2026-06-26): the canonical edge capability set across all four Volume-6 substrate docs (this doc, [docker-compose.md](#), [kubernetes.md](#)) and the [security privescan](#) is **{NET_ADMIN, NET_RAW}** + the **/dev/net/tun device** — no more, no less. This is the FACT the cited research points to (§11.4.6): [research-podman_k8s §2] confirms (3 sources) that the **kernel-mode WireGuard fast path** — the edge's primary data path — needs **both NET_ADMIN and NET_RAW** ("WireGuard needs

NET_RAW ... which Docker enables by default but Podman does not"); /dev/net/tun is added for the **userspace boringtun fallback** (module-less nodes). The earlier "ONLY NET_ADMIN" reading in the K8s/overview/security drafts was the under-specified side and has been corrected to this set everywhere; the privescan asserts **no caps beyond this canonical set**.

Note — build-org ≠ source-org. Container images are published under ghcr.io/helixdevelopment/* (the GHCR build/registry org) while the source submodules live under vasic-digital/* (the GitHub source org). This is intentional, not an inconsistency: the registry namespace and the source namespace differ. The same ghcr.io/helixdevelopment/* image coordinates are used by every Volume-6 substrate (quadlets, Compose, K8s).

3.3 helixd.container — the control plane (drops all caps)

```
# ~/.config/containers/systemd/helixd.container – CONTROL PLANE
[Container]
Image=ghcr.io/helixdevelopment/helixd:helix_vpn-1.0
Pod=helixvpn.pod
DropCapability=ALL                                # control plane needs no
                                                    extra caps
ReadOnly=true
ReadOnlyTmpfs=true
NoNewPrivileges=true
Environment=DATABASE_URL=postgres://helix_app@127.0.0.1/helix #
                                                    non-superuser → RLS enforced [API §6.5]
Environment=REDIS_URL=redis://127.0.0.1:6379
Secret=helix_db_pw,type=env,target=PGPASSWORD            #
                                                    podman secret, NEVER in unit text (§11.4.10)
HealthCmd=/helixd healthz                               #
                                                    → GET /readyz semantics [API §4.7]
[Service]
Restart=always
[Install]
WantedBy=default.target
```

The helix_app (non-superuser) DATABASE_URL is load-bearing: it makes Postgres Row-Level Security the floor even if RBAC is mis-wired (C8 [API §6.5]). The DB password is a **podman secret** injected as an env target, never written into the unit text (§11.4.10).

3.4 helix-pg.container + helix-redis.container

```
# ~/.config/containers/systemd/helix-pg.container
[Container]
Image=docker.io/library/postgres:16
Pod=helixvpn.pod
DropCapability=ALL
Volume=helix-pgdata.volume:/var/lib/postgresql/data:Z    #
                                                    SELinux relabel; data persists
Environment=POSTGRES_USER=helix_owner POSTGRES_DB=helix  #
                                                    owner role; helixd uses helix_app (§3.3)
Secret=helix_db_pw,type=env,target=POSTGRES_PASSWORD
```

```

[Install]
WantedBy=default.target

# ~/.config/containers/systemd/helix-redis.container
[Container]
Image=docker.io/library/redis:7
Pod=helixvpn.pod
DropCapability=ALL
ReadOnly=true
ReadOnlyTmpfs=true
# presence/events are EPHEMERAL (C2: losing Redis loses presence,
#   not identity) – no durable volume
[Install]
WantedBy=default.target

```

Postgres is **not** read-only-rootfs in practice (it writes its socket/lock state); it uses DropCapability=ALL + a named volume for PGDATA. Redis is read-only-rootfs because its state is ephemeral (presence TTL keys), so no durable volume is mounted [API §4.2 — Online from Redis].

3.5 helix.network + helix-pgdata.volume

```

# ~/.config/containers/systemd/helix.network  (optional explicit
#   network; the pod also supplies one)
[Network]
NetworkName=helixvpn
# pasta is the Podman 5.x rootless default backend (NAT-free,
#   faster – FACT [RPK §3])

# ~/.config/containers/systemd/helix-pgdata.volume
[Volume]
VolumeName=helix-pgdata

```

A pod already supplies a shared network namespace, so an explicit .network is optional for the single-pod model; it is included for operators who want a named bridge. The volume quadlet creates the named helix-pgdata referenced by helix-pg.container (§3.4) with an auto-generated dependency [RPK §1].

4. Rootless networking caveats (the load-bearing section for a VPN)

These are the FACT-grade caveats that bite a rootless VPN edge [RPK §3] — each is documented in the operator guide so a fresh init does not hit "connected but nothing loads":

Caveat	FACT	HelixVPN handling
:443 rootless bind	Rootless cannot bind ports < 1024 by default (net.ipv4.ip_unprivileged_port_start=1024) [RPK §3].	Lower the sysctl: net.ipv4.ip_unprivileged_port_start persisted in /etc/sysctl.d/99-helixvpn.conf

Caveat	FACT	HelixVPN handling
pasta backend	Podman 5.x rootless default backend is pasta (NAT-free, faster than slirp4netns); RHEL ≥ 9.5 defaults to it [RPK §3].	ports.conf.init prints this as a prerequisite. Verify via <code>podman info</code> → <code>host.networkBackend</code> . UNVERIFIED regression: a pasta throughput regression reported in Podman 5.8 [RPK §3, podman#28219] — the operator guides <i>benchmark your build</i> , never assumes.
No auto firewall/NAT	"Modifying the firewall requires root ... rootless Podman does not install iptables rules" [RPK §3].	A VPN edge that NATs client egress host-side forwarding OR runs the NAT inside the edge's own netns with NAT documented as a host prerequisite.
MTU pitfall	Tunnels frequently break on MTU ("connected but nothing loads") [RPK §3].	The edge sets the WG/quic MTU (<code>routing-and-addressing.md</code>); the guide lists the MTU knob.
Outbound interface pin	slirp4netns offers <code>--outbound-addr</code> to pin the egress interface [RPK §3].	When the gateway must egress a stream the operator pins it; UNVERIFIED exposes an equivalent on the target.

```
# host prerequisite printed by `helixvpnctl init` (rootless :443):
echo 'net.ipv4.ip_unprivileged_port_start=443' | sudo tee /etc/
    sysctl.d/99-helixvpn-ports.conf
sudo sysctl --system
# (sudo here is a HOST sysctl change, not container management –
    §11.4.161 forbids rootful CONTAINER ops, not host config)
```

§11.4.161 boundary (§11.4.6 honesty). Lowering the host sysctl and loading the WG kernel module are **host configuration** the operator does once, NOT rootful container management. §11.4.161 forbids rootful/sudo *container runtime* operations; it does not forbid the operator from configuring their own host kernel. The distinction is documented so the rootless mandate is not misread as "the host needs no privileged setup".

5. The on-demand-infra invariant (§11.4.76)

Operators (and tests) are **never** required to `podman machine / systemctl start infra` by hand for integration tests — the test entry point boots infra on demand via the containers submodule [05_OV §6.3, §11.4.76(3)]:

```
// helix-go/internal/store/integration_test.go (illustrative –
    FACT-grade pattern [05_OV §6.3])
import (
    "digital.vasic.containers/pkg/boot"
    "digital.vasic.containers/pkg/health"
)

func TestWatchNetworkMapDeltaStream(t *testing.T) {
    ctx := context.Background()
```

```

infra := boot.Compose(ctx, boot.Spec{                                     //
    rootless Podman (§11.4.161)
    Services: []boot.Service{
        {Name: "pg", Image: "docker.io/library/
postgres:16", Port: 5432},
        {Name: "redis", Image: "docker.io/library/
redis:7", Port: 6379},
    },
})
t.Cleanup(func() { infra.Down(ctx) })                                     //
    §11.4.14 quiescent teardown
health.WaitReady(ctx, infra, 30*time.Second)                             // real
    readiness probe, NOT a sleep
// ... drive enroll → advertise → policy → assert MapDelta
}

```

This is anti-bluff-critical (§11.4.76(5)): an integration test claiming to exercise the streamed WatchNetworkMap MUST actually boot Postgres+Redis via the submodule — a short-circuit fake that skips boot is a §11.4 violation. The production quadlets and the test infra both go through the containers submodule; no ad-hoc podman run exists in the codebase.

6. helixvpctl deploy quadlet render (the source of these units)

helixvpctl deploy quadlet --out ~/.config/containers/systemd/ renders the §3 units from one in-code spec via the containers submodule's pkg/compose primitives (§11.4.81 cross-platform- parity: one source, per-substrate render) [05_OV §6.3, helixvpctl §10]. The render is **parameterized** by the init inputs (domain → masque_sni, data-dir → volume host path, image tags → the helix_vpn-<version> release prefix [§11.4.151]). The three substrates (quadlet / compose / kube) derive from the **same** spec, so a change to the pod shape lands in all three, not three hand-maintained files.

7. Security posture summary

Control	Mechanism	FACT source
Least privilege	DropCapability=ALL on every member; AddCapability=NET_ADMIN NET_RAW ONLY on the edge	[RPK §1/§2]
Read-only rootfs	ReadOnly=true + ReadOnlyTmpfs=true (rw tmpfs on /dev /run /tmp) on edge/control/redis	[RPK §1]
Seccomp	SeccompProfile=/etc/containers/seccomp/helix-edge.json (custom, tighter than default)	[RPK §1]
No privilege escalation	NoNewPrivileges=true on control	[RPK §1]

Control	Mechanism	FACT source
Secrets	podman Secret= injected as env target; never in unit text	§11.4.10
Rootless	file-location-selected; no User=, no daemon, no sudo for container ops	§11.4.161, [RPK §1]
No shell/SSH in images	edge/control images carry no shell — reduces attack surface	[05_OV §7.1]
SELinux relabel	:Z on the pgdata volume mount	[RPK §1]

8. Test points (§11.4.169 comprehensive test-type coverage)

Every PASS cites captured evidence (§11.4.5/.69/.107); infra is the containers submodule (§11.4.76), rootless Podman (§11.4.161), never ad-hoc podman run.

#	Test type (§11.4.169)	Target	Concrete assertion + evidence
Q1	unit	render determinism	deploy quadlet twice → byte-identical units; image tag carries helix_vpn-prefix
Q2	integration	rootless one-pod boot	place units in ~/.config/containers/systemd/, systemctl --user start helixvpn-pod, all 4 members UP; status GREEN
Q3	integration	:443/udp reachable rootless	with the sysctl set, an agent completes MASQUE handshake to :443/udp; captured
Q4	security	capability set (§3.2)	podman inspect shows edge has ONLY NET_ADMIN+NET_RAW; control/pg/redis have none; mutation adding a cap → FAIL
Q5	security	read-only rootfs	a write to the edge rootfs is denied (EROFS); rw tmpfs on /run works
Q6	security	secret not in unit text	grep the rendered units → no PGPASSWORD value; only Secret= references
Q7	integration	on-demand-infra invariant (§5)	go test boots infra via boot.Compose; mutation removing boot.Compose → test FAILs (proves it was not a fake)
Q8	chaos	member restart	kill helix-pg, assert Restart=always recovers + control reconnects; presence (redis) loss degrades gracefully (C2)
Q9	e2e	three-substrate parity	quadlet + compose + kube boot the same image; helixvpctl status GREEN on each (recorded per §11.4.159)

#	Test type (§11.4.169)	Target	Concrete assertion + evidence
Q10	meta-test (§1.1)	gates paired	add a cap / inline a secret / skip read-only → each makes its gate FAIL

9. Decision callouts (options + recommendation — §11.4.66/.101)

id	Decision	Options	Recommendation
D-SUBSTRATE-DEFAULT	init default substrate	quadlet / compose / k8s	quadlet (§11.4.76/.161 canonical) — flag overrides
D-ROOTLESS-443	rootless :443 bind	(a) lower ip_unprivileged_port_start; (b) root-side proxy/authbind	(a) the clean path for :443/udp [RPK §3]; documented host prerequisite
D-WG-PATH	edge WG path	(a) kernel WG (NET_ADMIN+NET_RAW + host module); (b) userspace boringtun (+ /dev/net/tun)	(a) primary, (b) fallback where the host module is unavailable (Volume 2)

Sources verified

- v09-research/research-podman_k8s.md §1 (quadlet fundamentals, rootless search paths, User= hard constraint, [Container]/[Pod] directives, example skeleton), §2 (NET_ADMIN insufficient — also NET_RAW + host WG module + /dev/net/tun), §3 (rootless networking caveats: pasta default
 - 5.8 regression, no auto firewall/NAT, MTU, rootless :443 sysctl), §7 (FACT-grade carry-into- spec list) — verified against latest Podman docs 2026-06-25 per §11.4.99 — [RPK].
- 05-repo-layout-tooling-and-helix-ecosystem.md §7.1 (canonical quadlets, the pod/edge/control/pg unit sketches), §6.3 (containers submodule on-demand-infra + deploy render), §7.4 (substrate recommendation) — [05_0V].
- v03-control-plane/svc-api.md §4.7 (/readyz semantics for HealthCmd), §6.5 (helix_app non-superuser → RLS), §4.2 (Redis presence ephemeral, C2) — [API].
- Constitution anchors: §11.4.161 (rootless container runtime mandate — Podman rootless, no sudo/ rootful), §11.4.76 (containers submodule sole orchestration + on-demand-infra invariant), §11.4.10 (secrets as podman Secret, never in unit text), §11.4.81 (one-source three-substrate render), §11.4.112 (honest documented gap if a platform has no rootless option), §11.4.35 (visible reconciliation — rootless path vs the overview's /etc path), §11.4.6 (no-guessing

— UNVERIFIED pasta-regression/outbound-pin marks), §11.4.169 (test-type coverage), §1.1 (paired meta-test mutations).