

Challenges — HelixVPN nano-detail spec (Volume 8 · §11.4.169 type 5)

Challenges — HelixVPN nano-detail spec (Volume 8 · §11.4.169 type 5)

Revision: 2 Last modified: 2026-06-26T12:00:00Z

Reconciled (§11.4.35, 2026-06-26): the MVP Definition-of-Done has **NINE** acceptance criteria (AC1–AC9) per the canonical overview [10-testing-acceptance-and-qa.md §7.2](#). The earlier "eight ... (AC1–AC9)" wording (§1, §6) was a self-contradictory mis-count; a Challenge gate scoring 8/9 ACs is a §11.4.118 coverage bluff. Corrected to NINE.

Nano-detail expansion of [§5.5 of the Volume-8 overview](#). Challenges are the constitution's anti-bluff Challenge layer (§11.4.27(B)) — the vasic-digital/challenges submodule, a Go engine that scores a feature **PASS only** on captured evidence matching the §11.4.69 shape for that feature class, graded by the **engine re-reading the artifact**, never by the test's own exit code. This is the structural defence against bluff-classes **B1 (config-only)**, **B2 (absence-of-error)**, and **B3 (wrong-plane)**: the Challenge engine does not trust the harness's claim — it re-reads the pcap / counter / recording itself. The nine MVP Definition-of-Done acceptance criteria (AC1–AC9, [overview §7.2](#)) are authored as executable Challenges. This document fixes the bank model, the engine contract, the evidence scoring, the determinism, the gate, the paired §1.1 mutation, and a concrete bank YAML. Spec-only; unproven assumptions marked UNVERIFIED. Siblings: [unit.md](#), [integration.md](#), [e2e.md](#), [full-automation.md](#), [helixqa.md](#).

Table of contents

- [1. Scope — what a Challenge is, and what it covers](#)
 - [2. Harness — the challenges submodule engine](#)
 - [3. Fixtures — real on-rig drivers, engine-scored evidence](#)
 - [4. Evidence taxonomy — engine-re-read, not exit-code](#)
 - [5. Determinism — N-iteration identical evidence](#)
 - [6. Acceptance gate — when CHAL blocks a release](#)
 - [7. The paired §1.1 mutation \(anti-bluff proof\)](#)
 - [8. The MVP-DoD bank \(concrete YAML\)](#)
 - [9. Open decisions surfaced for QA](#)
 - [Sources verified](#)
-

1. Scope — what a Challenge is, and what it covers

A Challenge differs from a test in **who grades it**. A test asserts and exits with a code; a Challenge runs a driver, then the challenges engine **independently re-reads the captured evidence artifact** and scores PASS only if the artifact's content matches the declared §11.4.69 shape for that feature class. The test's exit code is *not* the verdict — the re-read is. This is the §11.4.27(B) anti-bluff layer: even a perfectly-written but lying harness cannot fake a Challenge, because the engine never sees the harness's claim, only its artifact.


Challenge	Feature class (§11.4.69)	What the engine re-reads	Bound AC
HVPN-CHAL-AC1-selfhost	boot_service	recorded terminal + helix_up==1 healthcheck	AC1
HVPN-CHAL-AC2-authorized-reach	network_connectivity	pcap: SYN→SYN-ACK from 10.10.0.20 + non-empty body hash	AC2
HVPN-CHAL-AC3-default-deny	network_connectivity (negative)	pcap: SYN out, zero SYN-ACK; edge drop++	AC3
HVPN-CHAL-AC4-masque-escalate	network_connectivity	tshark: HTTP/3 flow, no WG signature; transport=="masque-h3"	AC4
	counter delta		AC5

Challenge	Feature class (§11.4.69)	What the engine re-reads	Bound AC
HVPN-CHAL-AC5-reconcile		helix_reconcile_seconds p99 histogram bucket < 1 s	
HVPN-CHAL-AC6-revoke	counter delta	revoke-ts → edge-peer-removed-ts CSV, p99 < 1 s	AC6
HVPN-CHAL-AC7-killswitch-no-leak	killswitch_negative	host pcap: zero plaintext egress, zero :53 DNS	AC7
HVPN-CHAL-AC8-no-log-schema	negative (schema)	schemalint report: zero durable per-flow tables	AC8
HVPN-CHAL-AC9-three-apps	REC	per-app window-scoped MP4 + vision_engine verdict	AC9

Challenges cover the **DoD acceptance set** — the release-gating proof that the product works for the end user. They sit at the top of the test pyramid below REC ([overview §4](#)) and are driven by helix_qa ([helixqa.md](#)).

2. Harness — the challenges submodule engine

The challenges submodule (`digital.vasic.challenges`, Go) is consumed via a replace directive in development and a **pinned SHA** in release ([05 §6.1](#)), per §11.4.27(B)/§11.4.74 (reuse-don't-reimplement) + §11.4.28 (owned submodule, equal-codebase). A **bank file** (`challenges/banks/helixvpn_mvp_dod.yaml`) maps each AC to its on-rig driver + the evidence-class assertion the engine scores against.



- **Consumed by reference, never copied** (§11.4.28) — the engine lives in the submodule; the consumer registers its banks as data (the YAML), keeping the engine project-agnostic.
- **Engine-scored, never exit-code** — the engine runs the driver, captures the artifact, then **re-reads** it; `score: pass_only_if_evidence_matches` is the contract. A driver that exits 0 but produced no matching artifact scores **FAIL**.
- **Self-validated** (§11.4.107(10)) — a Challenge with `self_validated: true` ships a **golden-bad** artifact (a seeded-leak pcap, a WG flow

mislabelled H3) that the engine **must** score FAIL; if it scores PASS, the analyzer is the bluff and the meta-test FAILs (CM-ANALYZER-SELF-VALIDATED, defeats B5).

3. Fixtures — real on-rig drivers, engine-scored evidence

Fixture	Real or mock	Notes
AC drivers	real rig scripts (e2e.md §8 , §5.7 SEC)	rig/test_reach.sh, rig/killswitch_drop.sh, etc. — real packets
captured artifacts	real pcap/CSV/MP4/schemalint report	the engine re-reads these; they are not synthetic
golden-good fixture	real captured-clean artifact	the self-validation positive case
golden-bad fixture	real captured-bad artifact (seeded leak)	the self-validation negative case — must score FAIL
mocks	FORBIDDEN (§11.4.27)	Challenges are above the unit layer; the engine re-reads real evidence

The engine's re-read is the anti-mock guarantee: there is nothing to mock — the verdict is computed from a real pcap/CSV/MP4 captured during a real run. The only out-of-band step is the §11.4.10 credential bootstrap for AC1's init.

4. Evidence taxonomy — engine-re-read, not exit-code

The defining property: the Challenge **verdict** = **the engine's re-read of the artifact**, mapped to the §11.4.69 class shape:

Class	Engine assertion	Forbidden bluff it defeats
network_connectivity	tshark/pcap filter matches (SYN-ACK present) and body hash non-empty	B1 config-only — a config check produces no pcap
network_connectivity (negative)	filter matches zero rows (no SYN-ACK) + edge drop++	B2 absence-of-error — "no error"

Class	Engine assertion	Forbidden bluff it defeats
killswitch_negative	assert_zero: "ip && not loopback" + assert_zero_dns: "udp.port==53"	is not "zero SYN-ACK captured" B2 — the engine counts plaintext/DNS packets, not "kill-switch enabled"
counter delta	histogram bucket / timing CSV p99 < budget	B3 wrong-plane — a status field is not a measured delta
REC	vision_engine OCR verdict on the MP4 + core-FSM cross-check	B3 — a green shield while the FSM is Blocked FAILs

Per §11.4.83 the engine's per-AC verdict JSON lands under docs/qa/<run-id>/ (the committed evidence trail); per §11.4.69 a PASS without a matching artifact is impossible — the engine has nothing to score. The verdict JSON carries the artifact path (the §11.4.116 verdict-carries-evidence-path rule, [helixqa.md §4](#)).

5. Determinism — N-iteration identical evidence

Per §11.4.50 each Challenge runs N=3 (normal) / N=10 (cycle-validation) under the FA wrapper ([full-automation.md §5](#)):

1. The driver runs on the clean rig per iteration; the engine re-reads each iteration's artifact; the **structural hash** of the matched-evidence set is compared across runs; a divergence is auto-FAIL.
2. The body-hash / drop-count / schemalint-table-set are deterministic by construction (fixed sink page, fixed policy); timing fields are compared to a budget, not hashed.
3. A Challenge that scores PASS on run 1 and FAIL on run 2 is **not** a flake — it is a real defect surfaced (§11.4.50 mechanises §11.4.7); the §11.4.4 STOP fires.

6. Acceptance gate — when CHAL blocks a release

Gate	Bar	Layer
<code>make qa</code> (overview §9)	every DoD Challenge scores PASS via the engine, N=3	release sweep
CM-DOD-CHALLENGES-GREEN	a release tag is blocked while any DoD Challenge is non-PASS	pre-tag (§11.4.40)
CM-ANALYZER-SELF-VALIDATED (§11.4.107(10))	every <code>self_validated</code> Challenge's golden-bad artifact scores FAIL	pre-build meta-test

CHAL is a **release-gating layer**: the MVP ships only when all nine ACs' Challenges score PASS on a clean baseline (§11.4.40 + §11.4.108 runtime-signature on a fresh deploy). This is QA-D5 ([overview §10](#)) resolved: **challenges-scored, not raw exit codes** — the engine re-reads the evidence artifact, structurally defeating B1–B3; exit codes alone are bluffable (§11.4.27).

7. The paired §1.1 mutation (anti-bluff proof)

```
# §1.1 mutation for CM-ANALYZER-SELF-VALIDATED (the engine's self-validation is real)
- the AC7 kill-switch Challenge has self_validated: true with golden-bad = fixtures/killswitch_leaked_dns.pcap (a seeded :53 leak)
- MUTATE the engine's assertion: strip `assert_zero_dns: "udp.port==53"`
- assert: the engine now SCORES the leaked-DNS golden-bad pcap as PASS → meta-test FAILs
  (the analyzer became a bluff – B5 caught)
- restore the assertion; assert golden-bad scores FAIL again (§11.4.84 quiescence)

# §1.1 mutation for CM-DOD-CHALLENGES-GREEN (engine re-reads, not exit code)
- make rig/test_reach.sh `exit 0` but produce an EMPTY pcap (a lying harness)
- assert: the engine re-reads the empty pcap, finds no SYN-ACK → scores FAIL (not PASS)
  → proves the verdict comes from the artifact, not the exit code (defeats B1)
```

The second mutation is the structural proof of the Challenge model: a harness that *claims* success but produced no matching evidence scores FAIL, because the engine grades the artifact, never the claim.

8. The MVP-DoD bank (concrete YAML)

```
# challenges/banks/helixvpn_mvp_dod.yaml – one entry per DoD AC
# (§11.4.27(B)/.69/.83)
bank: helixvpn-mvp-dod
engine:
  digital.vasic.challenges          # consumed by-reference,
  pinned SHA in release (§11.4.28)
score_policy: pass_only_if_evidence_matches # NEVER trust a
  driver's exit code
evidence_root: docs/qa/$
  {RUN_ID}                          # §11.4.83 committed trail

challenges:
- id: HVPN-CHAL-AC2-authorized-reach
  feature_class: network_connectivity          # §11.4.69
  driver: rig/test_reach.sh pass
  evidence:
    kind: pcap
    assert: "frame.tcp.flags.syn==1 && frame.tcp.flags.ack==1
      from 10.10.0.20"
    and_body_hash_nonempty: true                # defeats B1
      (config-only)

- id: HVPN-CHAL-AC3-default-deny
  feature_class: network_connectivity_negative
  driver: rig/test_reach.sh deny
  evidence:
    kind: pcap
    assert_zero: "tcp.flags.syn==1 &&
      tcp.flags.ack==1" # ZERO SYN-ACK
    and_edge_counter: "drop > 0"                #
      defeats B2 (absence-of-error)

- id: HVPN-CHAL-AC4-masque-escalate
  feature_class: network_connectivity
  driver: rig/escalate.sh
  evidence:
    kind: pcap+status
    assert: "http3 present && wireguard absent" #
      tshark classification
    and_status: "transport == 'masque-h3'"      #
      defeats B3 (wrong-plane)

- id: HVPN-CHAL-AC6-revoke
  feature_class: counter_delta
  driver: rig/revoke.sh
```

```

evidence:
  kind: csv
  assert_p99_lt_seconds: 1.0 #
    revoke→enforcement < 1s (SL03)

- id: HVPN-CHAL-AC7-killswitch-no-leak
  feature_class: killswitch_negative
  driver: rig/killswitch_drop.sh
  evidence:
    kind: pcap
    assert_zero: "ip && not loopback" # zero plaintext
      egress on drop
    assert_zero_dns: "udp.port==53" # zero DNS leak
  self_validated: true # golden-bad
    seeded-leak pcap MUST score FAIL (§3.3)
  golden_good: fixtures/killswitch_clean.pcap
  golden_bad: fixtures/killswitch_leaked_dns.pcap


- id: HVPN-CHAL-AC8-no-log-schema
  feature_class: schema_negative
  driver: scripts/schemalint.sh
  evidence:
    kind: report
    assert: "forbidden_table_count == 0" # no durable per-
      flow table

- id: HVPN-CHAL-AC9-three-apps
  feature_class: REC # §11.4.153/.159
  driver: panoptic drive --app access,connector,console --flow
    enroll-connect-reach
  evidence:
    kind: mp4
    window_scoped: true # §11.4.154 –
      app window only, not desktop
    filename_prefix: helixvpn--- # §11.4.155
      project prefix
    vision_verdict: vision_engine # §11.4.163
      media-validation re-read
    expected_patterns:
      ["Connect","Connecting","Connected","shield:green"]
    cross_check: { core_fsm_at_green_shield: "Connected" } #
      defeats B3
  self_validated:
    true # golden-bad MP4 (green
      shield, FSM Blocked) MUST FAIL

```

8a. Challenge lifecycle (engine internals)

A Challenge moves through a fixed lifecycle inside the engine; understanding it is how the consumer reasons about a verdict:



Failure-mode catalog (what scores FAIL, and why that is correct):

Symptom	Engine behaviour	Why it is a FAIL not a flake
driver exits 0, empty pcap	NO_ARTIFACT → FAIL	the verdict is the re-read, not the exit code (B1)

Symptom	Engine behaviour	Why it is a FAIL not a flake
pcap present, no SYN-ACK row	SCORE → FAIL	reach not proven on the wire (§11.4.69)
transport=="masque-h3" but pcap has a WG signature	SCORE → FAIL	wrong-plane lie (B3) — status and wire disagree
golden-bad leaked-DNS pcap scores PASS	ANALYZER_BLUFF → FAIL	the analyzer cannot distinguish leak from clean (B5)
run 1 PASS, run 2 FAIL	§11.4.50 auto-FAIL	a real defect surfaced, never demoted to "flake" (§11.4.7)

The engine **never** converts an empty/unreachable artifact into a PASS-counting SKIP for a feature class that has a sink-side probe (§11.4.69 CM-NO-FAIL-OPEN-SKIP) — a genuinely-absent topology is an honest §11.4.3 SKIP with a closed-set reason surfaced on the sync channel ([helixqa.md §4](#)), never a silent pass.

9. Open decisions surfaced for QA

#	Decision	Options	Recommendation
C-D1	Challenge engine vs raw exit codes (QA-D5)	exit-code vs challenges-scored	challenges-scored (overview §10) — the engine re-reads the artifact, structurally defeating B1–B3
C-D2	bank consumption mode	replace (dev) vs pinned SHA (release)	replace in dev, pinned SHA in release (05 §6.1) — reproducible release, fast dev iteration
C-D3	golden-bad fixture authorship	hand-seed each vs derive from a real failed run	hand-seed the canonical bad artifact + add any real failed-run capture as a second golden-bad — broadens the self-validation surface (§11.4.107(10))

9a. Challenge-vs-test-vs-HelixQA — who owns what

The three top-of-pyramid layers are distinct and composing, not redundant:

Layer	Grades by	Owns	Spec
test (E2E/ INT)	its own exit code + an ab_pass_with_evidence call	the assertion	e2e.md , integration.md
Challenge	the challenges engine re-reading the artifact	the DoD AC verdict (B1– B3-proof)	this doc
HelixQA	the orchestrator driving + reconciling all of them, no human	the autonomous session + sync channel + tickets	helixqa.md

A Challenge consumes a test's *driver* but ignores its *verdict*; HelixQA consumes a Challenge's *verdict* and reconciles it into the coverage ledger + the workable-items DB. The chain — test drives, Challenge re-reads, HelixQA orchestrates — is three independent re-reads of the same captured evidence, which is precisely why a bluff must survive all three to ship, and by construction cannot (§11.4.27(B) / §11.4.165).

Why the engine is a submodule, not in-project (§11.4.74/.28). The re-read logic (pcap filters, counter parsers, OCR verdict) is reusable across every Helix project that needs anti-bluff acceptance scoring; reimplementing it in-project would duplicate the 80%+ that challenges already covers (§11.4.74 catalogue-first). The consumer contributes only **data** (the bank YAML) and any missing assertion kind is added by an upstream PR to vasic-digital/challenges (§11.4.74), never a local fork — keeping the engine project-agnostic and the consumer decoupled (§11.4.28).

UNVERIFIED: the `digital.vasic.challenges` engine's exact bank-YAML schema (`score_policy`, `assert_zero`, `self_validated`, `vision_verdict` keys) is the **intended** shape per the overview skeleton; confirm against the pinned `vasic-digital/challenges` SHA (§11.4.74 — extend upstream if a key is unsupported).

Sources verified

- Volume-8 overview §2 (taxonomy row CHAL), §5.5 (the bank YAML skeleton + engine-scored model), §5.6 (helix_qa drives banks), §3.3 (self-validated analyzers), §7.2 (AC1–AC9), §10 QA-D5 — read 2026-06-26.
- Sibling specs cross-referenced: e2e.md, helixqa.md, full-automation.md, ../05-repo-layout-tooling-and-helix-ecosystem.md — overview + repo doc read; the [05] ecosystem doc is the [P1] overview present at final/.
- Constitution: §11.4.27(B) (Challenges / no-fakes-beyond-unit), §11.4.69 (sink-side evidence taxonomy), §11.4.107(10) (self-validated analyzer / golden-good+bad), §11.4.83 (docs/qa/<run-id>/ evidence trail), §11.4.116 (verdict-carries-evidence-path), §11.4.28 (owned submodule consumed by reference), §11.4.74 (extend-upstream), §11.4.50 (determinism), §11.4.40 (pre-tag gate), §11.4.108 (runtime-signature on clean deploy), §11.4.153/.154/.155/.159/.163 (REC / window-scope / prefix / media-validation for AC9), §1.1 (paired mutation) — from CLAUDE.md in-context.
- The challenges engine bank-YAML schema is reproduced from the overview skeleton — **not independently fetched from vasic-digital/challenges** (UNVERIFIED; confirm at the pinned SHA).