

End-to-End Testing — HelixVPN nano-detail spec (Volume 8 · §11.4.169 type 3)

End-to-End Testing — HelixVPN nano-detail spec (Volume 8 · §11.4.169 type 3)

Revision: 2 Last modified: 2026-06-26T12:00:00Z

Reconciled (§11.4.35, 2026-06-26): the §2 harness topology + skeletons were rewritten onto the **canonical four-namespace rig** (client / censor / gateway / exit) of [test-rig.md](#), with the DPI/censor in a distinct censor-ns middlebox and the wire/DPI capture point on the censor egress cen2gw. The earlier inferior 3-namespace rig (client / gateway / lanA with the nft block bolted onto the gateway forward hook) is superseded — test-rig.md is the single source for the rig's namespace/interface names and capture points.

Nano-detail expansion of [§5.3 of the Volume-8 overview](#). E2E proves the **whole reachability slice on one box**: a real control plane issues a NetworkMap, a real helix-core client builds a real tunnel, real packets cross client → gateway → connector LAN, and a real curl returns the hello page — or, for the negative E2E, a SYN leaves and **zero** SYN-ACK returns (default-deny). The substrate is the Linux **netns + nftables-DPI + tc-netem rig** [04_P0 §3]. This is where bluff-classes **B1 (config-only)** and **B3 (wrong-plane)** are mechanically defeated by pcap + body-hash + wg-counter-advance evidence (§11.4.107). Spec-only; unproven assumptions marked UNVERIFIED. Siblings: [unit.md](#), [integration.md](#), [full-automation.md](#), [challenges.md](#), [helixqa.md](#).

Table of contents

- [1. Scope — the reachability slices E2E covers](#)
 - [2. Harness — the netns + nftables-DPI + tc-netem rig](#)
 - [3. Fixtures — real processes, simulated network, zero mocks](#)
 - [4. Evidence taxonomy + the §11.4.107 liveness battery](#)
 - [5. Determinism — N-iteration identical evidence](#)
 - [6. Acceptance gate — when E2E blocks a release](#)
 - [7. The paired §1.1 mutation \(anti-bluff proof\)](#)
 - [8. Test skeletons](#)
 - [9. Open decisions surfaced for QA](#)
 - [Sources verified](#)
-

1. Scope — the reachability slices E2E covers

E2E owns the user-visible security properties that **only a packet on a real tunnel can prove**. Each slice maps to a Phase-1 Definition-of-Done AC ([overview §7.2](#)).


Slice	Path	The user-visible property	Bound AC
Authorized reachability	client → gateway → connector LAN, curl http://10.10.0.20/ ⇒ 200	the user reaches a host the policy grants	AC2
Default-deny (negative E2E)	client → gateway, SYN to an unauthorized host	a host the policy does not grant is unreachable — SYN out, zero SYN-ACK	AC3
Transport escalation	plain-WG UDP blocked → auto-escalate to MASQUE/H3 → reach preserved	the tunnel survives a DPI UDP block, classified as HTTP/3 not WireGuard	AC4
Policy reconcile (live)	a policy edit reconfigures the running slice with no restart	a revoked/added grant takes effect on the live tunnel	AC5
Revoke (live)	revoke a device → its peer drops from every map → tunnel sealed	a revoked device loses reachability	AC6
Resilience		the tunnel holds goodput under	G1/G2 [04_P0]

Slice	Path	The user-visible property	Bound AC
	reach under 5% loss / 40 ms jitter (tc-netem)	impairment (adjacency to STRESS §11.4.85)	

The rig is the **reproducible substrate** for AC2/AC3/AC4 and survives Phase-0→1 ([overview §5.3](#)). E2E does **not** drive the Flutter UI — that is UI/REC via panoptic ([overview §5.14](#)); E2E drives the *core* and the *wire*.

2. Harness — the netns + nftables-DPI + tc-netem rig

The E2E substrate is the **canonical four-namespace rig** defined in [test-rig.md](#): four real Linux network namespaces — client, censor, gateway, exit — on one host. The adversary is split out of the path as its **own censor-ns middlebox** the packet must traverse, so the DPI/censor nft ruleset is a *distinct forwarding hop* (not an nft rule bolted onto an endpoint, [test-rig.md §0](#)); tc netem on the censor egress (cen2gw) injects loss/delay for the resilience slice. The rig is the Phase-0 deliverable HVPN-P0-022/023 ([06](#)). The rig scripts ([rig/netns_up.sh](#), [rig/netns_down.sh](#), the regime nft profiles, [rig/impair.sh](#), [rig/exercise.sh](#)) are defined canonically in [test-rig.md §§2–4](#) and shared by every skeleton below — same namespace names, same interface names, same cen2gw capture point.



The sudo for ip netns add is the **only** privileged step in the whole harness and is a documented, scoped CAP_NET_ADMIN exception — never a container-management escalation ([overview §9](#), §11.4.161). Container infra (the real control plane's PG/Redis behind the gateway) is still booted rootless via containers ([integration.md §2](#)).

```
# rig/netns_up.sh – the 4-ns rig (CANONICAL definition: test-
# rig.md §2)
# client(10.0.1.2) → censor → gateway(10.0.2.2:443) →
# exit(10.10.0.20)
set -euo pipefail
for ns in client censor gateway exit; do ip netns add "$ns" 2>/
dev/null || true; done
# veths client↔censor (10.0.1.0/30) · censor↔gateway (10.0.2.0/30)
· gateway↔exit (10.10.0.0/24),
```

```
# forwarding + routes, and the exit-ns sinks are wired in full by
  test-rig.md §2:
ip netns exec exit python3 -m http.server 80 --bind 10.10.0.20
& # the hello service (sink @10.10.0.20)

# rig/dpi_block.sh – DPI/censor sim on the CENSOR-ns forward hook
  (the middlebox), AC4 escalation E2E
ip netns exec censor nft add table inet censor
ip netns exec censor nft add chain inet censor fwd '{ type filter
  hook forward priority 0; policy accept; }'
ip netns exec censor nft add rule inet censor fwd udp dport 51820
  drop # kill plain WireGuard
ip netns exec censor nft add rule inet censor fwd udp dport 443
  accept # allow MASQUE/H3

# rig/impair.sh – loss/jitter on the censor egress cen2gw
  (§11.4.85 adjacency; CANONICAL: test-rig.md §4)
ip netns exec censor tc qdisc add dev cen2gw root netem loss 5%
  delay 40ms 10ms
```

Cleanup is non-negotiable: trap 'rig/netns_down.sh' EXIT deletes all four namespaces and leaves the host quiescent (§11.4.14); the orchestrator's post-test sanity check FAILs the run if any client/censor/gateway/exit namespace, the inet censor nft table, or a netem qdisc survives.

3. Fixtures — real processes, simulated network, zero mocks

Fixture	Real or mock	Notes
helix-core client	real build	builds a real WireGuard tunnel (boringtun or kernel WG, ../v02-data-plane/wireguard-core.md)
gateway + helix-edge	real build	real verdict-map packet classification
control plane (PG/Redis/coordinator)	real , booted via containers	issues the real NetworkMap the client applies
connector LAN sink	real python3 -m http.server	a real HTTP server — the reachability target
DPI block	real nftables rule	a real kernel drop of UDP/51820 (not a mock)
network impairment	real tc netem	real loss/jitter in the kernel
mocks	FORBIDDEN (§11.4.27)	E2E is post-unit; a mock socket here is a §11.4.27 violation

Everything on the path is real; only the *network topology* is simulated (netns + nftables + netem stand in for "a censored hostile network" reproducibly). This is exactly the §11.4.3 per-environment-topology dispatch: the rig **is** the topology, and absence of CAP_NET_ADMIN $\Rightarrow$ honest SKIP: topology_unsupported, never a fake PASS.

4. Evidence taxonomy + the §11.4.107 liveness battery

An E2E PASS cites **packet-level captured evidence** (§11.4.69 classes network_throughput / network_connectivity), never a status field (B3) or a config check (B1).

Slice	§11.4.69 evidence shape	Defeats
authorized reach (AC2)	pcap with a SYN $\rightarrow$ SYN-ACK from 10.10.0.20 and a non-empty curl body hash	B1 (config-only)
default-deny (AC3)	pcap showing SYN out, zero SYN-ACK; edge verdict-map drop++ counter	B1, B2
escalation (AC4)	StatusReport.transport=="masque-h3" and tshark classifies the flow as HTTP/3 with no WG signature	B3 (wrong-plane)
reach liveness (G1)	iperf3 -J JSON sum_received.bits_per_second > 0 over ≥ 10 s and wg show <if> transfer rx/tx advancing	B3
resilience (G2)	goodput under netem $\geq$ budget; pcap of decrypted inner packets	B3

The §11.4.107 liveness battery, instantiated for E2E (overview §3.2):

1. **Liveness over a window, not a point** — throughput sampled ≥ 10 s; a one-shot ping reply is a pre-filter, never the proof (§11.4.107(1)).
2. **Independent counter-advance** — goodput (iperf3) **and** the kernel WG counter (wg show ... transfer) must *both* move; a flat WG counter with moving iperf3 means a decoy/loopback path $\rightarrow$ FAIL (§11.4.107(2)).
3. **Not-stale cross-check** — after a transport escalation, the post-escalation pcap's first inner packet must not replay the pre-escalation flow's last packet (§11.4.107(4)).
4. **Loading is a distinct state** — a handshake in progress is Connecting, not Connected; liveness is judged only after the FSM reaches Connected (§11.4.107(3)); timeout+reachable $\Rightarrow$ FAIL, timeout+unreachable $\Rightarrow$ SKIP.

5. **Metamorphic relations** (no golden source) — the same LAN host reached over WG-UDP vs MASQUE returns the **same body hash**; a paused tunnel stops the WG counter; 2× clients ⇒ ~2× aggregate goodput (§11.4.107(8)).

Every analyzer (the tshark classifier, the leak detector) is **self-validated** (§11.4.107(10)) with a golden-good + golden-bad pcap pair ([overview §3.3](#)) — defeating bluff-class **B5**.

5. Determinism — N-iteration identical evidence

Per §11.4.50, E2E runs N=3 (normal) / N=10 (cycle-validation) and must produce identical *verdicts* with structurally-identical evidence:

1. **Body-hash equality** — the curl body hash is identical across runs (the sink serves a fixed page); a diverging hash is auto-FAIL.
 2. **Verdict equality** — pass→pass→pass or the run is a flake-FAIL; "first reach worked, retry it" is forbidden (§11.4.50 mechanises §11.4.7).
 3. **Deterministic impairment seed** — tc netem uses a fixed loss/delay so the resilience slice's goodput stays within a tolerance band across runs; the band, not an exact value, is the assertion.
 4. **Clean rig per iteration** — netns_up.sh/netns_down.sh bracket each iteration so no stale tunnel/route/qdisc shadows the next run (§11.4.108/.139 clean-artifact; defeats B4 stale-state).
-

6. Acceptance gate — when E2E blocks a release

Gate	Bar	Bound item
make e2e (overview §9)	positive reach + negative deny both GREEN with pcap	AC2 + AC3
G1 (Phase-0)	iperf3 ≥ 80% bare-link; curl 10.10.0.20:80 OK	E2E + BENCH
G2 (Phase-0)	core over MASQUE through a DPI UDP block ≥ 50% of plain; tshark = HTTP/3, no WG sig; survives 5% loss	E2E + SEC
G6 (Phase-0)	a static-map edit reconciles the slice with no restart	E2E + INT
AC4 / AC5 / AC6 (Phase-1)	escalation / reconcile / revoke proven on the live tunnel	E2E challenges

E2E is the **highest-authority reachability gate**; an E2E FAIL is a release blocker and triggers the §11.4.4 STOP + §11.4.114 known-good isolation. Per §11.4.132 risk-ordering the **security-floor E2E** (default-deny AC3, revoke AC6) runs *before* any convenience slice, and only after that set is GREEN with captured pcap does the rest run.

7. The paired §1.1 mutation (anti-bluff proof)

```
# §1.1 mutation for CM-DEFAULT-DENY-E2E (defeats the most
dangerous bluff: fail-open)
- in helix-edge verdict-map, change the default action:
    default: DROP                // correct: deny unless granted
    MUTATED: default: ACCEPT     // fail-OPEN – every host
    reachable
- assert: rig/test_reach.sh deny now sees a SYN-ACK from the
unauthorized host → E2E FAILs
- restore; assert deny PASS (zero SYN-ACK) again (§11.4.84
quiescence)

# §1.1 mutation for CM-ESCALATION-E2E (defeats B3 wrong-plane)
- force the escalation path to leave the StatusReport.transport at
"wg-udp" while the
  DPI block is active (the bug: reports masque but actually still
  tries WG)
- assert: tshark finds NO HTTP/3 flow + reach FAILs → the test
catches the lie
```

The leak/classifier analyzers' golden-bad fixtures (a seeded-leak pcap, a WG flow mislabelled as H3) **must** FAIL their analyzers — the §11.4.107(10) self-validation that defeats B5; stripping the discriminating tshark filter makes golden-bad pass → the meta-test catches the mutation.

8. Test skeletons

8.1 The reach assertion — positive + negative (the core E2E)

```
# rig/test_reach.sh – §11.4.50 N=3 deterministic; sources the
anti-bluff lib (§11.4.1)
run_reach() {                                # $1 =
    expected: pass|deny
    pcap="qa-results/e2e/$(date +%s)_reach.pcap"
    ip netns exec client tcpdump -i overlay0 -w "$pcap" & TPID=$!
    body=$(ip netns exec client curl -s --max-time 5 http://
    10.10.0.20/ || true)
    kill "$TPID" 2>/dev/null
    if [ "$1" = pass ]; then
```

```

[ -n "$body" ] || ab_fail "B1: no body – config-only PASS
  forbidden"
tshark -r "$pcap" -Y 'tcp.flags.syn==1 && tcp.flags.ack==1' |
  grep -q . \
  && ab_pass_with_evidence "authorized reach" "$pcap" \
  || ab_fail "no SYN-ACK captured – reach not proven on the
  wire"
else
[ -z "$body" ] || ab_fail "default-deny breached: got a body"
tshark -r "$pcap" -Y 'tcp.flags.syn==1 && tcp.flags.ack==1' |
  grep -q . \
  && ab_fail "deny breached: SYN-ACK present" \
  || ab_pass_with_evidence "default-deny (no SYN-ACK)" "$pcap"
fi
}

```

8.2 Liveness with independent counter-advance (defeats B3)

```

# rig/liveness.sh – both iperf3 goodput AND the kernel WG counter
  must advance (§11.4.107(2))
wg_rx_before=$(ip netns exec client wg show overlay0 transfer |
  awk '{print $2}')
THRPUT=$(ip netns exec client iperf3 -J -t 10 -c 10.10.0.20 | jq
  '.end.sum_received.bits_per_second')
wg_rx_after=$(ip netns exec client wg show overlay0 transfer |
  awk '{print $2}')
adv=$(( wg_rx_after - wg_rx_before ))
{ awk -v t="$THRPUT" 'BEGIN{exit !(t>0)}' && [ "$adv" -gt 0 ]; } \
  && ab_pass_with_evidence "live tunnel: goodput $THRPUT, wg+$adv
  bytes" "qa-results/e2e/liveness_$(date +%s)/" \
  || ab_fail "decoy path: iperf=$THRPUT but wg counter advanced
  $adv (B3 wrong-plane)"

```

8.3 Transport escalation through a DPI block (AC4)

```

# rig/escalate.sh – plain WG dropped → core auto-escalates to
  MASQUE/H3, reach preserved
rig/dpi_block.sh
  # nft drops udp/51820, accepts udp/443
ip netns exec client helix-core connect --target 10.10.0.20 --
  status-json /tmp/status.json
sleep 6
  # allow the ladder to escalate (FSM → Connected)
transport=$(jq -r '.transport' /tmp/status.json)
[ "$transport" = "masque-h3" ] || ab_fail "did not escalate:
  transport=$transport (B3)"
pcap="qa-results/e2e/$(date +%s)_escalate.pcap"
ip netns exec censor tcpdump -i cen2gw -w "$pcap" & TP=$! # tap
  the censor→gateway path where DPI classifies (test-rig.md
  §7)
ip netns exec client curl -s --max-time 5 http://10.10.0.20/ >/
  dev/null
kill "$TP" 2>/dev/null
tshark -r "$pcap" -Y 'http3' | grep -q . \

```



```

&& ! tshark -r "$pcap" -Y 'wg' | grep -q . \
&& ab_pass_with_evidence "AC4: escalated to MASQUE/H3, no WG
signature" "$pcap" \
|| ab_fail "escalation evidence missing (H3 absent or WG
signature present)"

```

8.4 Live revoke (AC6) — peer drops from the running tunnel

```

# rig/revoke.sh — revoke a device, assert the live tunnel is
sealed < 1 s (SL03 budget)
t0=$(date +%s.%N)
helixvpctl revoke --device device-a1 # control-
plane revoke
# poll the live edge until the peer is removed (loading is
distinct, §11.4.107(3))
until ! ip netns exec gateway wg show overlay0 peers | grep -q
"$DEV_A1_PUBKEY"; do
sleep 0.05; over_budget 1.0 "$t0" && ab_fail "revoke > 1s (SL03
missed)"
done
t1=$(date +%s.%N); echo "$t0,$t1" >> qa-results/e2e/
revoke_timing.csv
# and the revoked device can no longer reach the LAN host:
ip netns exec client curl -s --max-time 3 http://10.10.0.20/ |
grep -q . \
&& ab_fail "revoked device still reaches the LAN host" \
|| ab_pass_with_evidence "AC6: revoke enforced < 1s" "qa-
results/e2e/revoke_timing.csv"

```

8a. The rig's four-layer mapping (§11.4.108) + clean-artifact discipline

Every E2E slice crosses the §11.4.108 four layers; the rig makes each layer's assertion explicit so a SOURCE-green-but-RUNTIME-dead bug cannot pass (the §11.4.108 forensic class):

Layer	E2E assertion	Defeats
SOURCE	the policy/verdict-map source compiles to the expected routes (UNIT, <u>unit.md</u> §8.1)	—
ARTIFACT	the deployed edge binary's verdict-map bytes carry the route (not just the source)	B4 stale-state
RUNTIME-on-clean-target	a fresh netns_up.sh rig, no stale tunnel/route, runs the slice	B4
USER-VISIBLE		B1, B3

Layer	E2E assertion	Defeats
	the pcap shows the SYN-ACK / the absence thereof	

Clean-artifact discipline (§11.4.108/.139, defeats B4): the rig is brought down (`netns_down.sh`) and up fresh before each E2E iteration so validation never runs against a previous deploy's running daemon (a stale tunnel, a leftover route, an old verdict-map). The runtime-signature for an E2E slice is "the captured pcap on a freshly-deployed clean rig shows the expected wire behaviour" — a pcap captured against a stale rig is **invalid** and any PASS it produces is a §11.4 PASS-bluff.

8b. Bluff-class guard map (the five forbidden E2E bluffs)

Bluff	What it would look like in E2E	The rig assertion that forbids it
B1 config-only	"AllowedIPs contains the route" with no packet	[-n "\$body"] + SYN-ACK pcap row required (§8.1)
B2 absence-of-error	"kill-switch enabled, no error"	LEAK==0 && DNS==0 packet counts on a real drop (SEC §5.7)
B3 wrong-plane	transport=="masque-h3" but WG on the wire	independent wg-counter-advance + tshark H3-present/WG-absent (§8.2/§8.3)
B4 stale-state	validating an old running daemon	fresh <code>netns_up.sh</code> per iteration (§8a)
B5 unvalidated analyzer	the tshark classifier passes its golden-bad	self-validated golden-good/bad pcap pair (§11.4.107(10), §4)

9. Open decisions surfaced for QA

#	Decision	Options	Recommendation
E-D1	WG datapath in the rig	kernel WG vs boringtun userspace	match the platform-under-test (../v02-data-plane/wireguard-core.md); run both where both ship, evidence per-datapath (§11.4.3)

#	Decision	Options	Recommendation
E-D2	DPI sim fidelity	nftables UDP-port drop (simple) vs a real DPI engine	nftables port-drop for MVP (proves AC4 escalation); a real DPI engine is a Phase-2 SEC deepening (UNVERIFIED whether GFW-class active probing is reproducible locally)
E-D3	rig host	local dev box vs the parked nezha.local heavy node (../v06-deploy/remote-testing-infra.md)	local for MVP ; remote node PARKED until greenlit — the rig is host-portable by design

UNVERIFIED: the exact `StatusReport.transport` enum value "masque-h3" and the `wg show ... transfer output column index` are taken from the overview skeleton; confirm against the FFI surface ([../v04-client/ffi-surface.md](#)) and the deployed `wg(8)` version at rig build time.

Sources verified

- [Volume-8 overview §2](#) (taxonomy row E2E), §3.1/§3.2/§3.3 (evidence taxonomy + liveness battery + self-validated analyzers), §5.3 (the rig scripts), §7.1 (G1/G2/G6), §7.2 (AC2–AC6), §9 (Makefile e2e:) — read 2026-06-26.
- Sibling specs cross-referenced: [../v02-data-plane/wireguard-core.md](#), [../v02-data-plane/transport-masque-quic.md](#), [../v02-data-plane/transport-selection-ladder.md](#), [../v02-data-plane/routing-and-addressing.md](#), [../v03-control-plane/reconciliation-flow.md](#), [../v04-client/ffi-surface.md](#), [../v06-deploy/remote-testing-infra.md](#) — filenames confirmed present 2026-06-26 (v06 doc is [GEN]-planned per MASTER_INDEX).
- Constitution: §11.4.107 (liveness / no-frozen-frame / self-validated analyzer / metamorphic), §11.4.3 (topology dispatch / SKIP), §11.4.5/.68/.69 (captured / no-absence-of-error / sink-side evidence), §11.4.50 (determinism), §11.4.108/.139 (clean-artifact runtime signature), §11.4.114 (known-good isolation), §11.4.132 (risk-order), §11.4.14 (cleanup quiescence), §11.4.161 (rootless / scoped CAP_NET_ADMIN exception), §1.1 (paired mutation) — from CLAUDE.md in-context.
- Phase-0 spike rig **[04_P0 §3]** + WBS items HVPN-P0-022/023 cited via the overview — **not independently re-fetched** (UNVERIFIED for exact item numbering; the rig design is reproduced from the overview skeleton).