

The Network Test Rig — netns topology, nftables DPI/censor sim, tc-netem, iperf3, leak harness (§11.4.169 / §11.4.69)

The Network Test Rig — netns topology, nftables DPI/censor sim, tc-netem, iperf3, leak harness (§11.4.169 / §11.4.69)

Revision: 2 Last modified: 2026-06-26T12:00:00Z

Reconciled (§11.4.35, 2026-06-26): this document's four-namespace rig (client / censor / gateway / exit, DPI in a distinct censor-ns middlebox) is the **canonical** test-rig topology; [e2e.md §2](#) was rewritten onto it. §6.1 rig/killswitch_drop.sh is the **single canonical definition** of the kill-switch driver — capture on the censor / WAN-facing egress (cen2gw); the SEC view ([security.md §9](#)) cites it rather than redefining it.

Volume 8 (Testing & QA) nano-detail specification. This document deepens the [10 §5.3] netns E2E rig and [10 §4] topology into the **reproducible network test rig** that is the substrate for every E2E / SEC / BENCH / PERF / CHAOS claim in HelixVPN: a **four-namespace Linux topology** (client-ns [?] gateway-ns [?] censor-ns [?] exit-ns), an **nftables-based DPI/censorship simulator** that reproduces each [obj §1] censorship regime, **tc-netem** impairment (loss/latency/reorder/bursty loss), **iperf3** throughput bars, and the **leak-test harness** (kill-switch / DNS / IPv6, the §11.4.69 negative-evidence triad). It is **SPEC-ONLY**: it fixes the topology, the concrete ip netns / nft / tc / iperf3 commands, the per-transport × per-regime exercise matrix (tracing to the [obj §2] survival matrix), and the captured-evidence shape — it is

not the running rig. Sources cited inline: the Volume-8 overview [10] (§5.3 rig sketch, §3 evidence model, §5.7 SEC); the obfuscation/DPI spec [obf] (v02-data-plane/obfuscation-and-dpi.md — the regime taxonomy + survival matrix); the methodology research [**research-daita_test**] (v09-research/research-daita_test.md §5 — netns, nftables kill-switch/leak, tc-netem, iperf3, RES-cited). Unproven facts are flagged UNVERIFIED: per §11.4.6.

Table of contents

- 0. Why a four-namespace rig (and the §11.4.156 local-only constraint)
 - 1. The topology
 - 2. Namespace + veth bring-up (concrete commands)
 - 3. The censor-ns: nftables DPI/censorship simulator (per regime)
 - 4. tc-netem impairment profiles
 - 5. iperf3 throughput bars + the goodput evidence
 - 6. The leak-test harness (kill-switch / DNS / IPv6)
 - 7. Per-transport × per-regime exercise matrix (traces to [obf §2])
 - 8. Self-validated analyzers + determinism + cleanup
 - 9. On-demand boot, rootlessness, and the one sanctioned sudo
 - 10. Edge cases & honest gaps
 - Sources verified
-

0. Why a four-namespace rig (and the §11.4.156 local-only constraint)

Every HelixVPN security property is a **path property** — default-deny, kill-switch, DPI-evasion, transport escalation all assert what does or does not cross a network path under an adversary. Proving them needs a real, reproducible path with a real adversary in the middle, on **one host, with no remote CI** (§11.4.156). Linux network namespaces give exactly that: isolated network stacks joined by veth links, fully reproducible, no real WAN, no shared runner [research-daita_test §5.1].


The [10 §5.3] overview used a **three**-process rig (client / gateway / connector). This document refines it into a **four-namespace** rig by splitting the adversary out of the path as its own censor-ns, so the DPI/censorship simulator is a **distinct middlebox** the packet must traverse — not an nft rule bolted onto an endpoint. This matters because the [obf] survival claims are about a *middlebox* that inspects and drops/mangles in transit; modelling

it as a separate namespace is what makes "plain WG is dropped but MASQUE survives" a genuine forwarding-path test rather than an endpoint-local filter ([obf §1], §3 here).

The four namespaces map to the real roles ([10 §5.3], [obf §1]):

- **client-ns** — runs `helix-core` (the Rust client), builds the tunnel.
- **gateway-ns** — runs `helix-edge` (the `:443` listener + probe-classifier [obf §3]) and the WG gateway.
- **sensor-ns** — the adversarial middlebox: forwards client↔gateway, runs the `nftables` DPI/censorship rules (§3) and `tc-netem` impairment (§4).
- **exit-ns** — the "authorized LAN host" / synthetic internet behind the gateway: the hello HTTP service (the reachability sink) + a controlled DNS resolver (the leak-test oracle, §6).

1. The topology



Addressing plan (fixed so captures + nft rules are stable across runs):

Link	Subnet	client side	other side
client [?] censor	10.0.1.0/30	10.0.1.2 (client-ns)	10.0.1.1 (sensor-ns)
sensor [?] gateway	10.0.2.0/30	10.0.2.1 (sensor-ns)	10.0.2.2 (gateway-ns, public :443)
gateway [?] exit	10.10.0.0/24	10.10.0.1 (gateway-ns)	10.10.0.20 (exit-ns hello/DNS sink)
overlay (tunnel inner)	fd7a:115c:a1e0::/48 ULA	client overlay addr	exit overlay addr

The sensor-ns is the **only** forwarding path client→gateway; client-ns has its default route via 10.0.1.1 (sensor-ns), which forwards to 10.0.2.2. There is no side path, so a packet the sensor drops is genuinely gone — the negative-evidence captures (§6) tap on the sensor-ns egress interface, observing exactly what would reach the real WAN.

2. Namespace + veth bring-up (concrete commands)

rig/netns_up.sh (idempotent, set -euo pipefail, sources the [10 §9] anti-bluff lib; one sanctioned sudo per §9). Snake_case per §11.4.29.

```
#!/usr/bin/env bash
# rig/netns_up.sh - bring up the 4-ns rig [research-daita_test
#                    §5.1, 10 §5.3]
set -euo pipefail

for ns in client censor gateway exit; do ip netns add "$ns" 2>/
dev/null || true; done

mkveth() { # $1 a-ns $2 a-name $3 a-cidr $4 b-ns $5 b-name $6
b-cidr
    ip link add "$2" netns "$1" type veth peer name "$5" netns "$4"
    ip -n "$1" addr add "$3" dev "$2"; ip -n "$1" link set "$2" up
    ip -n "$4" addr add "$6" dev "$5"; ip -n "$4" link set "$5" up
}
mkveth client cli2cen 10.0.1.2/30 censor cen2cli 10.0.1.1/30
mkveth censor cen2gw 10.0.2.1/30 gateway gw2cen 10.0.2.2/30
mkveth gateway gw2exit 10.10.0.1/24 exit exit2gw 10.10.0.20/24

# censor-ns forwards client<->gateway (it is the adversarial
# middlebox).
ip netns exec censor sysctl -w net.ipv4.ip_forward=1 >/dev/null
ip netns exec gateway sysctl -w net.ipv4.ip_forward=1 >/dev/null
ip -n client route add default via 10.0.1.1 # client
-> censor
ip -n censor route add 10.0.2.0/30 dev cen2gw # censor
-> gateway
ip -n gateway route add 10.0.1.0/30 via 10.0.2.1 # return
path

# exit-ns sinks: the reach target + the controlled DNS resolver
# (leak oracle §6).
ip netns exec exit python3 -m http.server 80 --bind 10.10.0.20 \
>qa-results/rig/hello.log 2>&1 & #
reachability sink
ip netns exec exit dnsmasq -k -p 53 -a 10.10.0.20 \
--no-resolv --address=/example.test/93.184.216.34 \
>qa-results/rig/dns.log 2>&1 & #
controlled resolver

echo "rig up: client(10.0.1.2) -> censor -> gateway(10.0.2.2:443)
-> exit(10.10.0.20)"
```

Teardown (rig/netns_down.sh) deletes all four namespaces (which reaps veths + child processes) and is wired to a trap '... EXIT' in every rig consumer so the host is left quiescent (§11.4.14):

```
#!/usr/bin/env bash
# rig/netns_down.sh - §11.4.14 quiescence; idempotent
set -uo pipefail
for ns in client censor gateway exit; do ip netns pids "$ns" 2>/
    dev/null | xargs -r kill 2>/dev/null || true; done
for ns in client censor gateway exit; do ip netns del "$ns" 2>/
    dev/null || true; done
```

3. The censor-ns: nftables DPI/censorship simulator (per regime)

The censor-ns reproduces each [obf §1] Regime bit as a concrete nftables ruleset on its forward hook, so a transport's [obf §2] survival claim becomes a falsifiable forwarding-path test. Each profile is a function the rig loads before the per-transport exercise (§7). nft matches sit on inet censor table, forward chain.

3.1 WG_FINGERPRINT — drop plain WireGuard by signature

The GFW/DPI key on the WG datagram: 1-byte type $\in \{1..4\}$ at offset 0 + the invariant 148 B (init) / 92 B (response) handshake sizes [obf §5.1, research-hysteria2 §4]. The simulator approximates this with a port+size match (an exact payload-byte match needs nft ... @ih raw-payload matching; the port+handshake-size form is the rig's tractable approximation, UNVERIFIED: exact byte-offset fidelity vs a real GFW classifier — the [obf §7 T-WGFP] test asserts the behavioural outcome, not classifier-internal parity):

```
# rig/regimes/wg_fingerprint.nft -> WG_FINGERPRINT
table inet censor { chain forward { type filter hook forward
    priority 0; policy accept;
    udp dport 51820 udp length 148 drop          # WG handshake_init
        invariant size [obf §5.1]
    udp dport 51820 udp length 92 drop
        # WG handshake_response invariant size
    udp dport 51820 drop
        # residual: plain WG to the canonical port
} }
# Expected: plain-udp FAILs handshake; lwo (header+size mangled,
    [obf §5.2]) survives.
```

3.2 UDP_BLOCK — hard UDP drop (only TCP-carried survive)

```
# rig/regimes/udp_block.nft -> UDP_BLOCK
table inet censor { chain forward { type filter hook forward
    priority 0; policy accept;
```

```

meta l4proto udp drop                                # ALL UDP dropped
    (incl. QUIC/443) [obf §1 UDP_BLOCK]
} }
# Expected: plain-udp/lwo/masque-h3/connect-ip ALL fail;
           shadowsocks + udp-over-tcp connect.

```

3.3 QUIC_SNI_FILTER — QUIC permitted, Initial-SNI extracted + blocklisted

The GFW decrypts the QUIC Initial (RFC 9001 header-derivable key), extracts the SNI, matches a blocklist [obf §1.1, research-hysteria2 §5(a)]. The rig models this with a helper that inspects the first UDP/443 datagram's QUIC-Initial SNI (a small userspace nfqueue classifier, since nftables alone cannot decrypt a QUIC Initial); a blocked SNI → drop, an unblocked/fronted SNI → accept:

```

# rig/regimes/quic_sni_filter.nft + rig/quic_sni_classifier.py
  (nfqueue) -> QUIC_SNI_FILTER
table inet censor { chain forward { type filter hook forward
    priority 0; policy accept;
    udp dport 443 queue num 0                # hand QUIC/443 to the
    SNI classifier
} }
# rig/quic_sni_classifier.py: derive QUIC-Initial key, extract
  SNI; verdict per blocklist.
#   - does NOT reassemble Initials split across >1 datagram ->
  models the GFW weakness
#   Gecko exploits [obf §5.3, research-hysteria2 §5(a)]; T-
  GECKO-FRAG asserts this.
# Expected: masque-h3 with PlainBlockable SNI dropped; fronted SNI
  + hysteria2/Salamander survive.

```

3.4 TLS_SNI_FILTER, ACTIVE_PROBING, RESIDUAL_BLOCK

```

# rig/regimes/tls_sni_filter.nft -> TLS_SNI_FILTER (TCP/443
  ClientHello SNI blocklist via nfqueue)
table inet censor { chain forward { tcp dport 443 queue num
    1; } } # SNI classifier on TCP

# ACTIVE_PROBING -> a scripted prober (rig/active_probe.sh)
  connects to the gateway :443
#   with forged first bytes; the edge MUST return decoy/silent/
  benign-RST [obf §3], proved by T-PROBE.

# RESIDUAL_BLOCK -> stateful: after a trigger, block the
  (src,dst,sport,dport) tuple for a window.
table inet censor {
    set blocked4 { type ipv4_addr . inet_service . ipv4_addr .
    inet_service; flags timeout; }
    chain forward { type filter hook forward priority 0; policy
    accept;
    ct state new udp dport 51820 add @blocked4 { ip saddr . udp
    sport . ip daddr . udp dport timeout 200s }

```

```

        ip saddr . udp sport . ip daddr . udp dport @blocked4 drop
        # ~180s GFW window [obf §4.4]
    } }
# Expected: mid-session drop -> ladder cools the tuple + rotates
port/escalates [obf §4.4]; T-RESIDUAL.

```

The regime profiles compose (the GFW = WG_FINGERPRINT | QUIC_SNI_FILTER | ACTIVE_PROBING | RESIDUAL_BLOCK): the rig loads the union ruleset and asserts the ladder converges on the surviving transport ([obf §4.3], §7 T-LADDER).

4. tc-netem impairment profiles

Applied **inside the censor-ns** on the cen2gw veth so only the test path is impaired [research-daita_test §5.2]. Profiles model realistic adversary/network conditions; combined with §3 nft rules they reproduce throttle + loss regimes.

```

# rig/impair.sh - tc-netem profiles on the censor egress
  (IFACE=cen2gw, in censor-ns)
ns="ip netns exec censor"
case "$1" in
  clean)      $ns tc qdisc del dev cen2gw root 2>/dev/null ||
              true ;;
  loss5)      $ns tc qdisc add dev cen2gw root netem loss 5% delay
              40ms 10ms ;; # uniform loss + jitter
  burst)      $ns tc qdisc add dev cen2gw root netem loss gmodel
              2% 10% 70% 0.1% ;; # Gilbert-Elliott bursty [research-
              daita_test §5.2]
  reorder)    $ns tc qdisc add dev cen2gw root netem delay 20ms
              reorder 25% 50% ;; # reordering
  throttle)   $ns tc qdisc add dev cen2gw root tbf rate 8mbit burst 32kb
              latency 400ms ;; # bandwidth throttle (THROTTLE regime)
  highrtt)    $ns tc qdisc add dev cen2gw root netem delay 150ms
              30ms distribution normal ;;
esac

```

The burst (Gilbert-Elliott) profile is the realistic-wireless loss model [research-daita_test §5.2] used for the resilience PERF tests (T-LOSS [obf §7]); the throttle profile reproduces the THROTTLE regime where Hysteria2/Brutal sustains goodput and MASQUE collapses (T-THROTTLE [obf §7]).

5. iperf3 throughput bars + the goodput evidence

iperf3 server in `exit-ns`, client in `client-ns` driving traffic **through the tunnel** (over the overlay), measured against the bare-link baseline [research-daita_test §5.3]. The G1 gate is $\geq 80\%$ bare-link plain-UDP; G2 is $\geq 50\%$ of plain over MASQUE [10 §7.1].

```
# rig/bench.sh - N=3 deterministic goodput bars (§11.4.50) [10 §5.13]
ip netns exec exit iperf3 -s -1 -D #
    one-shot server, exit-ns
bare=$(ip netns exec client iperf3 -J -c 10.10.0.20 -t 10 | jq
    '.end.sum_received.bits_per_second')
tun=$( ip netns exec client iperf3 -J -c <overlay-exit-addr> -t
    10 | jq '.end.sum_received.bits_per_second')
pct=$(awk -v t="$tun" -v b="$bare" 'BEGIN{printf "%.1f", 100*t/
    b}')
printf '%s,%s,%s,%s,%s\n' "$TRANSPORT" "$REGIME" "$tun" "$bare"
    "$pct" \
    >> qa-results/bench/goodput.csv #
    the evidence (min/max/mean/p95 per §11.4.24)
ab_pass_with_evidence "goodput $TRANSPORT/$REGIME = ${pct}% bare-
    link" "qa-results/bench/goodput.csv"
```

The §11.4.107 liveness battery applies: goodput is sampled over a ≥ 10 s window (not a one-shot ping), and an **independent** WG counter (`wg show <if> transfer rx/tx` advancing) must also move — a moving iperf3 with a flat WG counter is a decoy/loopback path and FAILs (B3, [10 §3.2]). The metamorphic relation ($2\times$ clients $\Rightarrow \sim 2\times$ aggregate goodput; same exit reached over WG-UDP vs MASQUE $\Rightarrow$ same body hash) is the no-golden-source oracle ([10 §3.2]).

6. The leak-test harness (kill-switch / DNS / IPv6)

The §11.4.69 negative-evidence triad — three properties whose proof is **the absence of a packet** captured on the censor-ns egress (the path to the synthetic WAN) [research-daita_test §5.4/§5.5]. Each uses a host-side pcap on `cen2gw` (the only path out) so the capture observes exactly what would reach a real ISP.

6.1 Kill-switch (S4) — seal on tunnel drop

Canonical rig/killswitch_drop.sh. This is the single source of truth for the kill-switch driver; the SEC view ([security.md §9](#)) cites this definition rather than redefining it. The capture taps the **censor / WAN-facing egress (cen2gw)** because AC7's §11.4.69 negative-evidence (zero plaintext / zero :53 leak) is only valid on the WAN path, not on the client-local interface.

```
# rig/killswitch_drop.sh - S4, the anti-bluff kill-switch test
  (defeats B2) [10 §5.7]
set -euo pipefail
pcap="qa-results/sec/$(date +%s)_killswitch.pcap"
ip netns exec censor tcpdump -i cen2gw -w "$pcap" & TPID=$!      #
  tap the path to the WAN
ip netns exec client curl -s --max-time 30 http://<overlay-exit>/
  >/dev/null & # traffic in flight
sleep 2
rig/force_tunnel_drop.sh      # core FSM -> Blocked, firewall
  seals (fwmark reject) [research-daita_test §5.4]
sleep 5
ip netns exec client nslookup example.test 2>/dev/null || true #
  try to leak a DNS query
kill "$TPID" 2>/dev/null
# PASS only if, AFTER the drop (t>2s), ZERO non-loopback packets
  AND ZERO :53 left the host:
LEAK=$(tshark -r "$pcap" -Y 'frame.time_relative>2 && ip && not
  (ip.addr==127.0.0.1)' | wc -l)
DNS=$( tshark -r "$pcap" -Y 'frame.time_relative>2 &&
  udp.port==53' | wc -l)
[ "$LEAK" -eq 0 ] && [ "$DNS" -eq 0 ] \
  && ab_pass_with_evidence "kill-switch sealed, no DNS leak"
  "$pcap" \
  || ab_fail "LEAK=$LEAK DNS=$DNS - plaintext/DNS escaped the
  seal"
```

The kill-switch mechanism is the WireGuard fwmark pattern: tunnel traffic is tagged with an fwmark, an OUTPUT rule rejects any egress **not** matching it, so a tunnel drop fails closed [research-daita_test §5.4].

6.2 DNS leak (S5) — queries forced through the tunnel resolver

The exit-ns dnsmasq (§2) is the **only** legitimate resolver; the censor-ns has no resolver. A DNS query observed on cen2gw to anything other than the tunnelled resolver is a leak. Defense-in-depth per [research-daita_test §5.5]: tunnel DNS config → nftables drop of :53 outside the tunnel → server-side resolver. The test asserts **zero** udp.port==53 || tcp.port==53 on the non-tunnel path during a resolve + during a reconnect.

6.3 IPv6 leak (S5b) — no v6 escapes the seal

IPv6 is the common gap [research-daita_test §5.5]: either route v6 through the tunnel or block v6 egress at nftables. The rig asserts **zero** IPv6 packets on cen2gw when the policy is v6-block, or all v6 inside the overlay when v6-tunnelled:

```
# IPv6-leak assertion (the censor-ns blocks/observes v6 egress)
V6=$(tshark -r "$pcap" -Y 'ipv6 && not (ipv6.addr==::1)' | wc -l)
[ "$V6" -eq 0 ] && ab_pass_with_evidence "no IPv6 leak" "$pcap"
|| ab_fail "IPv6 leaked: $V6 pkts"
```

Each leak test is self-validated (§8): a golden-bad pcap with a seeded :53 / v6 packet MUST score FAIL, else the detector is the bluff (§11.4.107(10), [10 §3.3]).

7. Per-transport × per-regime exercise matrix (traces to [obf §2])

This is the rig's load-bearing artifact: every [obf §2.1] survival-matrix cell is a rig run (transport, regime) whose outcome is captured. The rig drives each transport under each nft regime profile (§3) + netem profile (§4) and asserts the expected connect/fail, tracing each expectation to the [obf] cell-source. ✓ = expected-connect (capture goodput), ✗ = expected-fail (capture the DPI drop).

Transport	empty	WG_FINGERPRINT	UDP_BLOCK	QUIC_SNI_FILTER (plain SNI)	QUIC_S (fronto)
plain-udp	✓	✗	✗	✗ (not QUIC; UDP path)	✗
lwo	✓	✓ (header+size mangle)	✗	n/a	n/a
masque-h3	✓	✓ (no WG sig on wire)	✗	✗ (SNI blockable)	✓ (fronto)
shadowsocks	✓	✓	✓ (TCP)	n/a	n/a
udp-over-tcp	✓	✓	✓ (TCP)	n/a	n/a
hysteria2+Salamander ¹	✓	✓			✓

Transport	empty	WG_FINGERPRINT	UDP_BLOCK	QUIC_SNI_FILTER (plain SNI)	QUIC_SNI_FILTER (fronted)
			× (needs UDP)	✓ (SNI scrambled)	

¹ Phase-2 / feature-gated ([obf §6.3]); the MVP rig exercises plain-udp, masque-h3, shadowsocks, udp-over-tcp, lwo (the MVP transport set [10 §2]).

The UDP-block carve-out (load-bearing, [obf §2.2]). Under UDP_BLOCK every UDP-carried transport (plain-udp, lwo, masque-h3, connect-ip, hysteria2) **must** fail — the rig asserts this explicitly so a mask bug cannot claim the impossible (§11.4.6). Only the TCP-carried shadowsocks / udp-over-tcp connect.

The ladder run (T-LADDER, [obf §4.3]). The rig loads the composed GFW ruleset (WG_FINGERPRINT | QUIC_SNI_FILTER | ACTIVE_PROBING | RESIDUAL_BLOCK) and asserts the detector-driven ladder dials plain-udp (drops) → lwo (drops) → masque-h3 fronted (**connects**), commits the QUIC_SNI_FILTER regime estimate, and pins masque-h3 in per-network memory — the captured TunnelStatus trace is the evidence, run -count=3 deterministic (§11.4.50, [obf §7 T-LADDER]).

```
# rig/exercise.sh - drive one (transport, regime) cell and capture
load_regime "$REGIME" # rig/regimes/
    <regime>.nft + nfqueue classifier
load_impair "$NETEM" # §4
pcap="qa-results/e2e/${TRANSPORT}_${REGIME}.pcap"
ip netns exec censor tcpdump -i cen2gw -w "$pcap" & TPID=$!
if rig/dial.sh "$TRANSPORT"; then status=connected; else
    status=failed; fi
kill "$TPID" 2>/dev/null
assert_outcome "$TRANSPORT" "$REGIME" "$status" "$pcap"
    # checks against the §7 matrix; ab_pass_with_evidence
```

8. Self-validated analyzers + determinism + cleanup

Every rig analyzer (the pcap classifier, the leak detector, the QUIC-SNI classifier, the goodput parser) ships with a **golden-good + golden-bad** fixture pair wired into the meta-test sweep (§11.4.107(10), [10 §3.3]) — an analyzer that passes its golden-bad fixture is itself the bluff and FAILs CM-ANALYZER-SELF-VALIDATED:

Analyzer	golden-good (PASS)	golden-bad (MUST FAIL)	mutation
leak detector	killswitch_clean.pcap	killswitch_leaked_dns.pcap (seeded :53)	strip the udp.port==53 clause → golden-bad passes → meta-test FAILs
WG-fingerprint sim	wg_blocked.pcap (plain WG dropped)	wg_passed_through.pcap	weaken the nft size match → WG passes → T-WGFP FAILs
QUIC-SNI classifier	sni_blocked.pcap	sni_missed.pcap (blocklisted SNI passed)	break the SNI extractor → blocked SNI passes → T-QUIC-SNI FAILs
goodput parser	iperf_live.json (advancing)	iperf_frozen.json (flat counter)	accept the flat counter → B3 → FAILs

Determinism (§11.4.50). Every rig PASS runs N=3 (normal) / N=10 (cycle-validation) against the same artifact MD5 + same rig, identical evidence-hashes; a divergent run is auto-FAIL via `ab_run_n_times` ([10 §5.4]). The fixed addressing plan (§1) + fixed regime rulesets (§3) make the rig reproducible.

Cleanup (§11.4.14). Every rig consumer wires `trap 'rig/netns_down.sh' EXIT` so the four namespaces + their child processes (`http.server`, `dnsmasq`, `iperf3`, `tcpdump`) are reaped on every exit path; the orchestrator post-test sanity-checks that no `client/censor/gateway/exit` namespace survives and FAILs the just-completed test if one does (§11.4.14 orphan-state guard).

9. On-demand boot, rootlessness, and the one sanctioned sudo

Per [10 §9] the rig boots **on demand** — make `e2e` runs `rig/netns_up.sh`, the exercise, then `rig/netns_down.sh`; the developer never hand-creates namespaces. The container-backed infra the INT layer needs (Postgres +

Redis for the control plane behind the gateway) boots rootless via the containers submodule (containers/pkg/boot, §11.4.76/.161) — no rootful Docker, no container-management sudo.

The **only** sudo in the whole harness is rig/netns_up.sh / rig/netns_down.sh: network-namespace + veth creation needs CAP_NET_ADMIN, a documented, scoped exception ([10 §9]) — never a container-management escalation, never host-power, never disabling session managers (§12).
UNVERIFIED: whether rootless netns via a user namespace (ip netns inside a unshare -Un) fully substitutes for the CAP_NET_ADMIN path on every target host — the spec records the privileged path as canonical and the rootless-netns path as a Phase-1 investigation (§11.4.81 cross-platform-parity adjacency for macOS, where netns does not exist and the rig SKIPS with topology_unsupported per §11.4.3).

10. Edge cases & honest gaps

#	Edge case	Required behavior	Cite
R1	nft size-match approximates, not byte-exact, the GFW WG classifier	assert the behavioural outcome (plain-udp drops, lwo survives), flag the classifier-internal parity UNVERIFIED:	§3.1, §11.4.6
R2	QUIC-Initial SNI needs decryption nftables can't do	offload to an nfqueue userspace classifier; model the GFW no-reassembly weakness explicitly (Gecko relies on it)	§3.3, [obf §5.3]
R3	A regime profile leaves a stale nft table between runs	load_regime flushes inet_censor first; teardown deletes the namespace (drops the table)	§3, §8
R4	iperf3 measures the bare link, not the tunnel, by mistake	the client targets the overlay exit addr; the independent WG counter must also advance (B3 guard)	§5
R5	A leak test "passes" because no traffic was generated at all	the test asserts traffic was in flight before the drop (a non-empty pre-drop capture window) — absence-of-traffic ≠ no-leak	§6, §11.4.68
R6	macOS / non-Linux host has no netns	honest §11.4.3 SKIP: topology_unsupported + tracked migration item; never a fake PASS	§9, §11.4.81
R7	The host lacks CAP_NET_ADMIN	SKIP: topology_unsupported; the rig cannot run unprivileged on	§9

#	Edge case	Required behavior	Cite
		that host (investigate rootless-netns)	

Honest boundary (§11.4.6). The rig proves the **path-level** security and survival properties on a Linux netns substrate that *models* each censorship regime; it does **not** prove behaviour against a *real* nation-state DPI deployment (the nft/nfqueue simulator is a faithful behavioural model, not the GFW itself — classifier-internal parity is UNVERIFIED: and the [obf §7] tests assert the behavioural outcome, not byte-for-byte adversary parity). The on-wire DPI-resistance claims (a masque-h3 flow classifies as HTTP/3 with no WG signature) are additionally proven by the tshark-on-capture wire-fingerprint check + a window-scoped MP4 vision verdict (T-G2-WIRE [obf §7], §11.4.159). Real-world DPI-evasion confidence ultimately rests on field evidence the rig cannot supply — the rig empties the *reproducible-regime* defect class, not the *all-adversaries* class.

Sources verified

- [10] docs/research/mvp/final/10-testing-acceptance-and-qa.md §3 (evidence model)
 - liveness battery), §4 (pyramid/topology), §5.3 (netns E2E rig), §5.7 (SEC invariants + killswitch_drop.sh), §5.13 (bench.sh), §9 (local gate harness + sanctioned sudo) — read 2026-06-26.
- [obf] docs/research/mvp/final/v02-data-plane/obfuscation-and-dpi.md §1 (regime taxonomy bits), §2 (regime → transport survival matrix + UDP-block carve-out), §3 (probe-classifier / decoy edge), §4 (regime-aware ladder + residual-block), §5 (wire byte-layouts), §7 (T-id test points) — read 2026-06-26.
- [research-daita_test] docs/research/mvp/final/v09-research/research-daita_test.md §5.1 (netns + veth topology), §5.2 (tc-netem incl. Gilbert-Elliott bursty loss), §5.3 (iperf3), §5.4 (nftables kill-switch fwmark + DNS-leak port-drop + DPI sim), §5.5 (DNS/kill-switch/IPv6 leak testing + defense-in-depth order) — read 2026-06-26 (all sources access-dated 2026-06-25 in the dossier).
- Constitution: §11.4.69 (sink-side negative-evidence taxonomy), §11.4.68 (positive sink-side / no absence-of-error), §11.4.107 (liveness battery + self-validated analyzers), §11.4.50 (determinism), §11.4.14 (cleanup/quiescence), §11.4.156 (no remote CI — local rig), §11.4.76/.161 (containers/rootless), §11.4.81 (cross-platform parity — macOS netns SKIP), §11.4.169 (test-type mandate) — CLAUDE.md, read 2026-06-26.